

Course MongoDB

Topics (1)

Database models

SQL vs NoSQL

Introduction to MongoDB

CRUD operation

How to design schema of MongoDB

Improve performance

Topics (2)

Administration tasks
Scalable
Security
Monitoring and Operation

Database Models

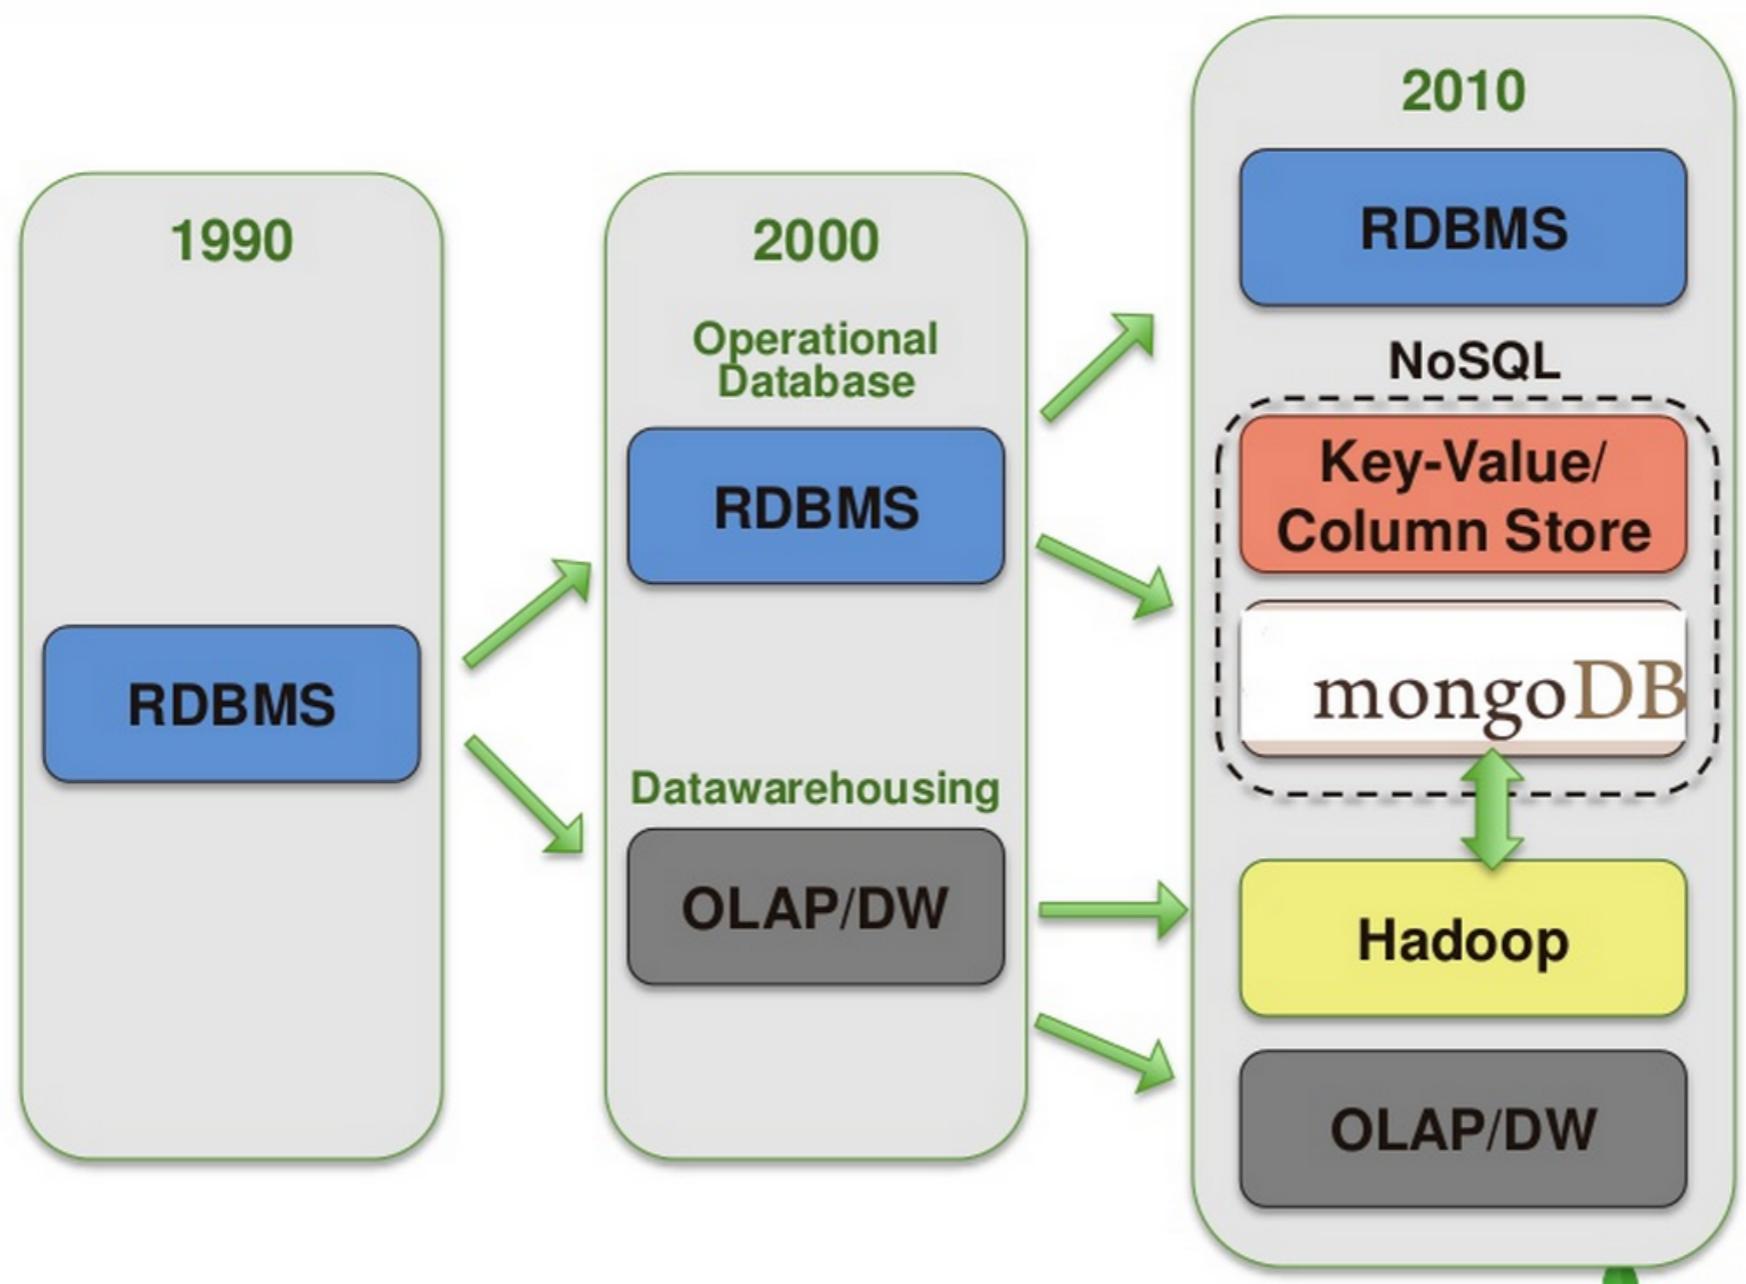
Database ranking

381 systems in ranking, November 2021

Rank			DBMS	Database Model	Score		
Nov 2021	Oct 2021	Nov 2020			Nov 2021	Oct 2021	Nov 2020
1.	1.	1.	Oracle	Relational, Multi-model	1272.73	+2.38	-72.27
2.	2.	2.	MySQL	Relational, Multi-model	1211.52	-8.25	-30.12
3.	3.	3.	Microsoft SQL Server	Relational, Multi-model	954.29	-16.32	-83.35
4.	4.	4.	PostgreSQL	Relational, Multi-model	597.27	+10.30	+42.22
5.	5.	5.	MongoDB	Document, Multi-model	487.35	-6.21	+33.52
6.	6.	↑ 7.	Redis	Key-value, Multi-model	171.50	+0.15	+16.08
7.	7.	↓ 6.	IBM Db2	Relational, Multi-model	167.52	+1.56	+5.90
8.	8.	8.	Elasticsearch	Search engine, Multi-model	159.09	+0.84	+7.54
9.	9.	9.	SQLite	Relational	129.80	+0.43	+6.48
10.	10.	10.	Cassandra	Wide column	120.88	+1.61	+2.13

<https://db-engines.com/en/ranking>

Evolution of database

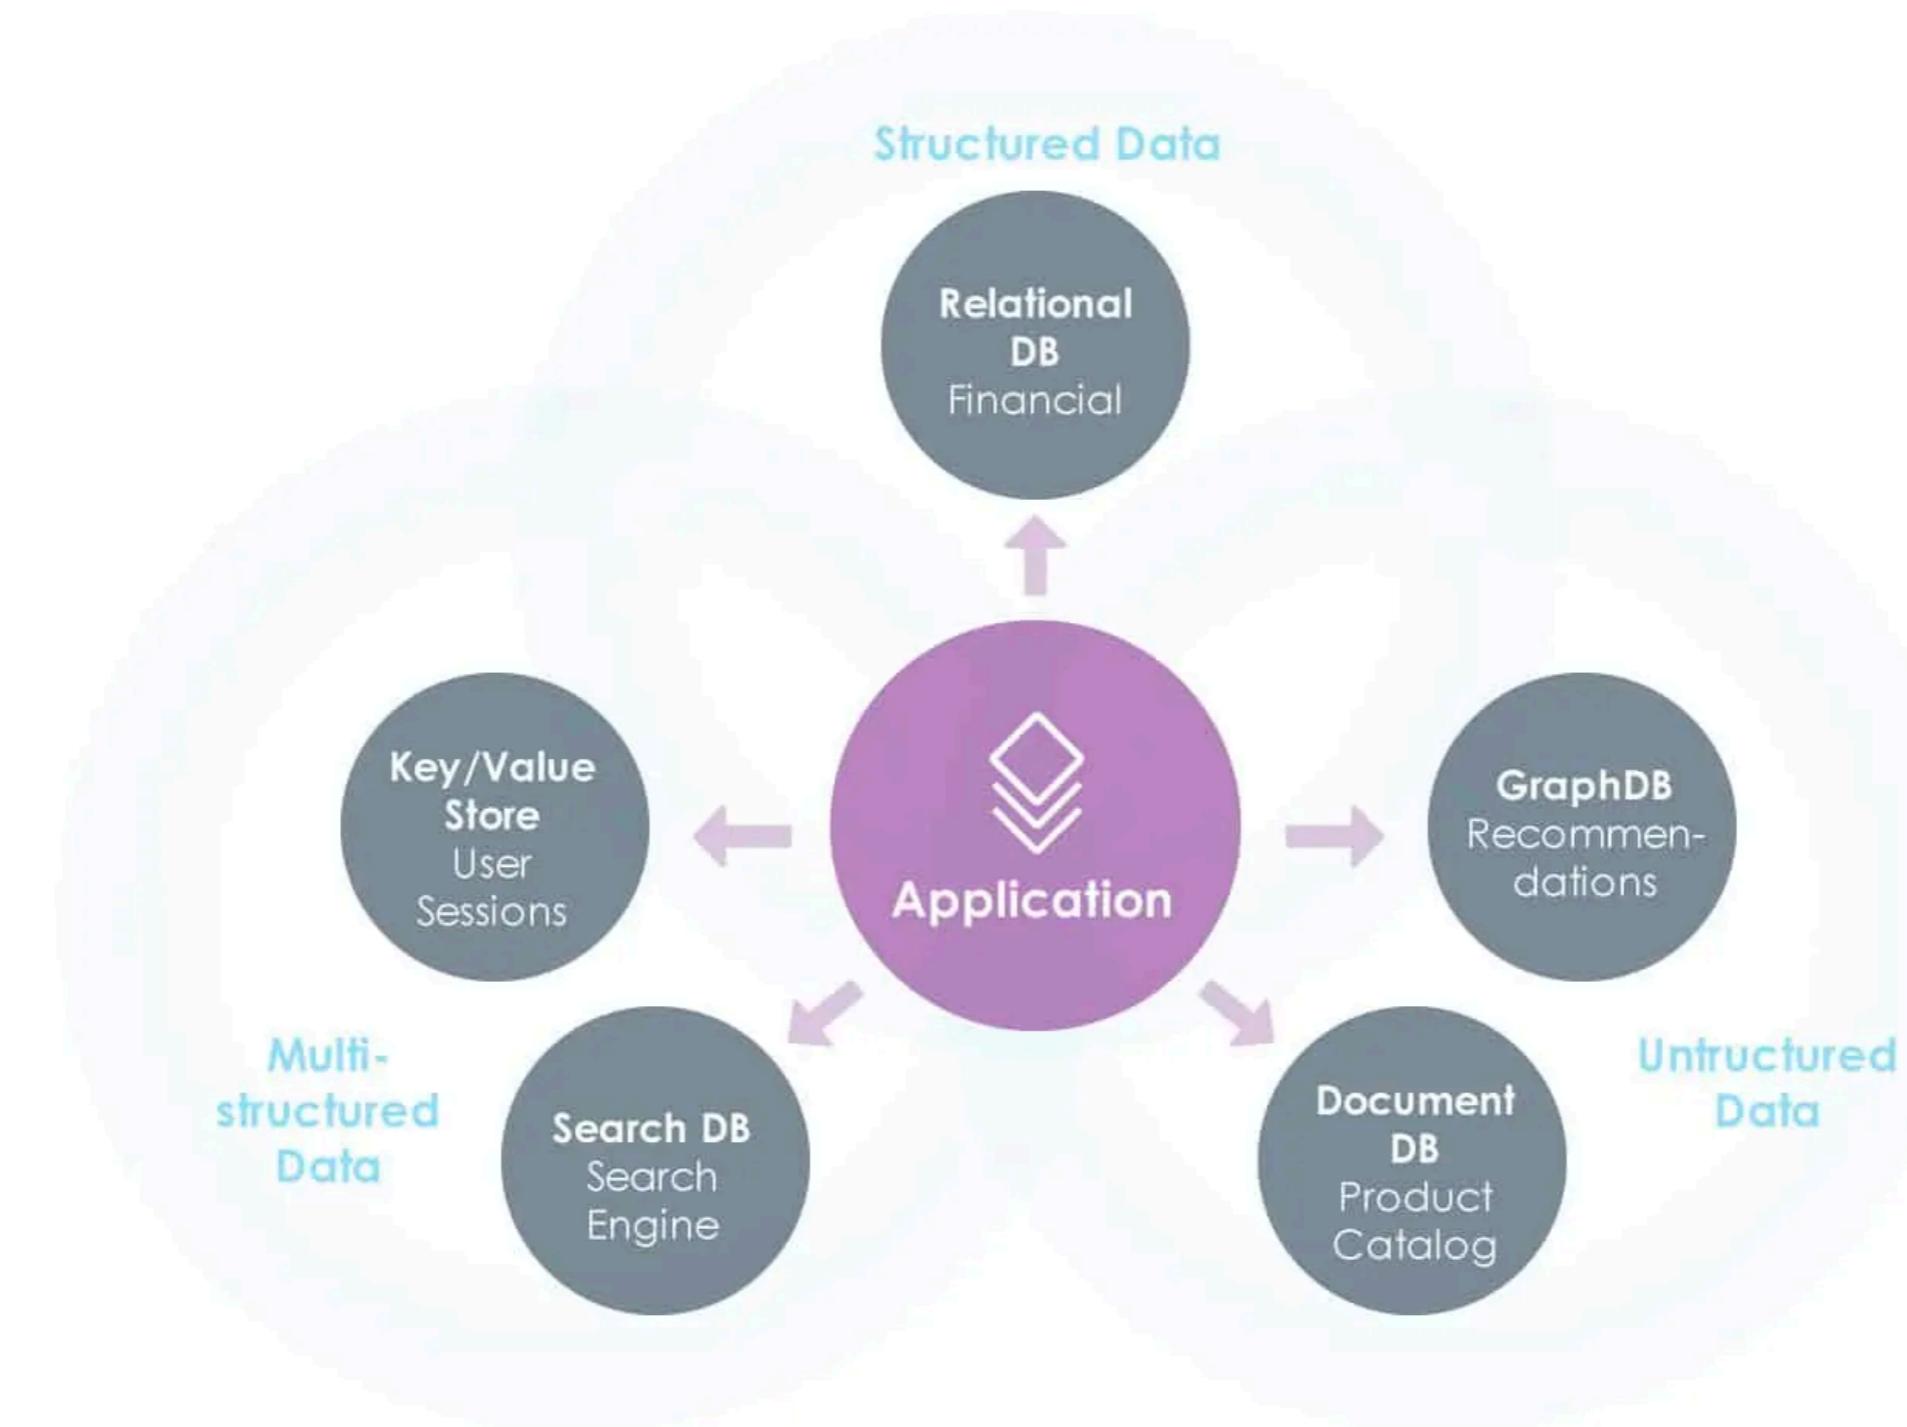


Polyglot Database



<https://martinfowler.com/bliki/PolyglotPersistence.html>

Polyglot Database



NoSQL

Not Only SQL

Non-relation

No join

No complex transaction

Focus on horizontal scalable

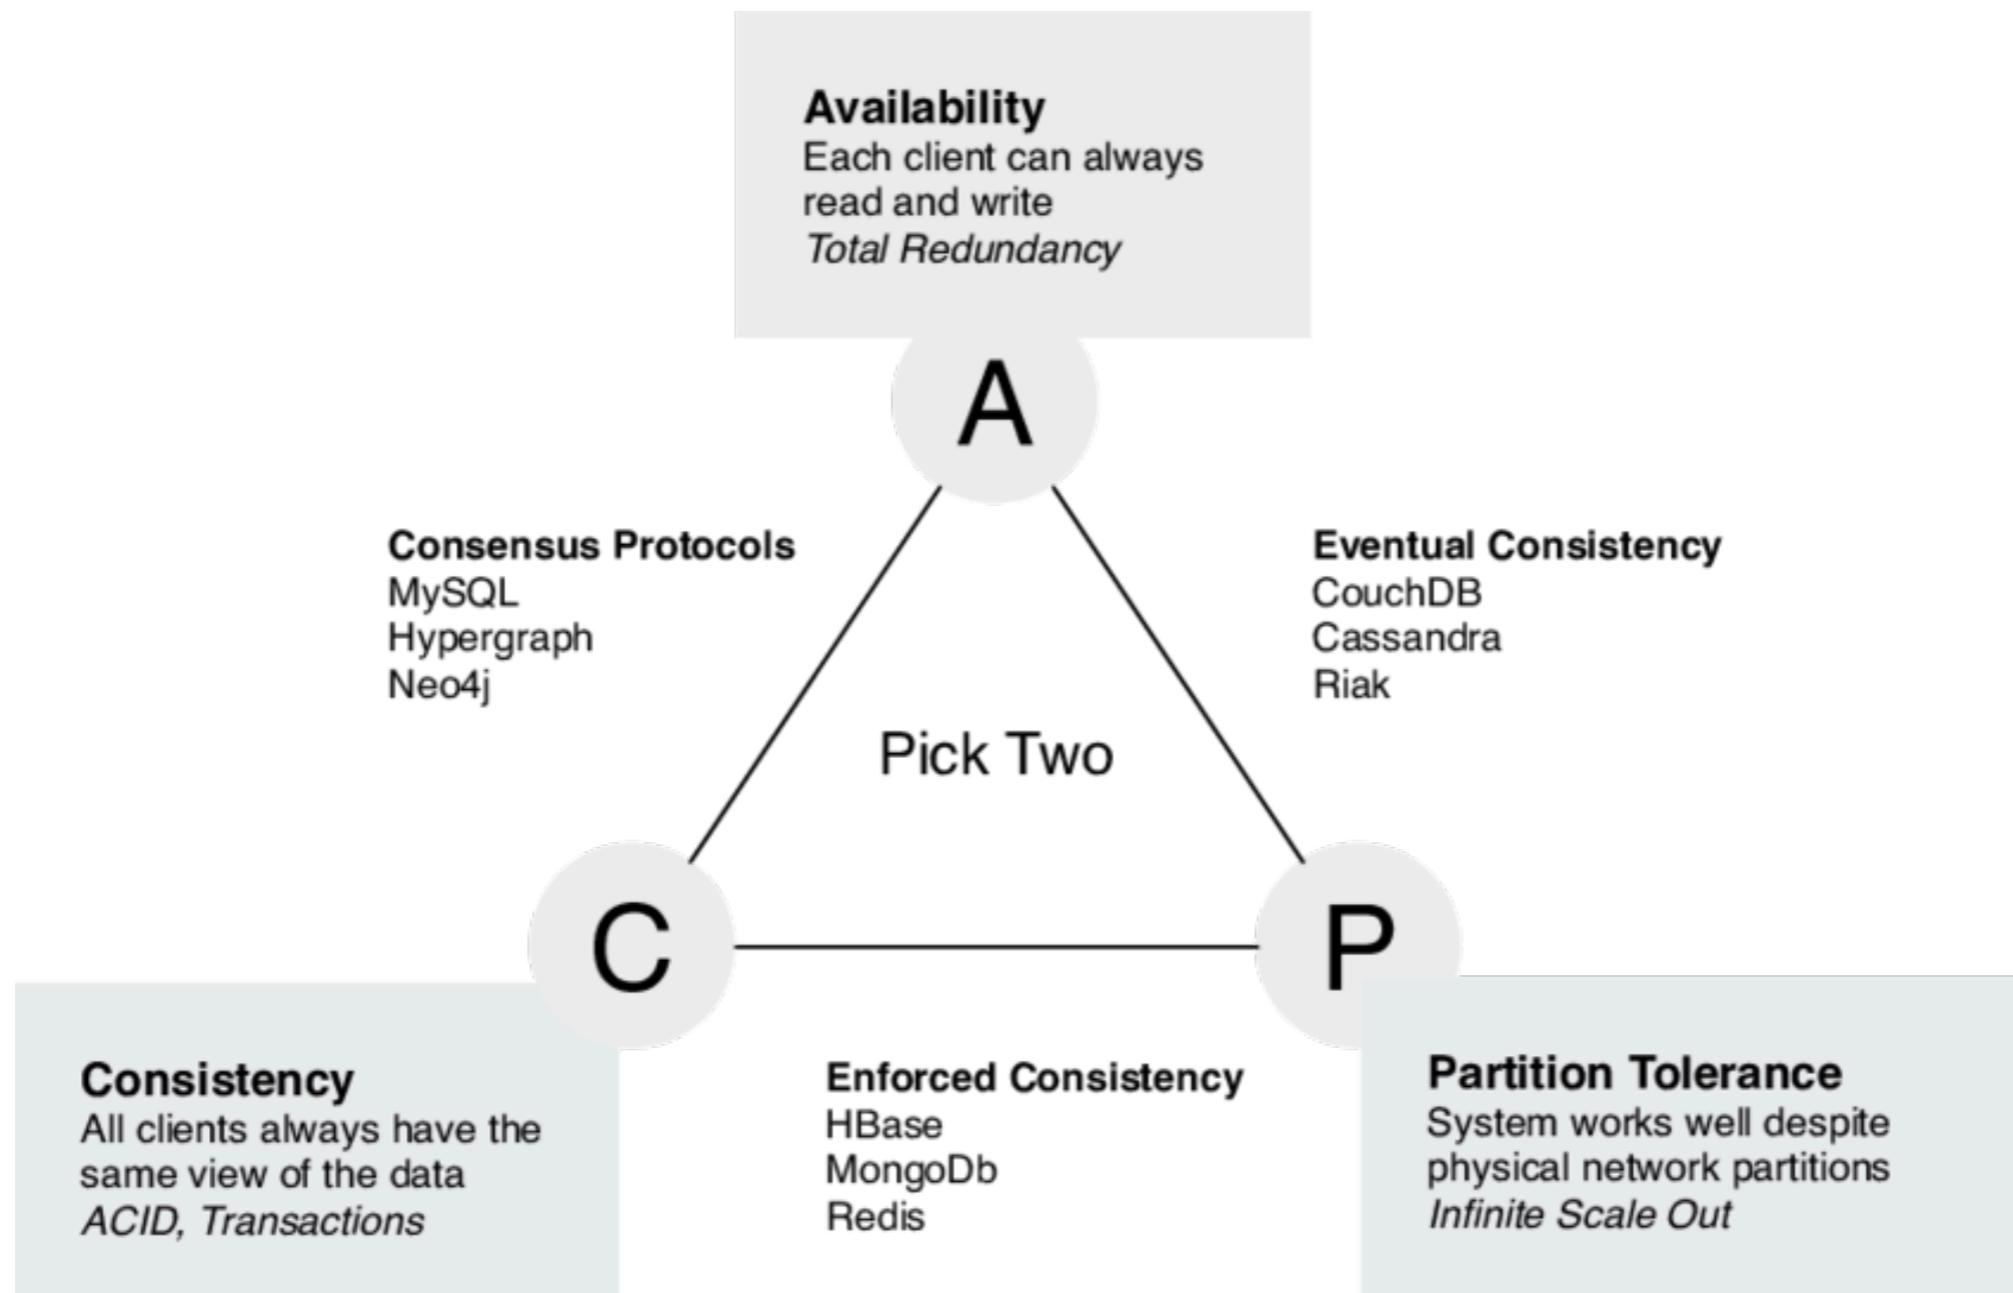
Need new database model

NoSQL World

Document Database	Graph Databases
  	 
Wide Column Stores	Key-Value Databases
  	    Amazon SimpleDB

@cloudbtxt <http://www.aryannava.com>

CAP Theorem



MongoDB

MongoDB

Humongous

hu·mon·gous
/(h)yoo'ಮäNGgəs,-'məNG-/

adjective NORTH AMERICAN *informal*

huge; enormous.
"a humongous steak"

MongoDB

Open sources
Scalable
High-performance
Schema-free, dynamic schema
Document-oriented database

MongoDB

NoSQL documentDB
Stored in collections

Document

Way to organize and store data
Set of field-values pairs

```
{  
  "id": 1,  
  "name": "mongodb",  
  "age": 30  
}
```

JSON format

JavaScript Object Notation

Start and end with { }

Separate each key and value with :

Separate each key-value with ,

Text-based format

JSON use cases

APIs

Configuration files

Log message

Database storage

JSON as database

Text-based format (parsing is very slow)

Less space-efficient

Only support a limited number of data types

BSON

Binary JSON

Bridges the gap between binary and JSON
format

<https://bsonspec.org/>

BSON Goals

Speed
Space
Flexibility

Designed for high performance

JSON vs BSON

	JSON	BSON
Encoding	UTF-8 string	Binary
Data types	String, Boolean, Number, Array	String, Boolean, Number, Array, Date, Binary
Readability	Human and Machine	Machine only

<https://www.mongodb.com/json-and-bson>

Summary

BSON => store

MongoDB stores data in BSON, internally and over the network

JSON => view

Can be natively stored and retrieved in MongoDB

<https://bsonspec.org/>

RDBMS vs MongoDB

Database

Table

Row

Column

Database

Collection

Document
(JSON, BSON)

Field

<https://docs.mongodb.com/manual/reference/sql-comparison/>

Collection in MongoDB

An organized store of documents

Collection name = order

Document 1

Document 2

Document 3

Collections types

Default
Views

On-demand material views
Capped collection (fixed-size)
Time series collection

Installation

Choices

On premise

On cloud

On premise

The screenshot shows the MongoDB website's deployment section. At the top, there are four options: 'Atlas' (MongoDB as a service), 'On-premises' (MongoDB locally, which is selected and highlighted with a green border), 'Tools' (Boost productivity), and 'Mobile & Edge' (Realm Datastore). Below this, a dropdown menu is open for 'MongoDB Enterprise Server'. Underneath, the 'MongoDB Community Server' section is visible, featuring a brief description of its capabilities and a note about its availability as a managed service via MongoDB Atlas. A 'Download' button is located in the 'Available Downloads' sidebar, which also includes dropdown menus for Version (5.0.3 current), Platform (macOS), and Package (tgz). The sidebar also lists 'Current releases & packages', 'Development releases', 'Archived releases', 'Changelog', and 'Release Notes'. A small circular icon with a speech bubble is in the bottom right corner.

MongoDB. Products Solutions Resources Company Pricing

Choose which type of deployment is best for you

Atlas MongoDB as a service

On-premises MongoDB locally

Tools Boost productivity

Mobile & Edge Realm Datastore

MongoDB Enterprise Server

MongoDB Community Server

The Community version of our distributed database offers a flexible document data model along with support for ad-hoc queries, secondary indexing, and real-time aggregations to provide powerful ways to access and analyze your data.

The database is also offered as a fully-managed service with [MongoDB Atlas](#). Get access to advanced functionality such as auto-scaling, serverless instances (in preview), full-text search, and data distribution across regions and clouds. Deploy in minutes on AWS, Google Cloud, and/or Azure, with no downloads necessary.

Give it a try with a free, highly-available 512 MB cluster.

Available Downloads

Version 5.0.3 (current)

Platform macOS

Package tgz

Download Copy Link

Current releases & packages

Development releases

Archived releases

Changelog

Release Notes

<https://www.mongodb.com/try/download/community>

On cloud (Atlas)

The screenshot shows the MongoDB Atlas web interface. At the top, there's a navigation bar with icons for demo, Access Manager, Billing, All Clusters, Get Help, and a user account named somkiat. Below the navigation is a secondary header with a folder icon for demo01, a three-dot menu, an Atlas icon, a Realm icon, a Charts icon, and user profile icons.

The main content area has a sidebar on the left with sections for DEPLOYMENT (Databases selected), SECURITY (Database Access, Network Access, Advanced), and a bottom section with a back arrow icon.

The main panel title is "Database Deployments" under "DEMO > DEMO01". It features a search bar with "Find a database deployment..." placeholder text. In the center is a large green icon of two stacked cylinders with a plus sign, labeled "Create a database". Below it is the text "Choose your cloud provider, region, and specs." and a prominent green button labeled "Build a Database".

At the bottom of the main panel, a note says: "Once your database is up and running, live migrate an existing MongoDB database into Atlas with our [Live Migration Service](#)". To the right of this note is a circular notification icon with a red border and the number "2".

CRUD operations

In mongo shell

CRUD operations

Operation	Example
Create	<code>insertOne()</code> <code>insertMany()</code>
Read	<code>find()</code> - criteria, limit, sorting
Update	<code>updateOne()</code> <code>updateMany()</code> <code>replaceOne()</code>
Delete	<code>deleteOne()</code> <code>deleteMany()</code>
Upsert (Update or insert)	<code>updateOne(... "upsert": true)</code>

Build write operations

Write one or more documents

Transaction per document

Ordered vs Unordered operation

Use bulkWrite()

Upsert ?

If upsert is **true**

YES

Is there a match?

Update the matched document

NO

Is there a match?

Insert a new document

Workshop

MongoDB Query Language (MQL)

Query and projection operator
Update operator
Comparison operator
Aggregation pipeline

Comparison operator

Operator	Description
\$eq	Equal to
\$ne	Not equal to
\$gt	Greater than
\$gte	Greater than or equal
\$lt	Less than
\$lte	Less than or equal

Update operator

Modify data in database
\$inc, \$set and \$unset

Aggregation pipeline

Another way to query data in MongoDB

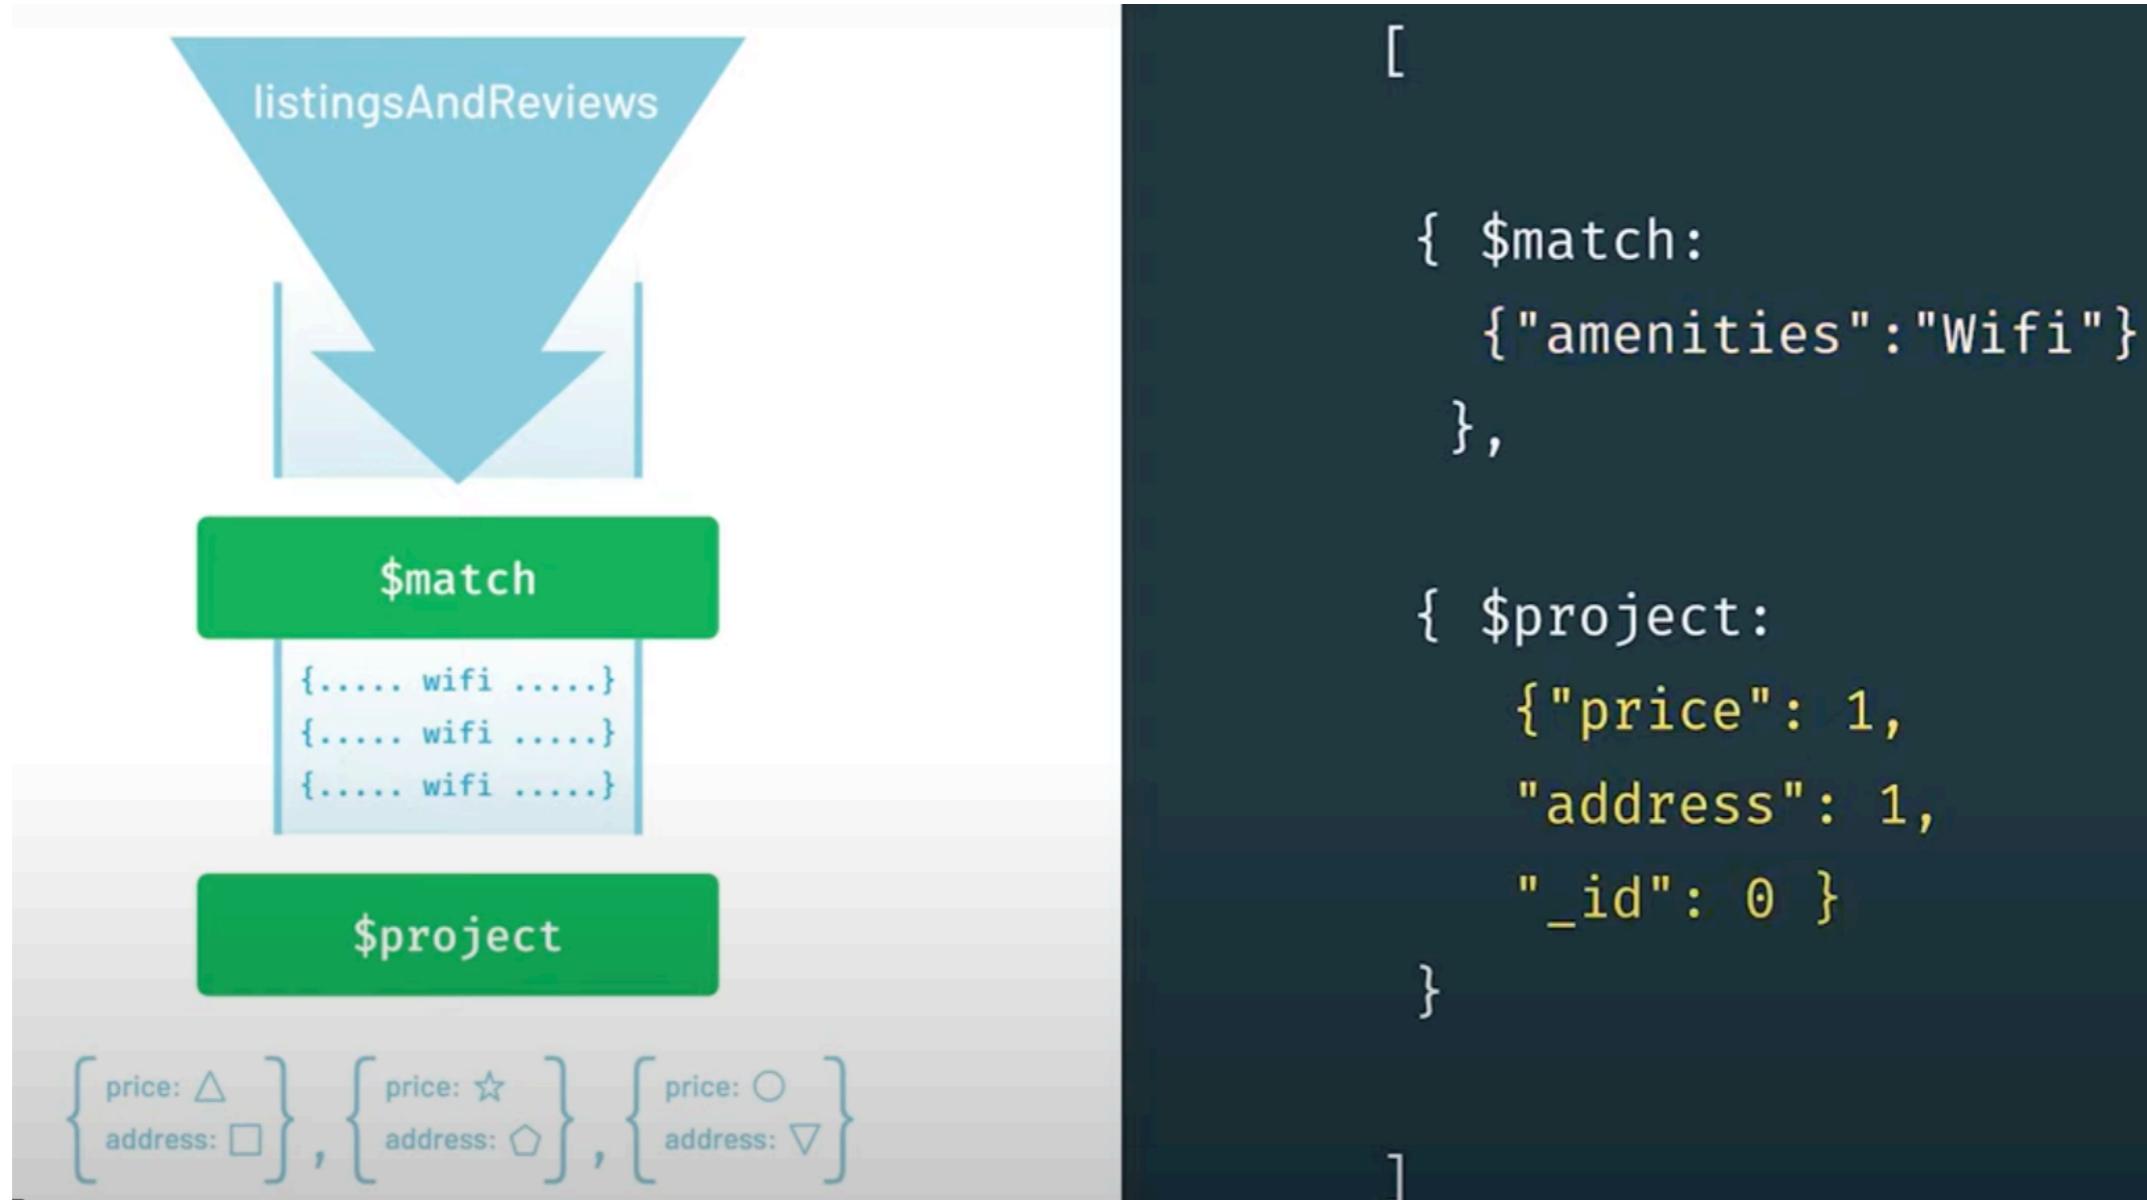
Powerful to transform data

Break down a complex query in to easier stages



<https://docs.mongodb.com/manual/core/aggregation-pipeline/>

Aggregation pipeline



Syntax

```
Cluster0-shard-0:PRIMARY> db.solarSystem.aggregate([{
...   $match: {
...     atmosphericComposition: { $in: [/O2/] },
...     meanTemperature: { $gte: -40, $lte: 40 }
...   }
... }, {
...   $project: {
...     _id: 0,
...     name: 1,
...     hasMoons: { $gt: ["$numberOfMoons", 0] }
...   }
... }], { allowDiskUse: true}
... )
```

<https://docs.mongodb.com/manual/meta/aggregation-quick-reference/>

Workshop

Data modeling

<https://docs.mongodb.com/manual/core/data-modeling-introduction/>

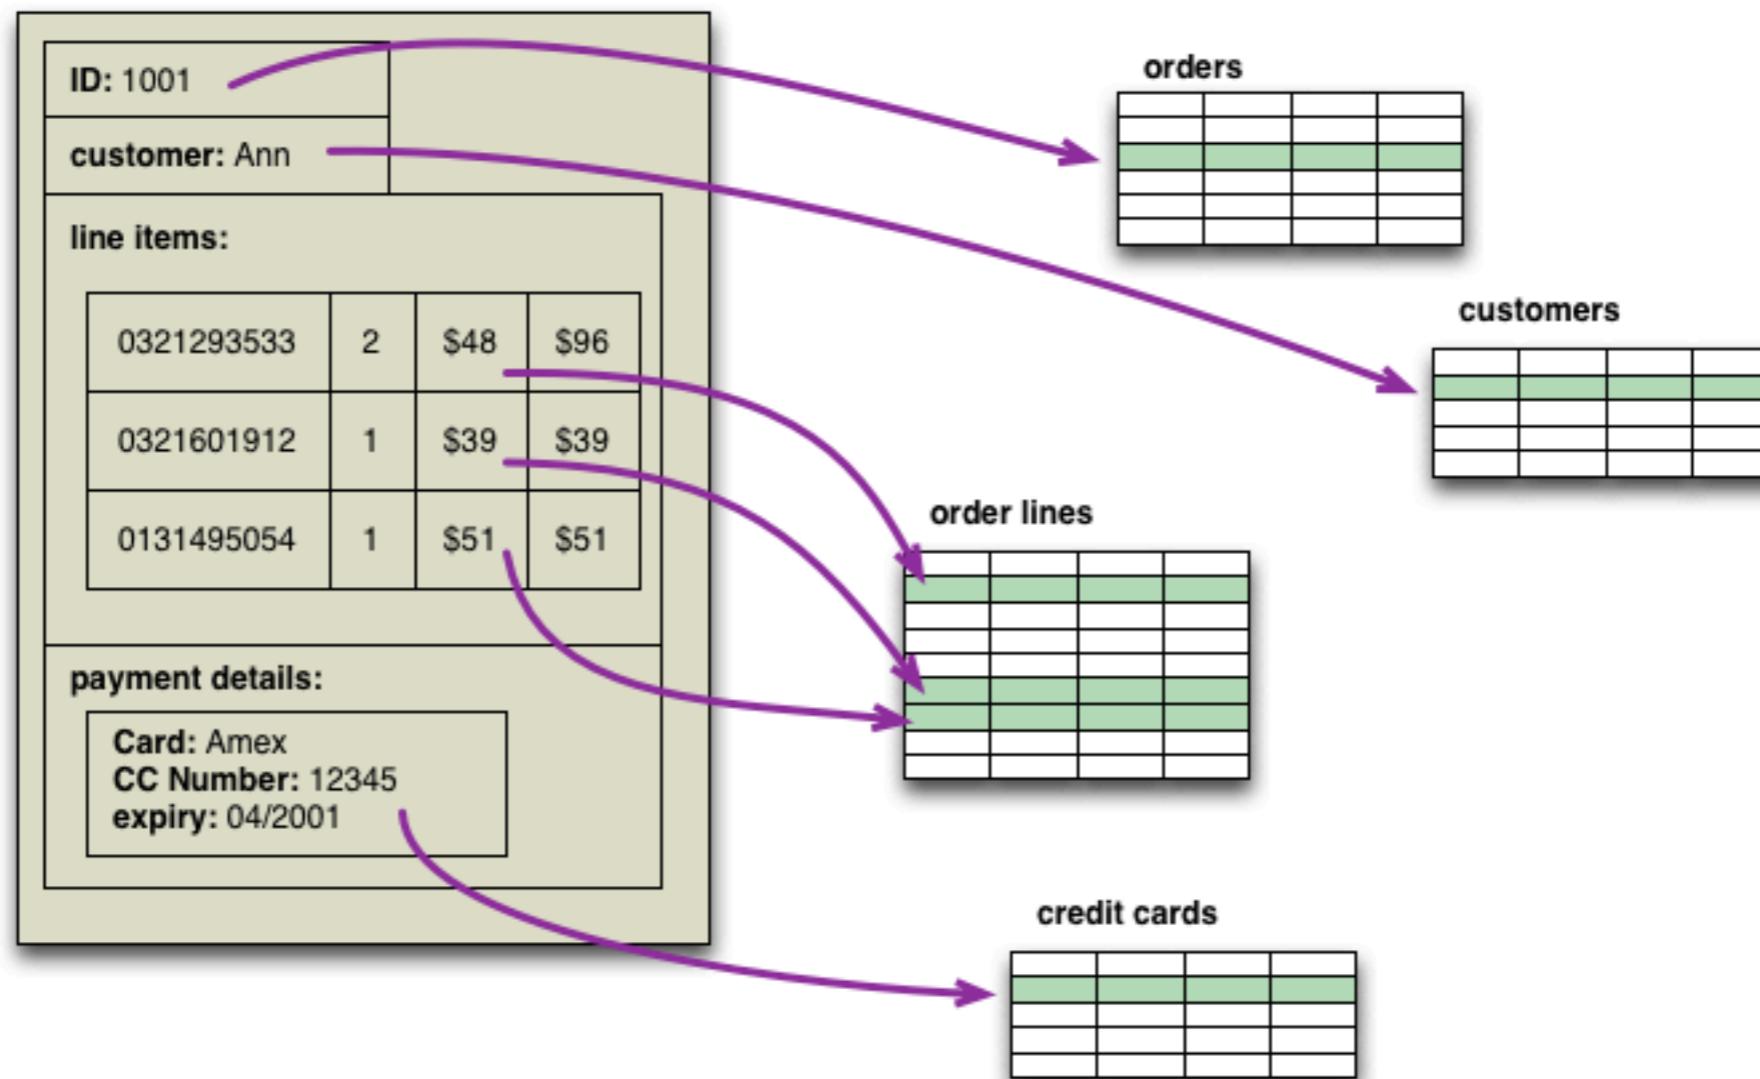
Data modeling in MongoDB

Way to organize fields in a document
To support your application **performance**
And **query** capabilities

Design for read !!

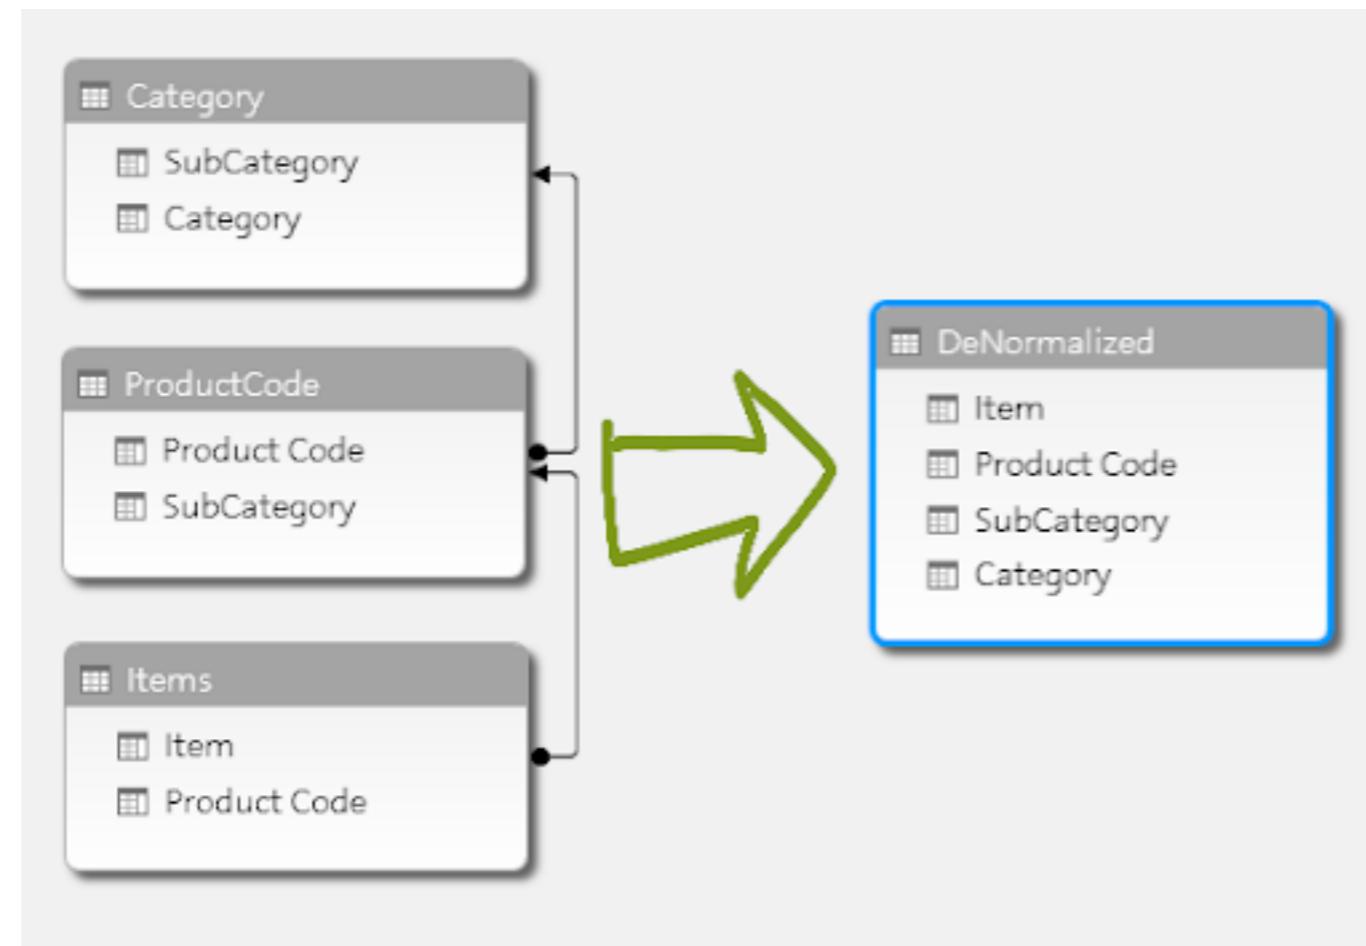
Normalization

Design for write or read ?



Denormalization

Design for write or read ?



Comparison

Normalization	Denormalization
Non-redundant and consistency data	Combine data that can be query more speed
More tables and joins	Decrease tables and joins
Optimized usage	Wastage
No redundant data	More redundant data
Design for write	Design for read

MongoDB Rule

*Data is stored in the way
that it is used*

More questions ?

How it will be queried ?

Who is using our application ?

Ratio between reads and writes ?

Design for users

Keep in mind

Data that is used together should be stored together

Evolving application implies an evolving data model too

Performance issues

Data modeling phases

Workload (data size, read/write)

Relationships

Patterns (optimization)

Document structures

Embedded data
References

Relationships

One-to-one

One-to-many

Many-to-many

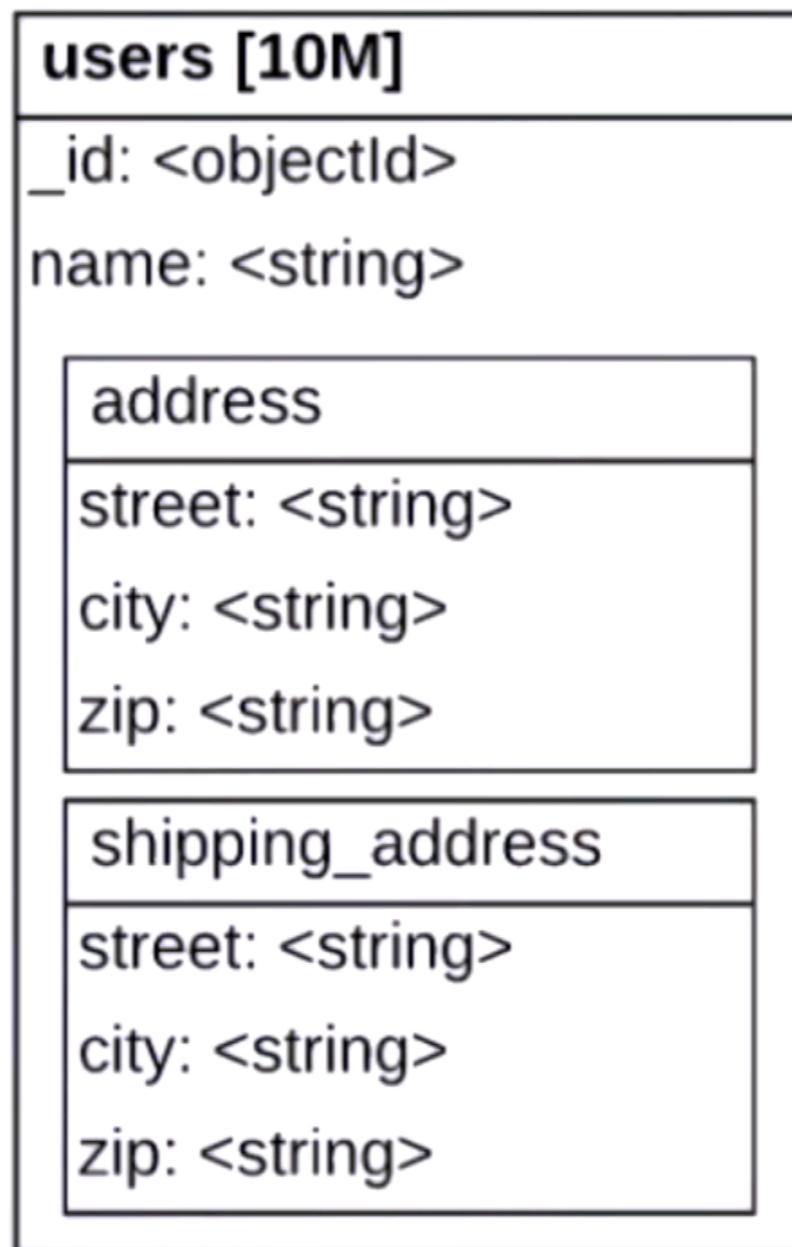
One-to-One

One-to-One



Embed vs Reference ?

One-to-One



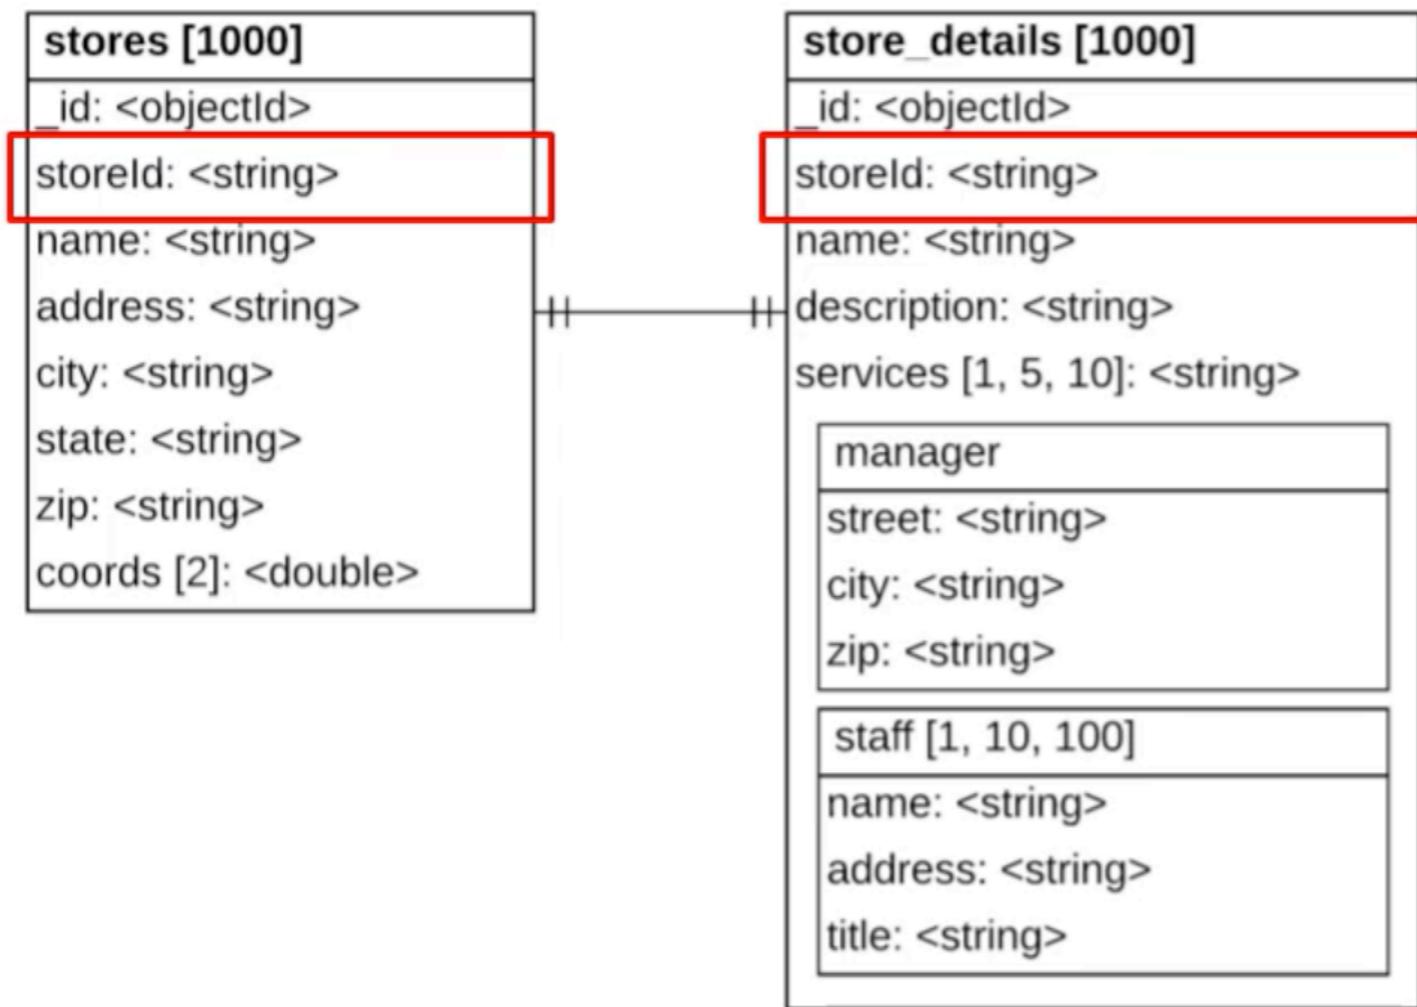
Embed vs Reference ?

One-to-One



Embed vs Reference ?

One-to-One



Embed vs Reference ?

One-to-One

Prefer embedding over reference for simplicity

Use subdocuments to organize the fields

Use reference for optimization purpose

One-to-Many

One-to-many !!

A mother and children ?
A person and address ?

Embed vs Reference ?

One-to-many

One-to-zillion

One-to-many (low)

One-to-Many



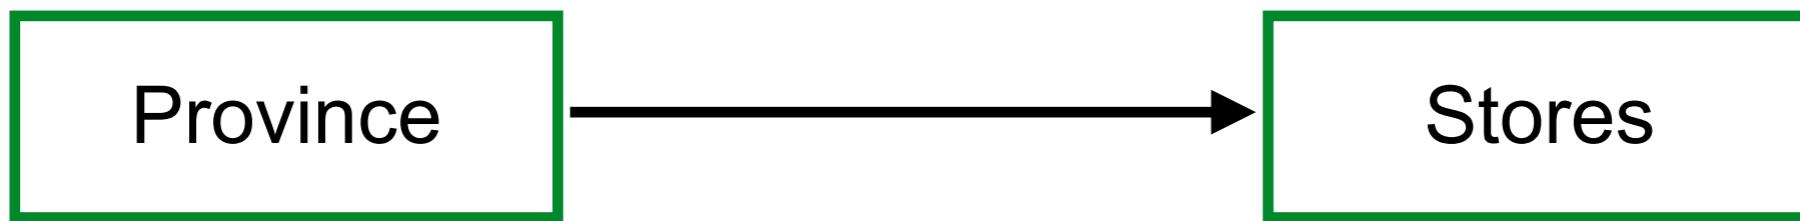
Embed vs Reference ?

One-to-Many



Embed vs Reference ?

One-to-Many



Embed vs Reference ?

One-to-Many



Stores = [1,2,3,4]

Embed vs Reference ?

One-to-Many



Embed vs Reference ?

One-to-Many



Stores = [1,2,3,4]

provinceld = 1

Embed vs Reference ?

One-to-Many

A Lot of choices

Choose the side between **one** and **many**

Duplication when use embedding

Embed on the side of the most query

Prefer referencing when associated documents
are not always needed with often query
documents

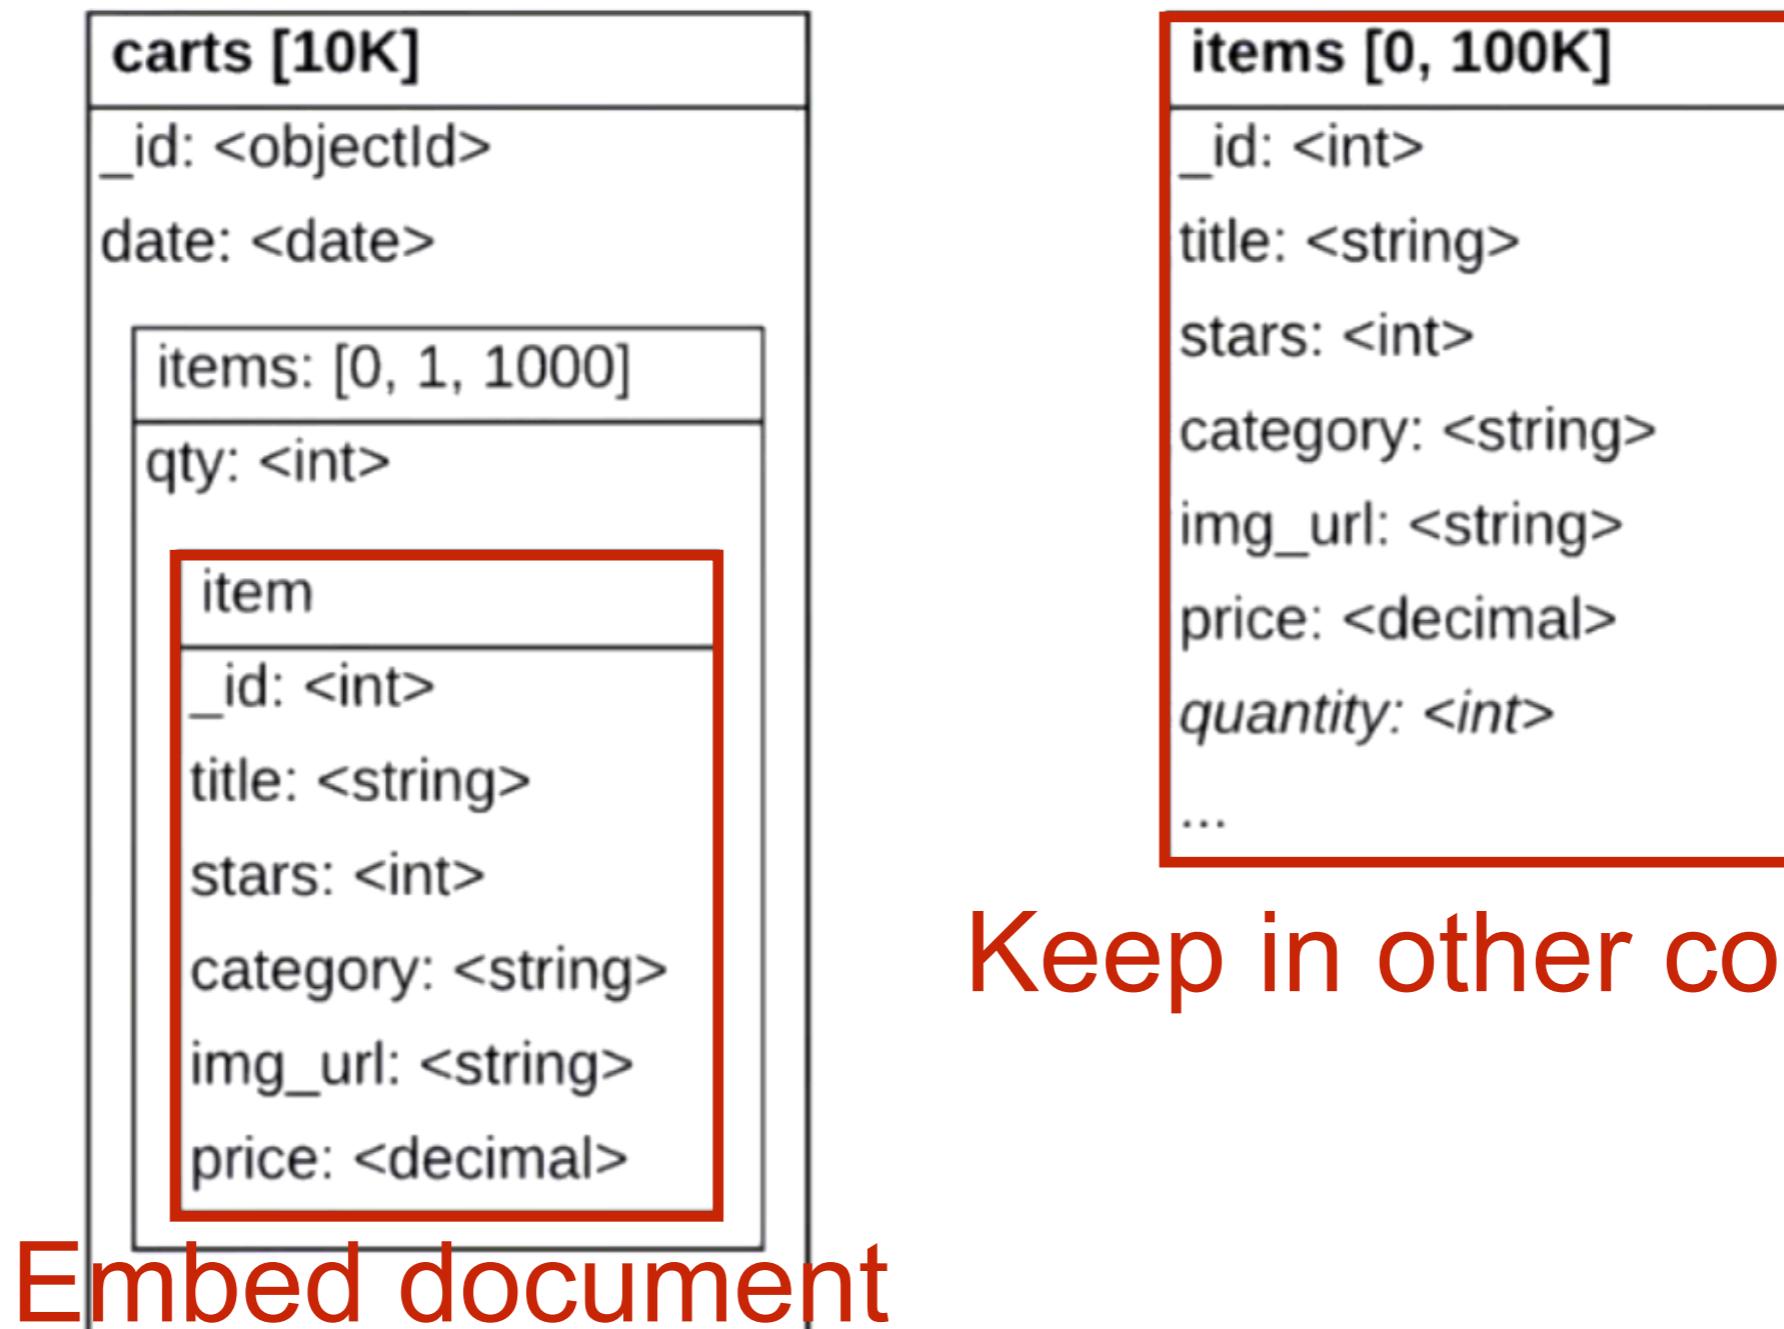
Many-to-Many

Many-to-Many



Embed vs Reference ?

Many-to-Many



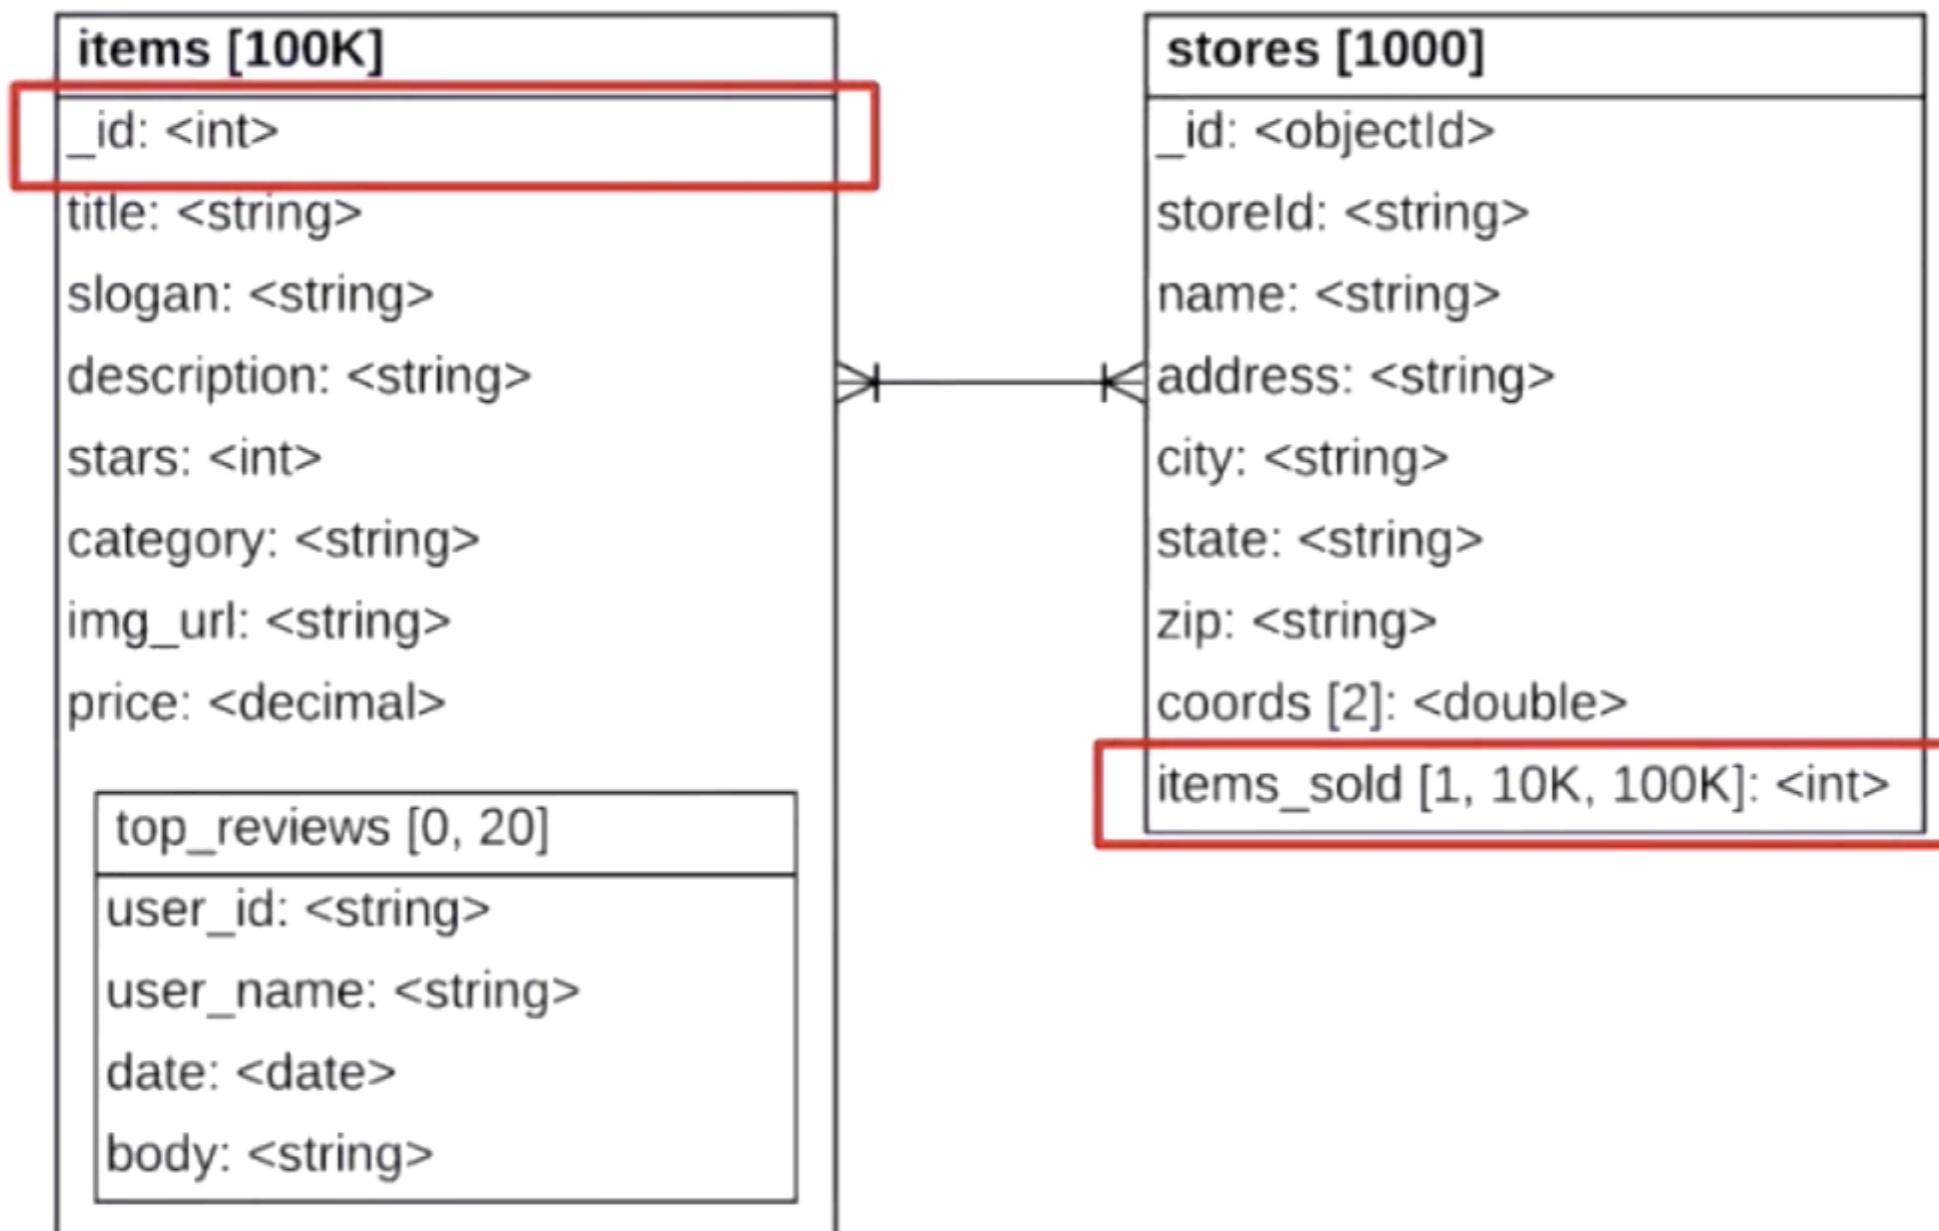
Many-to-Many



Find all items sold in each store ?

Embed vs Reference ?

Many-to-Many



Many-to-Many

Prefer reference over embedding to avoid
duplication

Prefer embedding on the most query side

Workshop

MongoDB Deployment

Goals

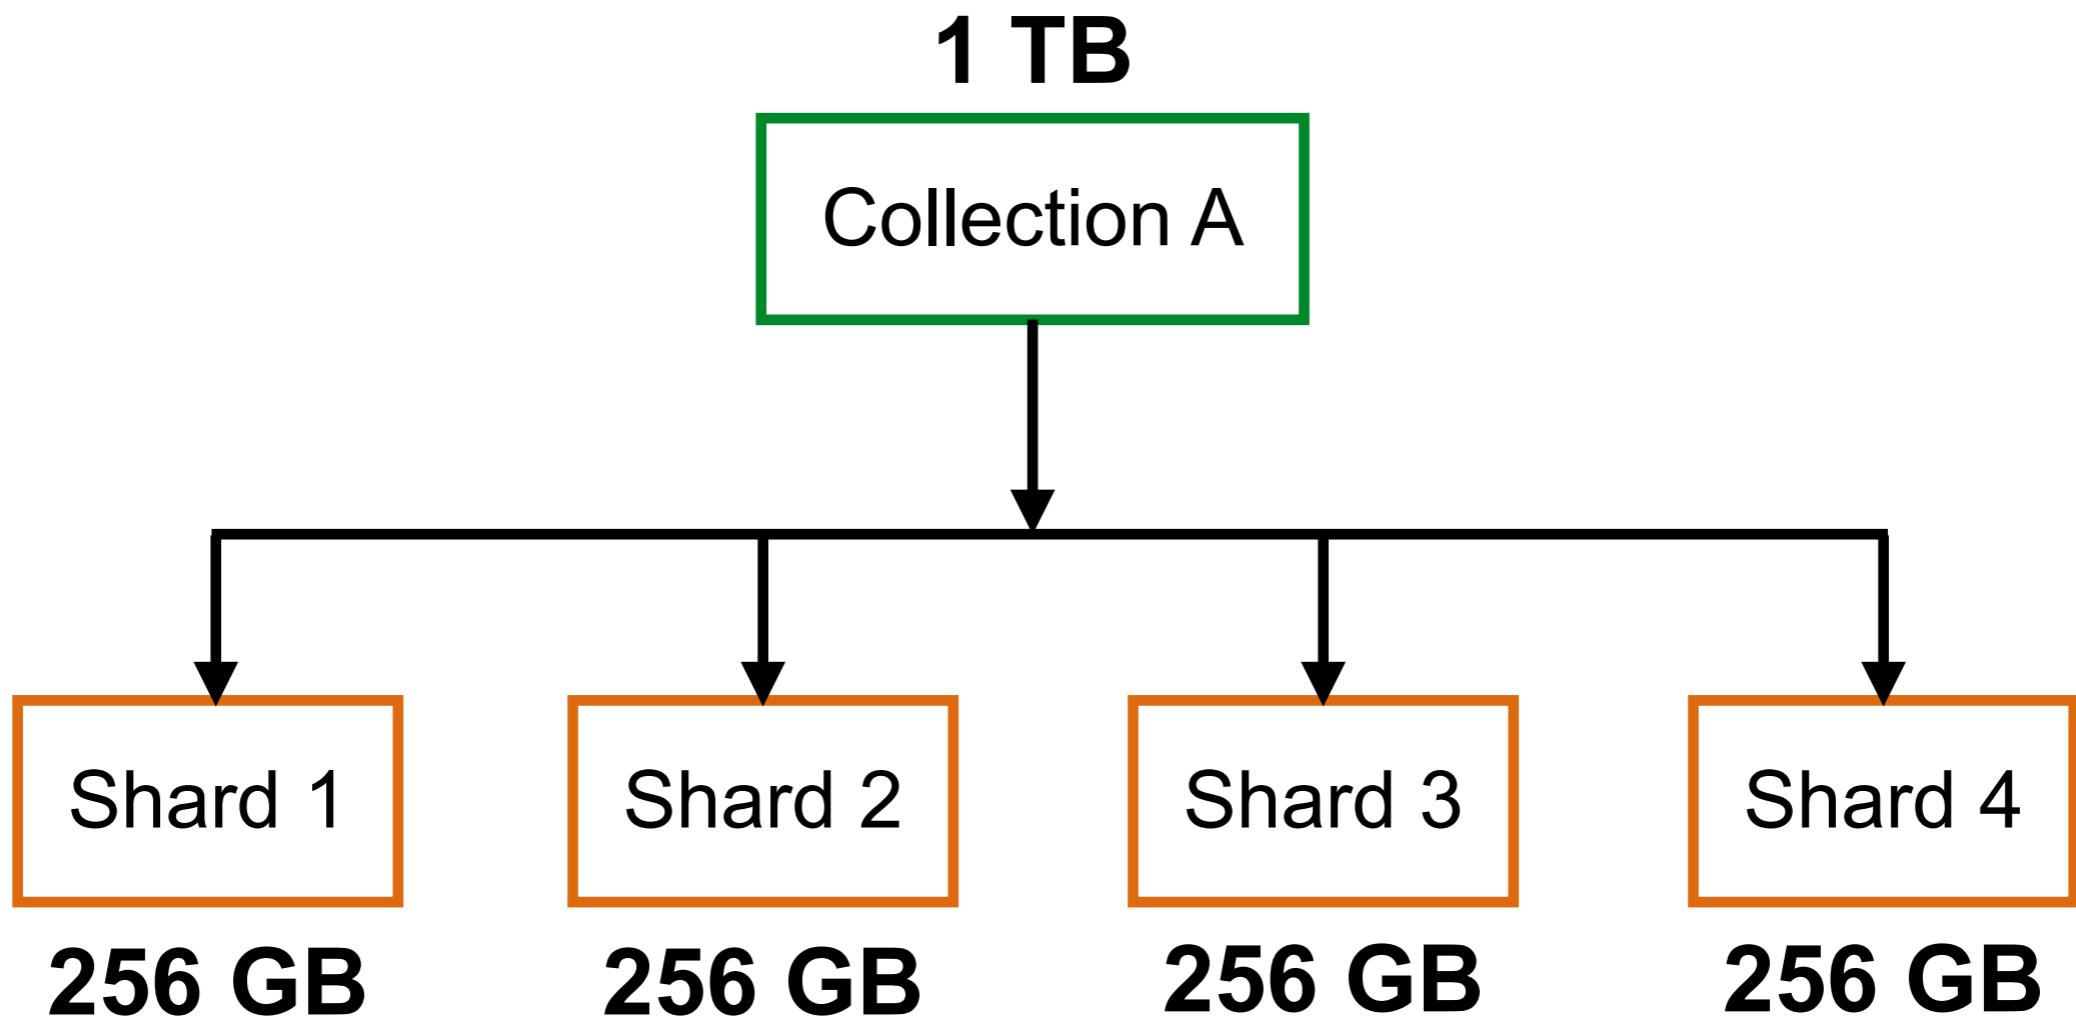
Architecture of a sharded cluster
Query handling in a sharded cluster
Data distribution method

Sharding

Distributing data across multiple machines
Horizontal scaling

In MongoDB with large data sets,
needed **high throughput** operations

Sharding



When to use Shards ?

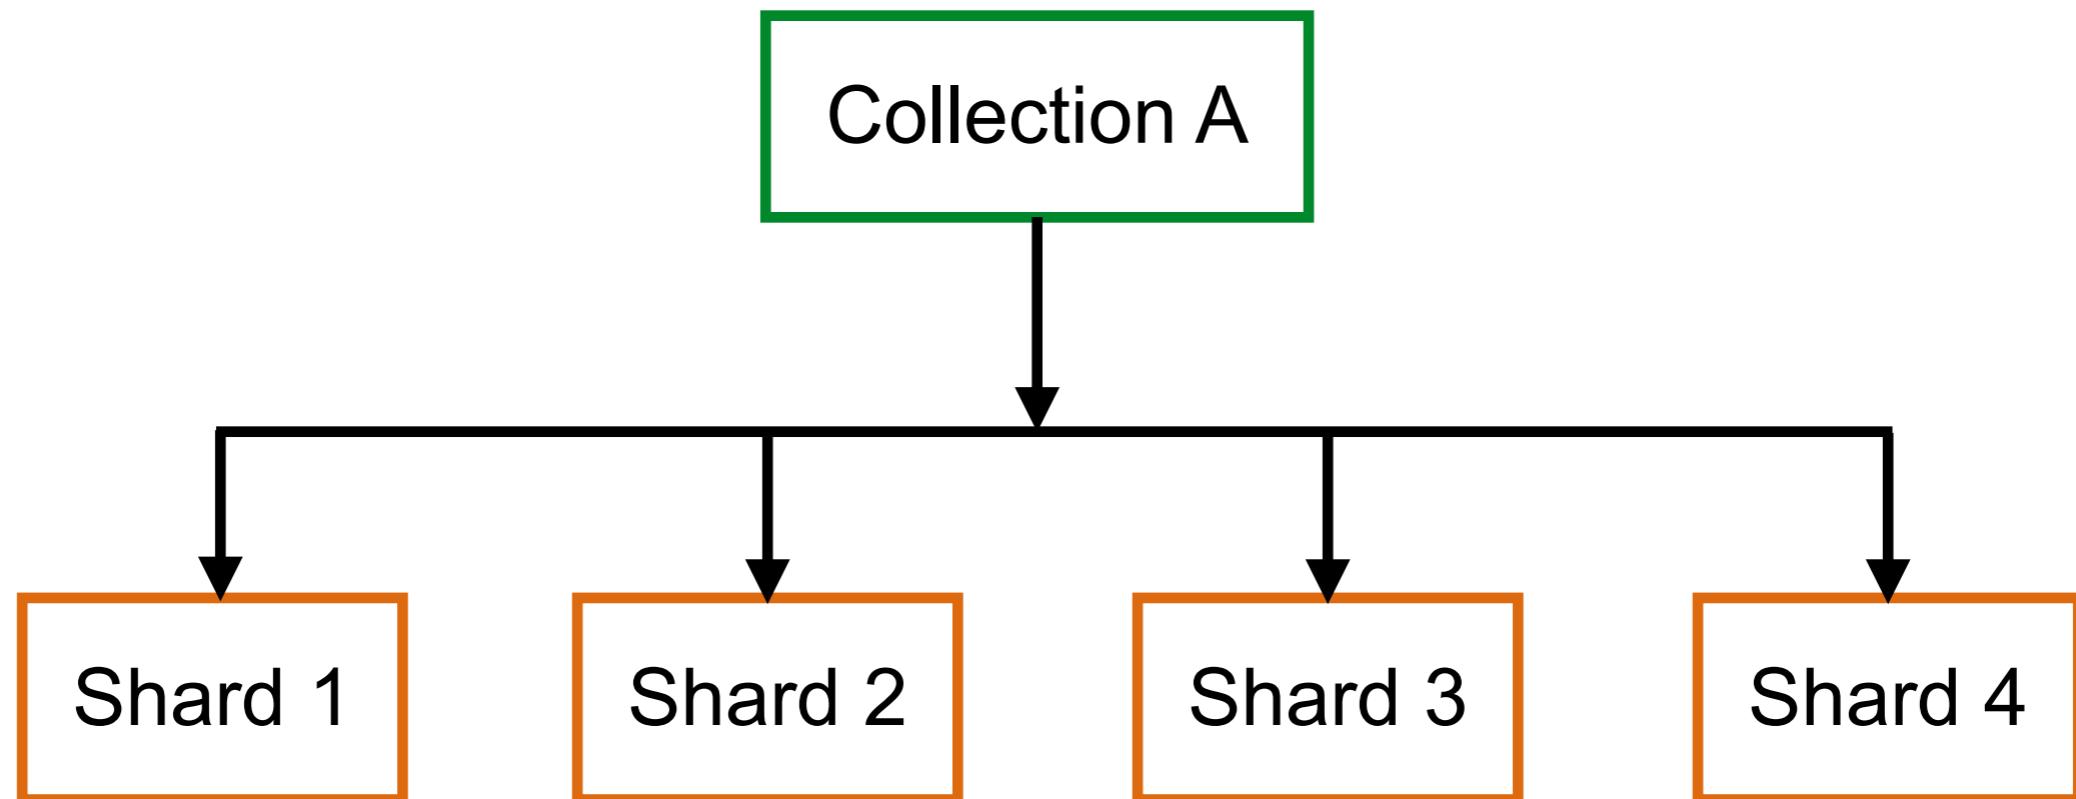
Write workload on a single server exceed that server's capacity

Working set no longer fits into RAM

Single server can no longer handle the size of data set

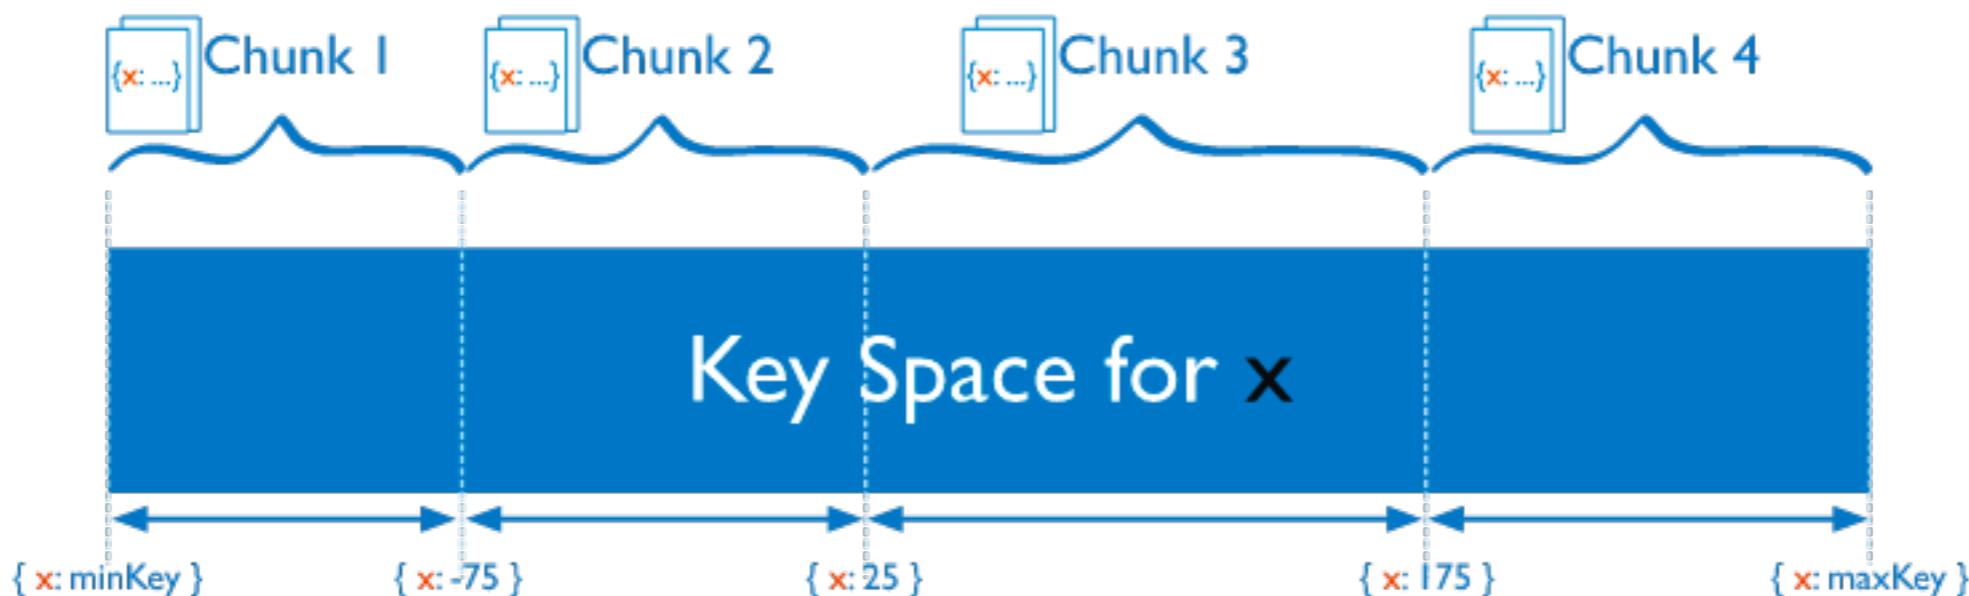
Shard key

Use to determine how documents in a collection are distributed across the shards



Shard key

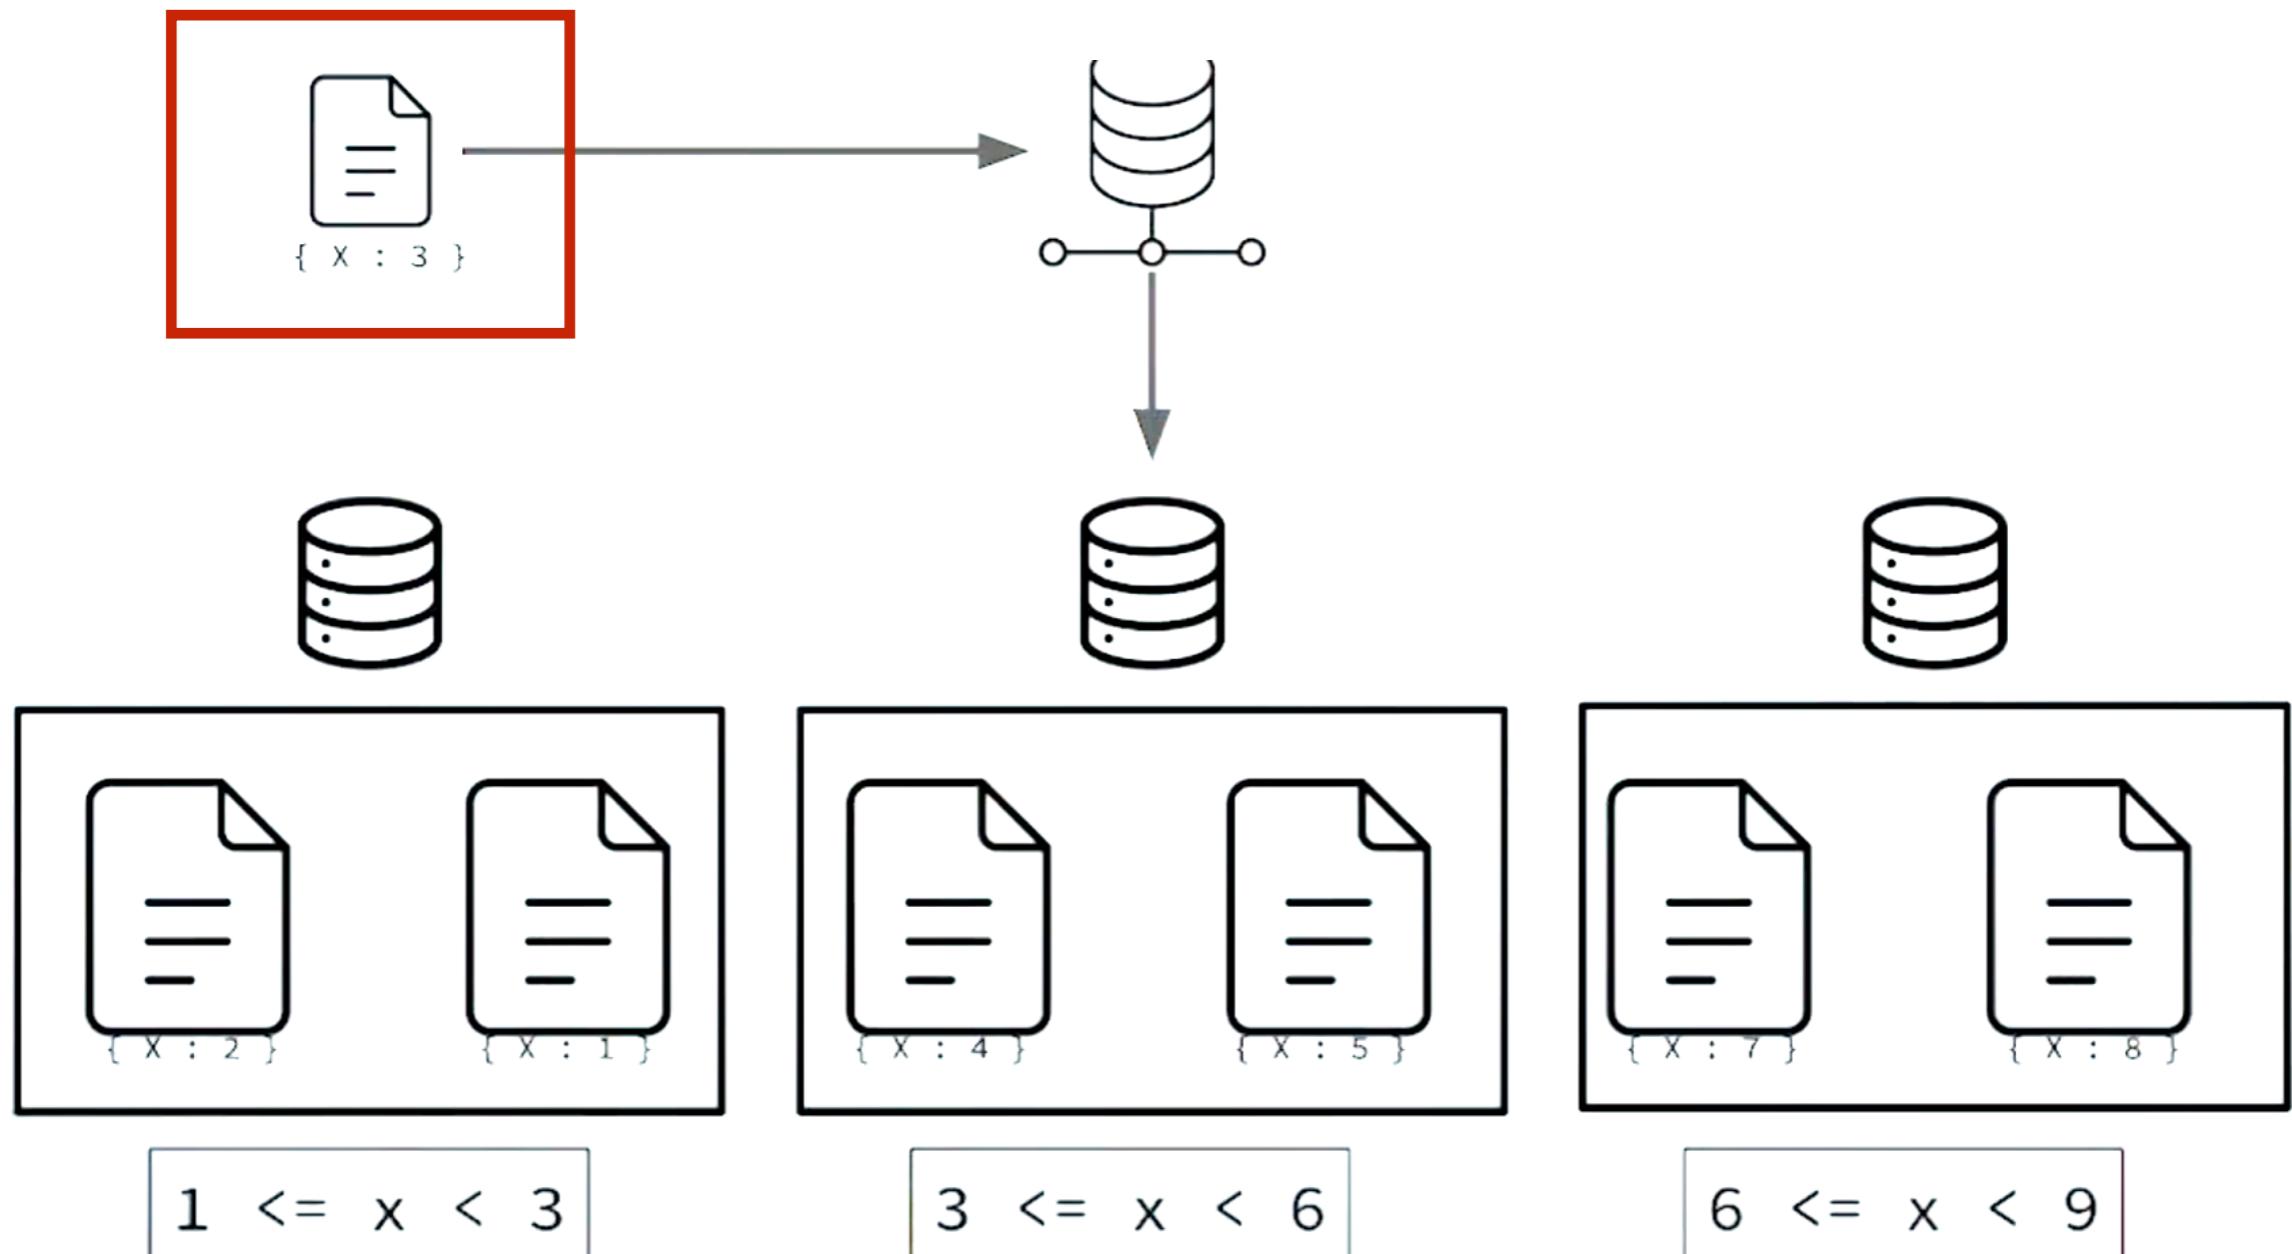
Use to determine how documents in a collection are distributed across the shards



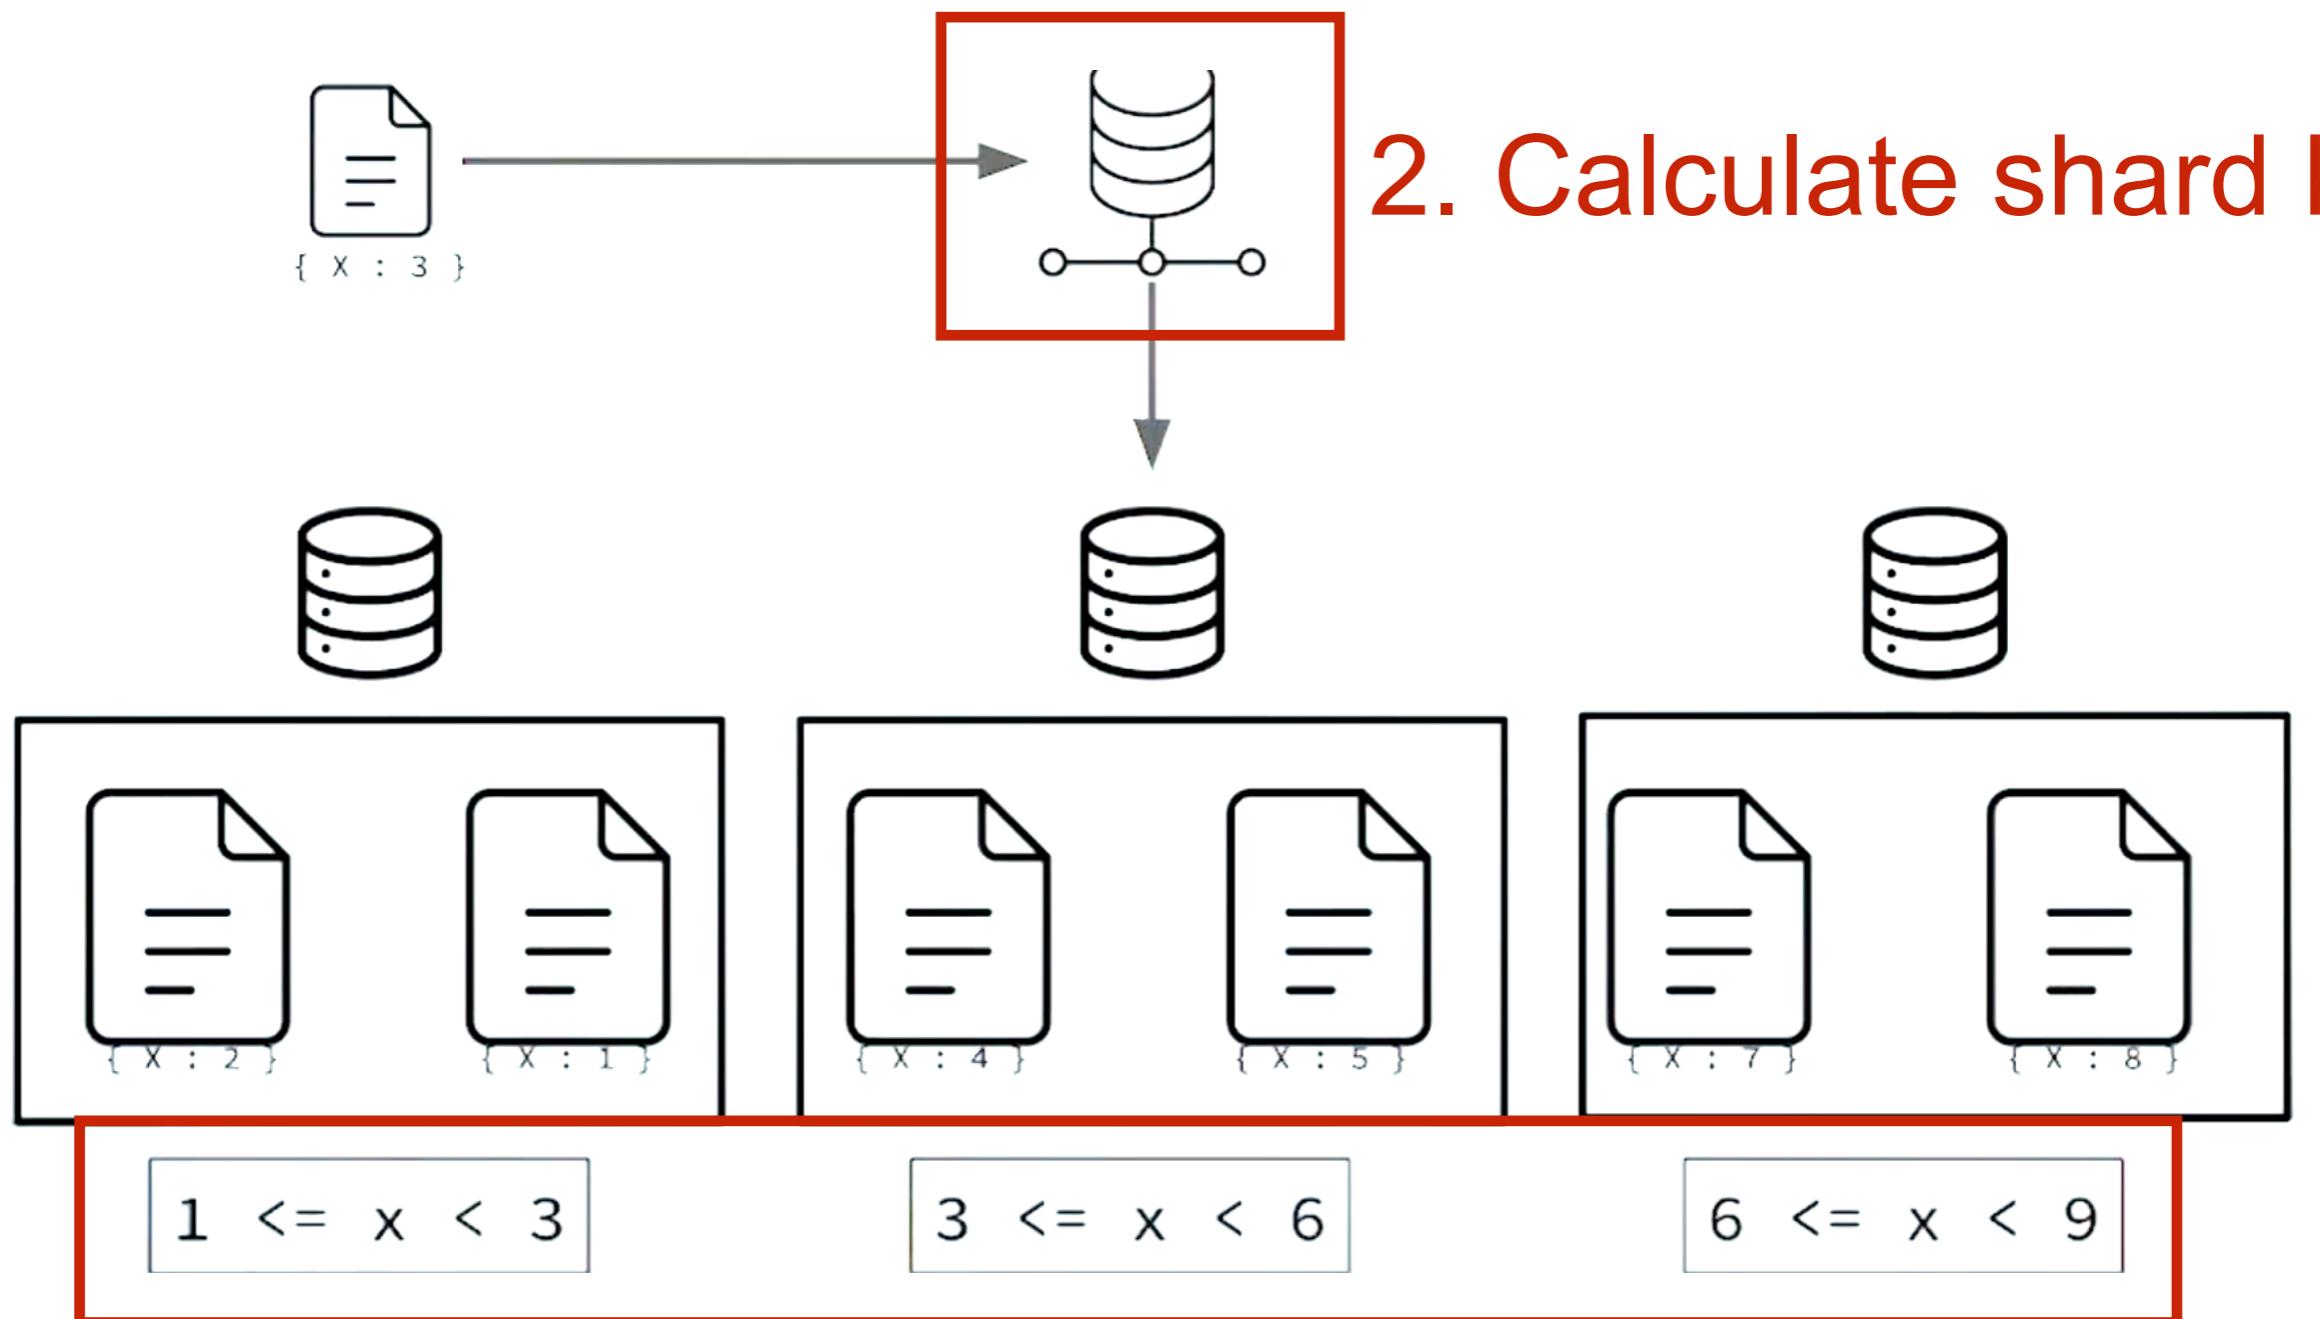
<https://docs.mongodb.com/manual/core/sharding-shard-key/>

Shard key

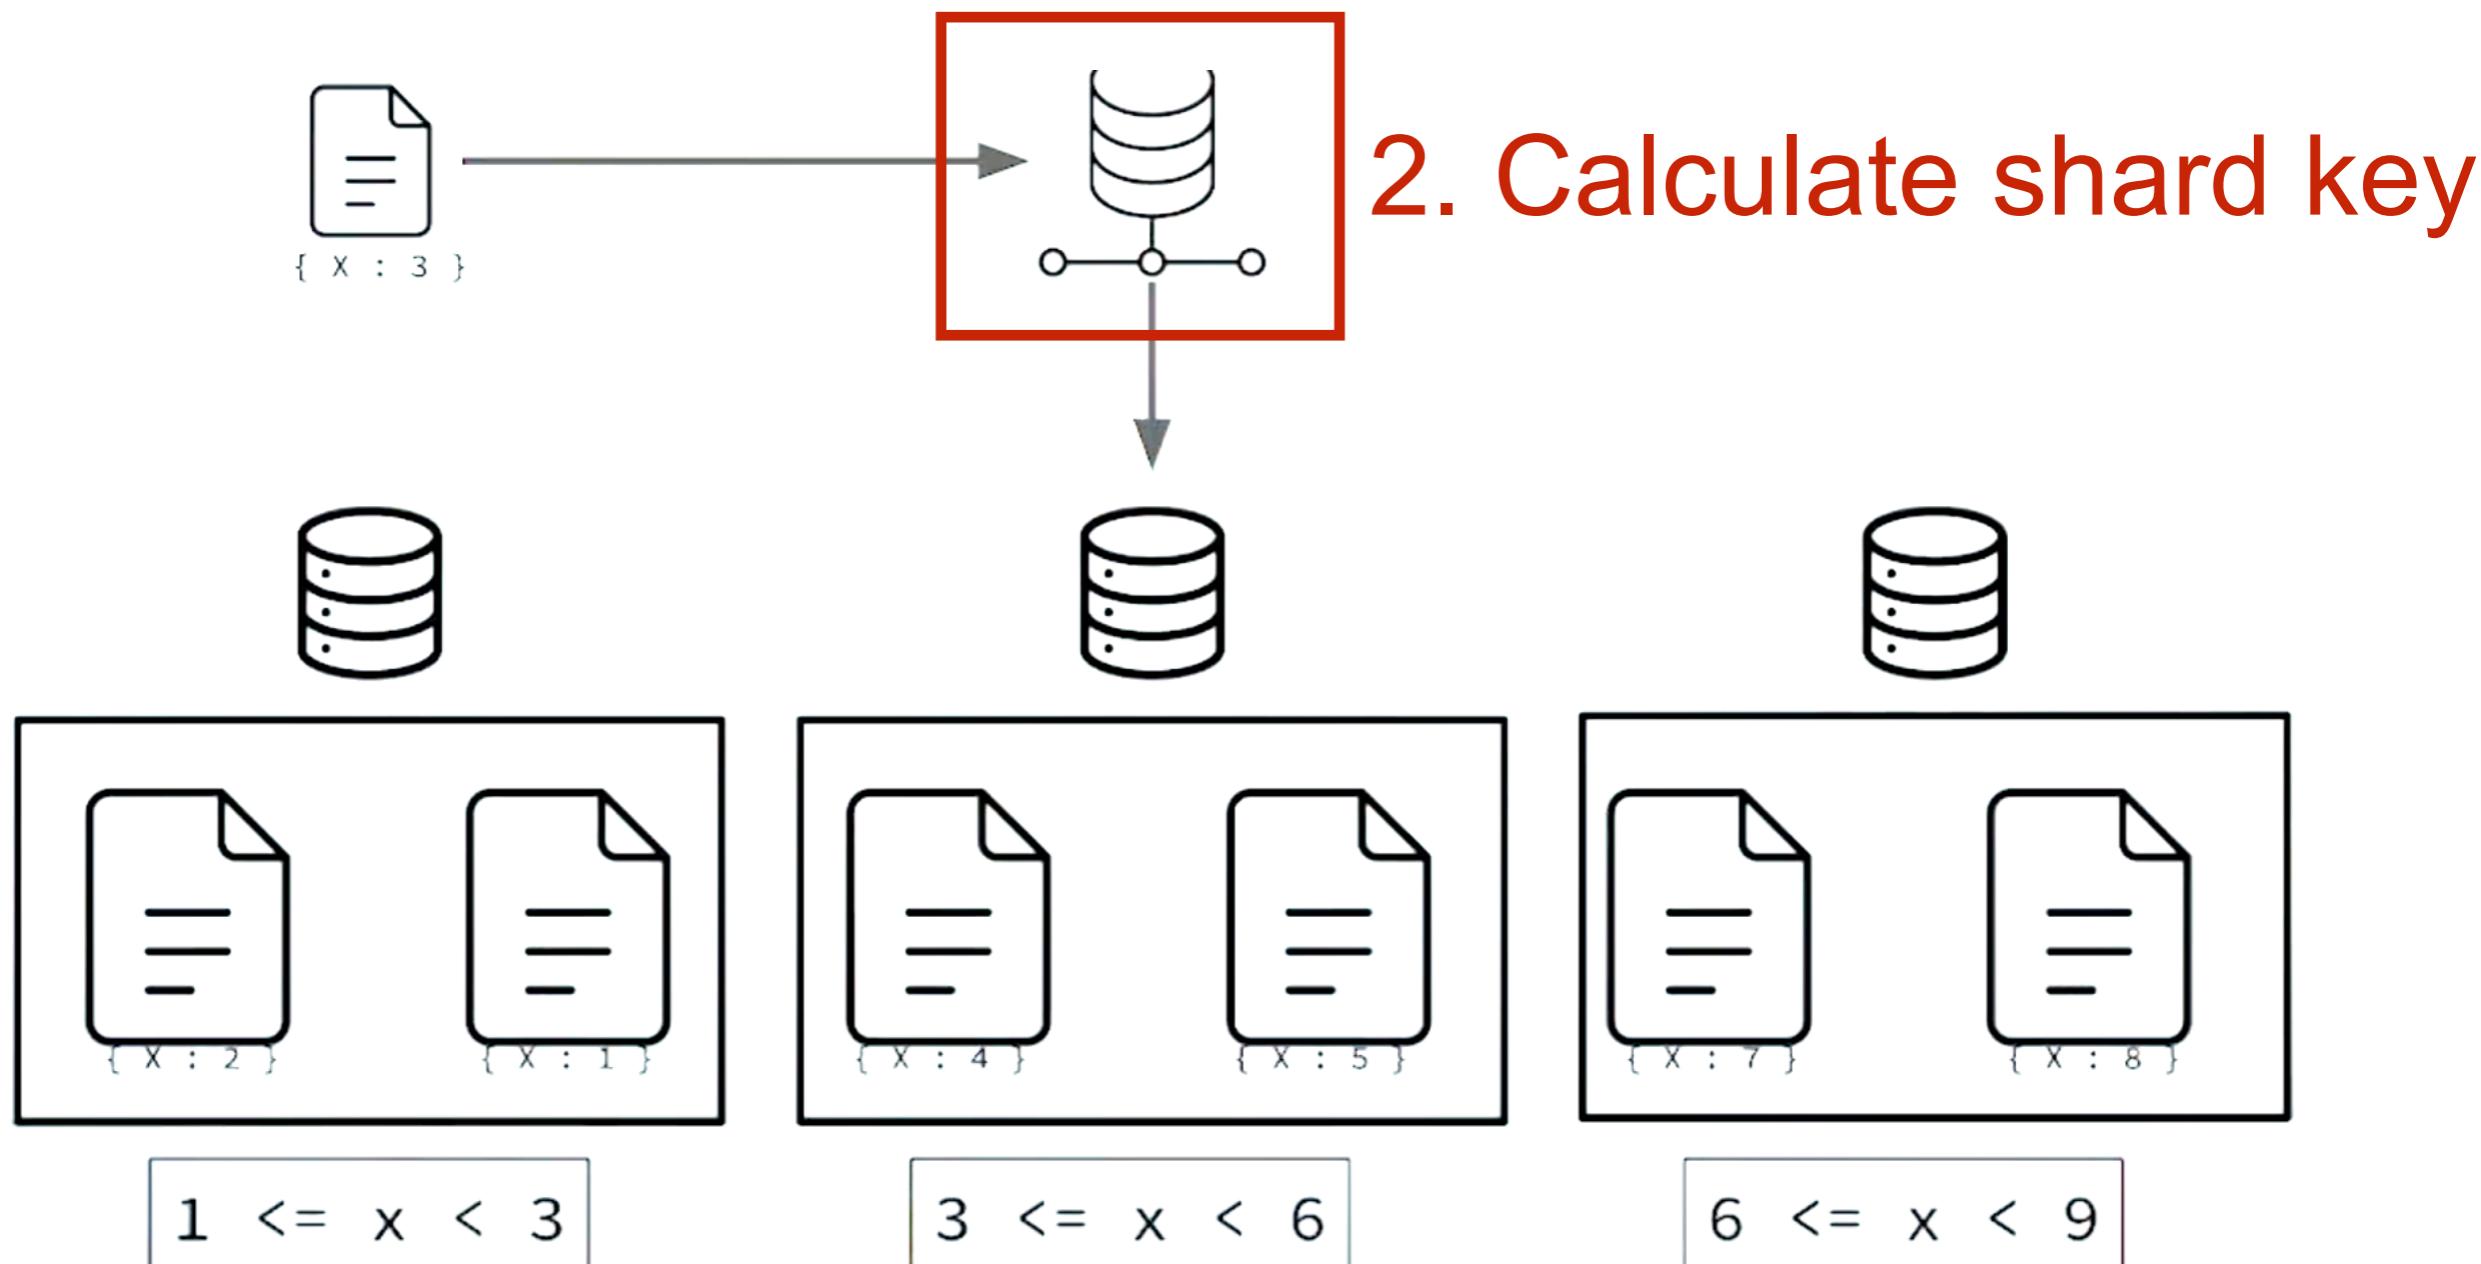
1. Create a document



Shard key



Shard key



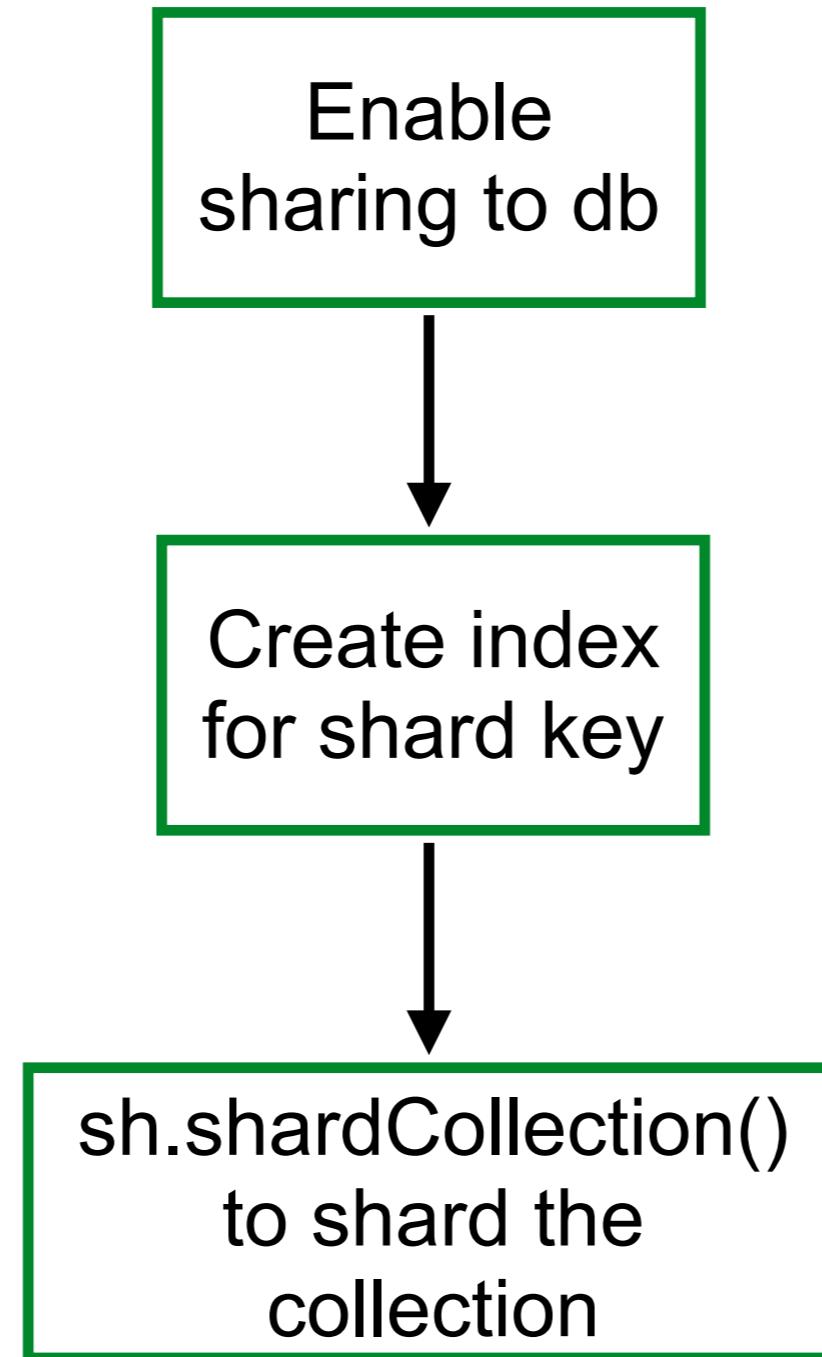
Shard key

Must be indexed

Immutable (can't be change)

Permanent

How to Shard ?



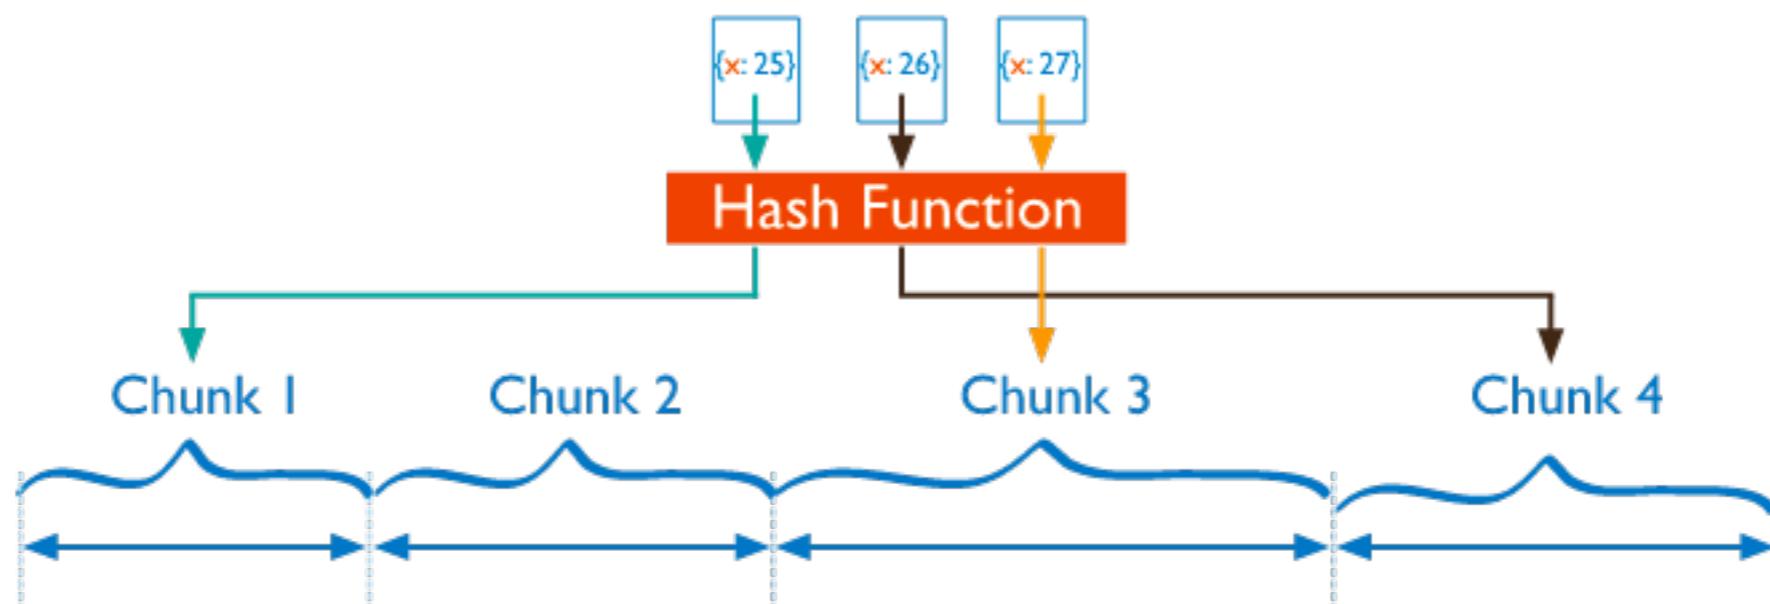
Choose shard key ?

The goal is a shard key whose values provides
good write distribution

- Cardinality**
- Frequency**
- Monotonic change**

Hashed Shard key

Single field hashed index
Compound hashed index



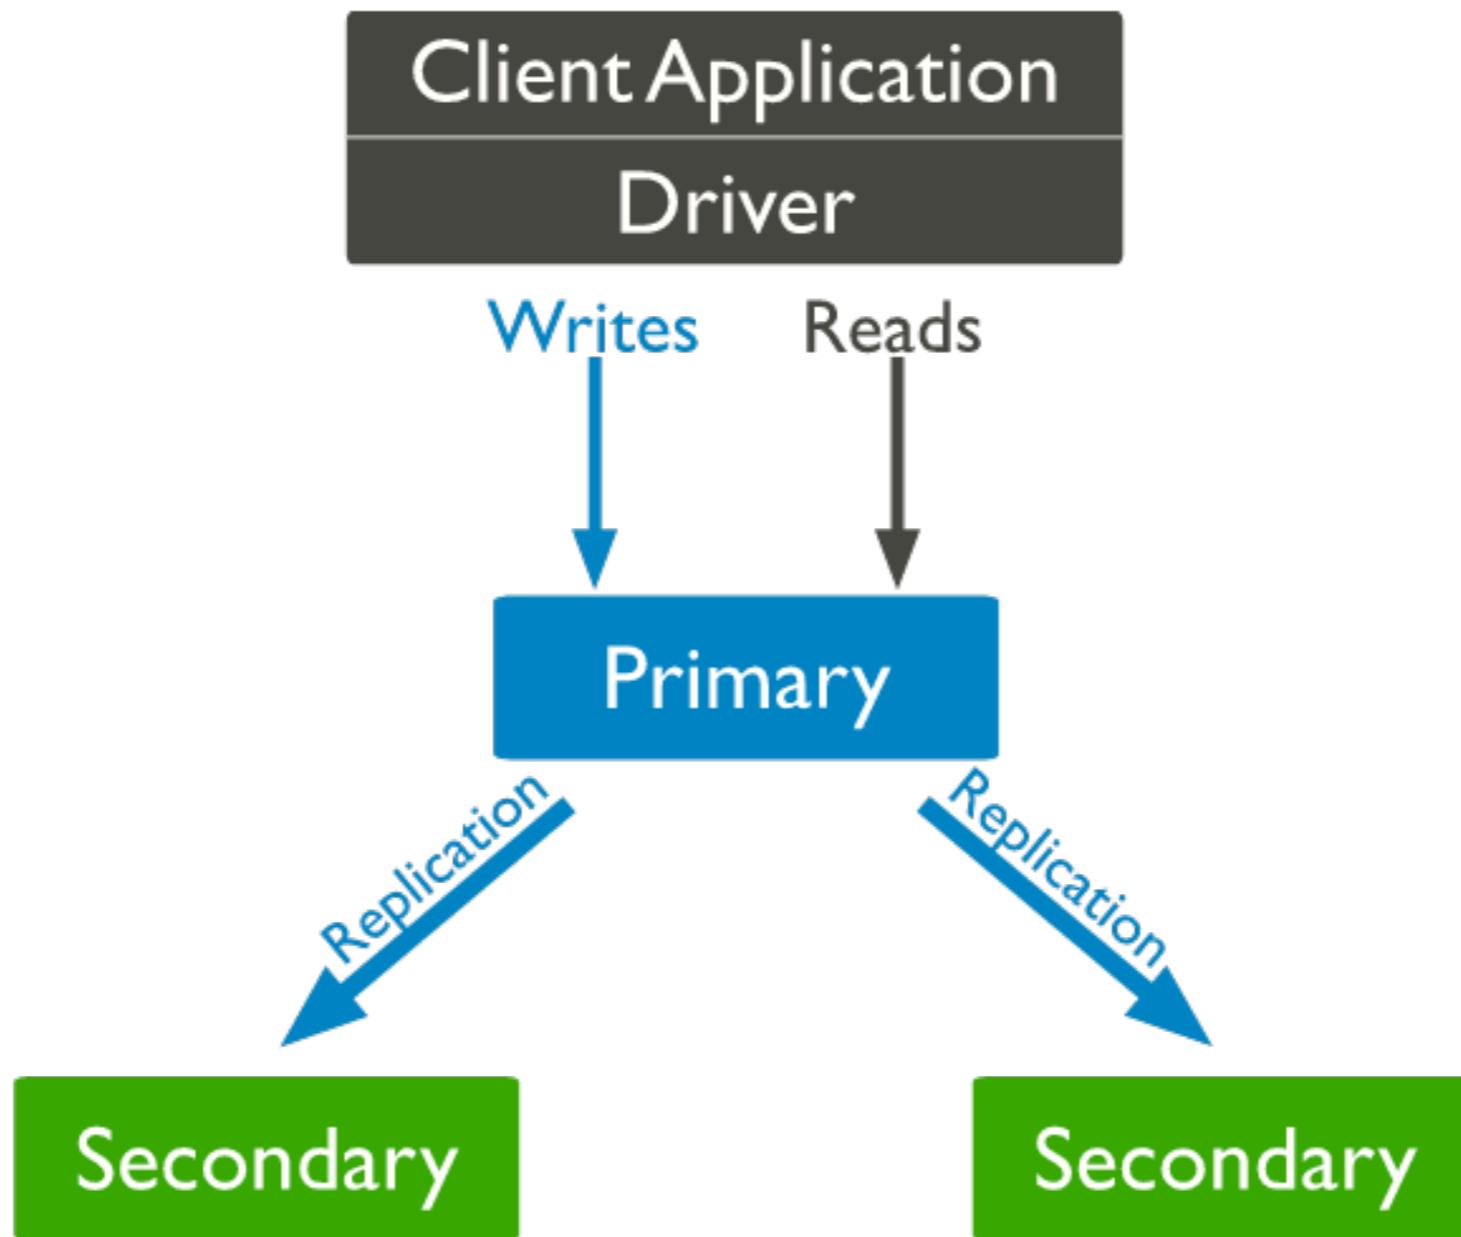
<https://docs.mongodb.com/manual/core/hashed-sharding/>

Workshop

MongoDB Deployment

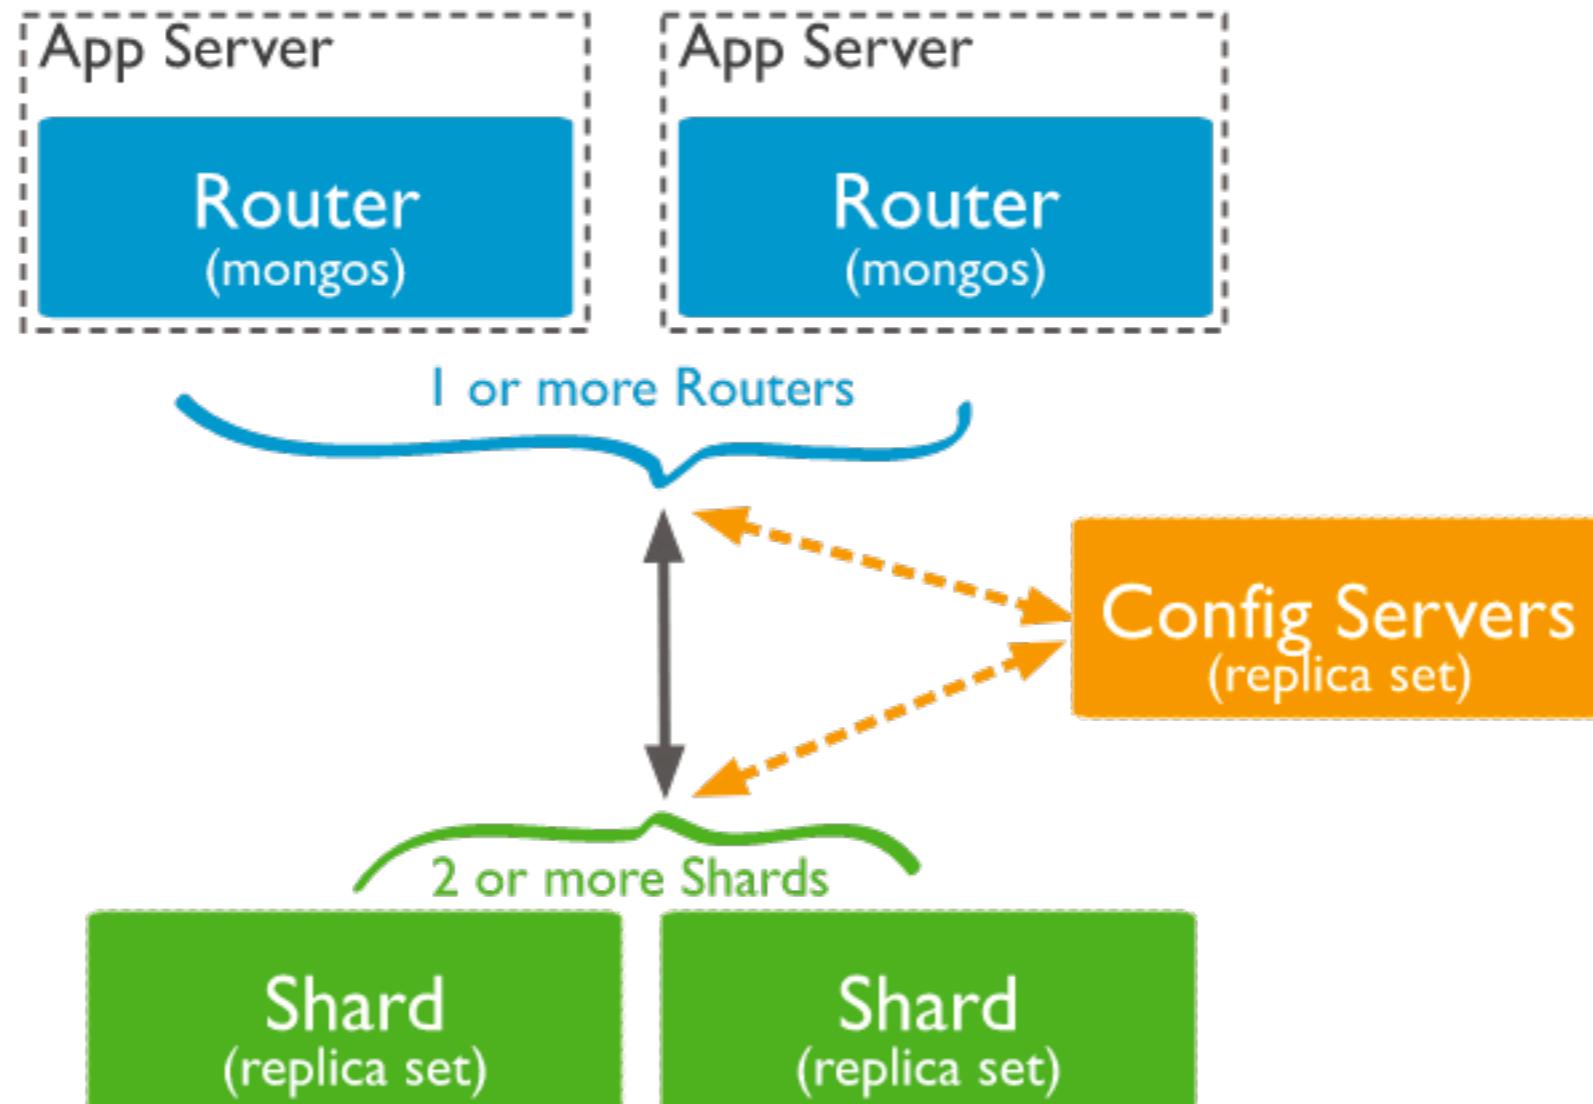
Standalone or Cluster
Replication
Shard cluster

Replication Deployment



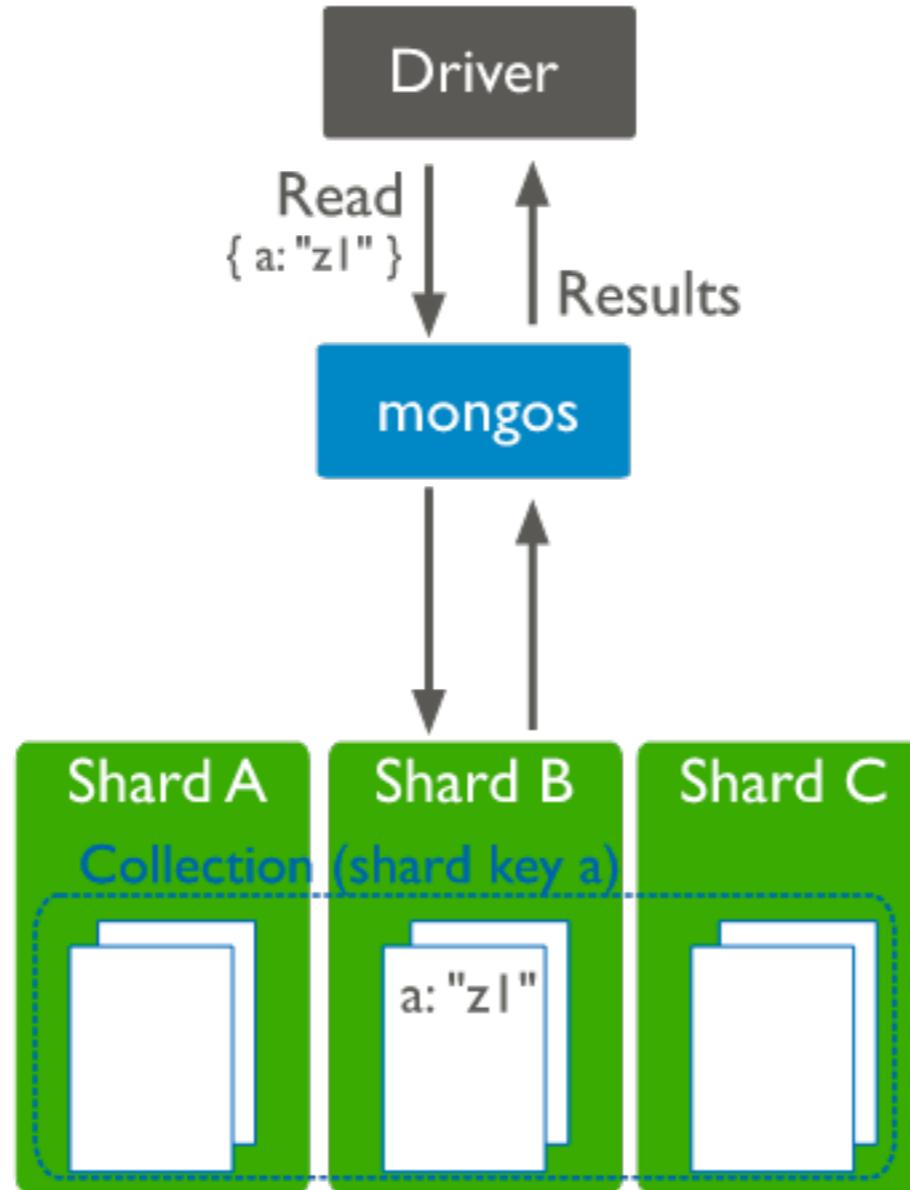
<https://docs.mongodb.com/manual/replication/>

Shard Cluster Deployment



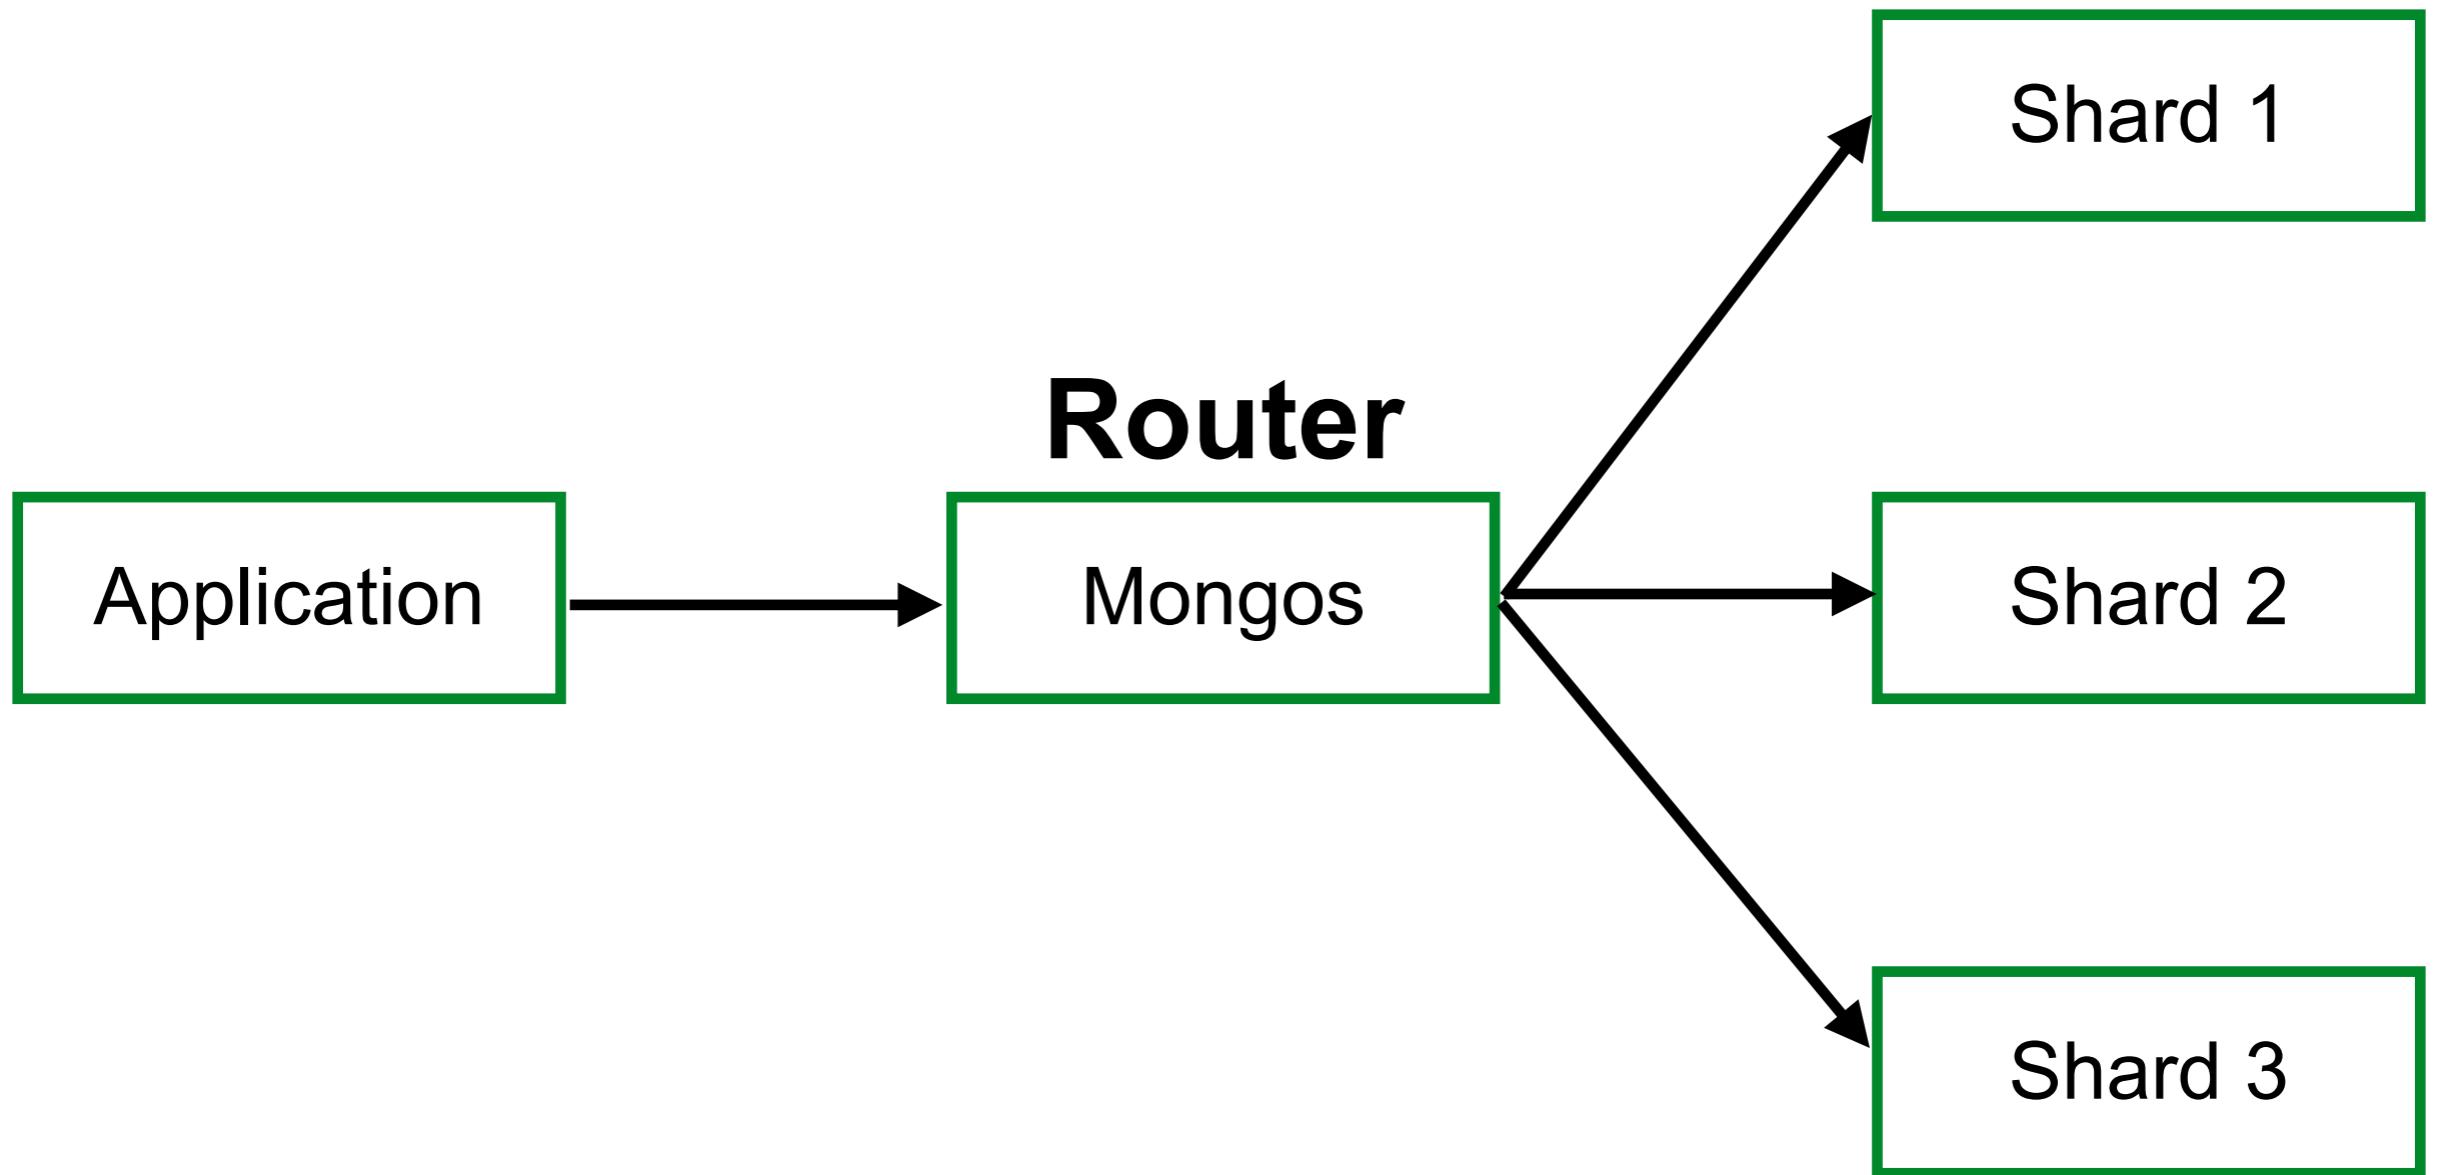
<https://docs.mongodb.com/manual/core/sharded-cluster-components/>

Shard Cluster Deployment

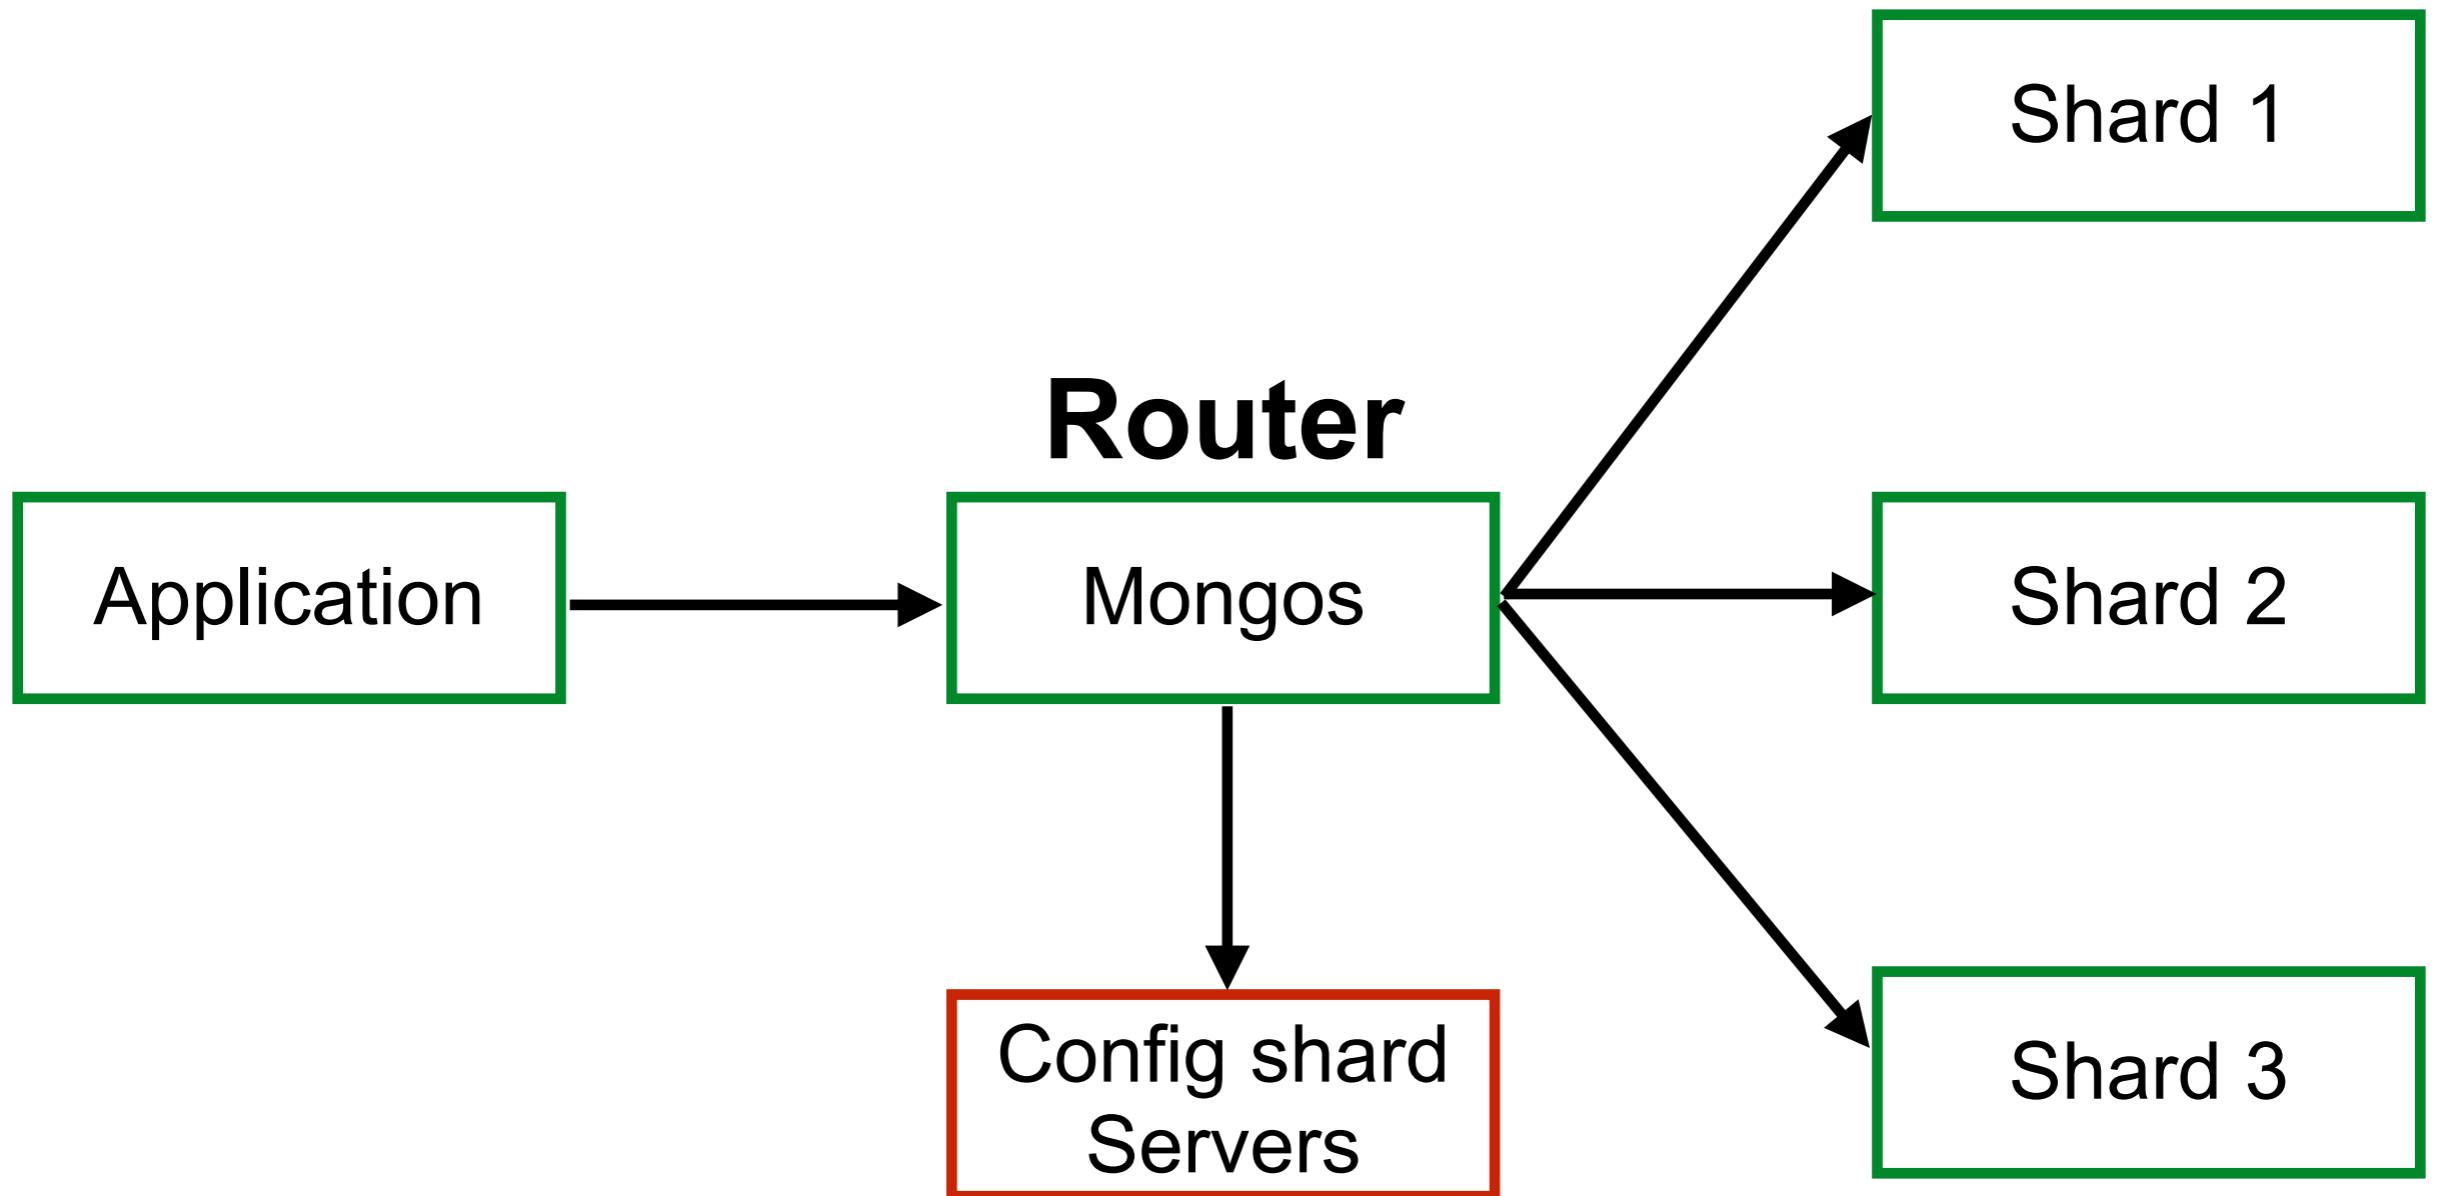


<https://docs.mongodb.com/manual/core/sharded-cluster-components/>

Shard Cluster Deployment



Shard Cluster Deployment



Query from shard cluster

sort()

limit()

skip()

<https://docs.mongodb.com/manual/core/distributed-queries/>

Query from shard cluster

sort()

The mongos pushes the sort to each shard and merge-sorts the result

limit()

The mongos passes the limit to each targeted shard, then re-applies the limit to the merge set of results

skip()

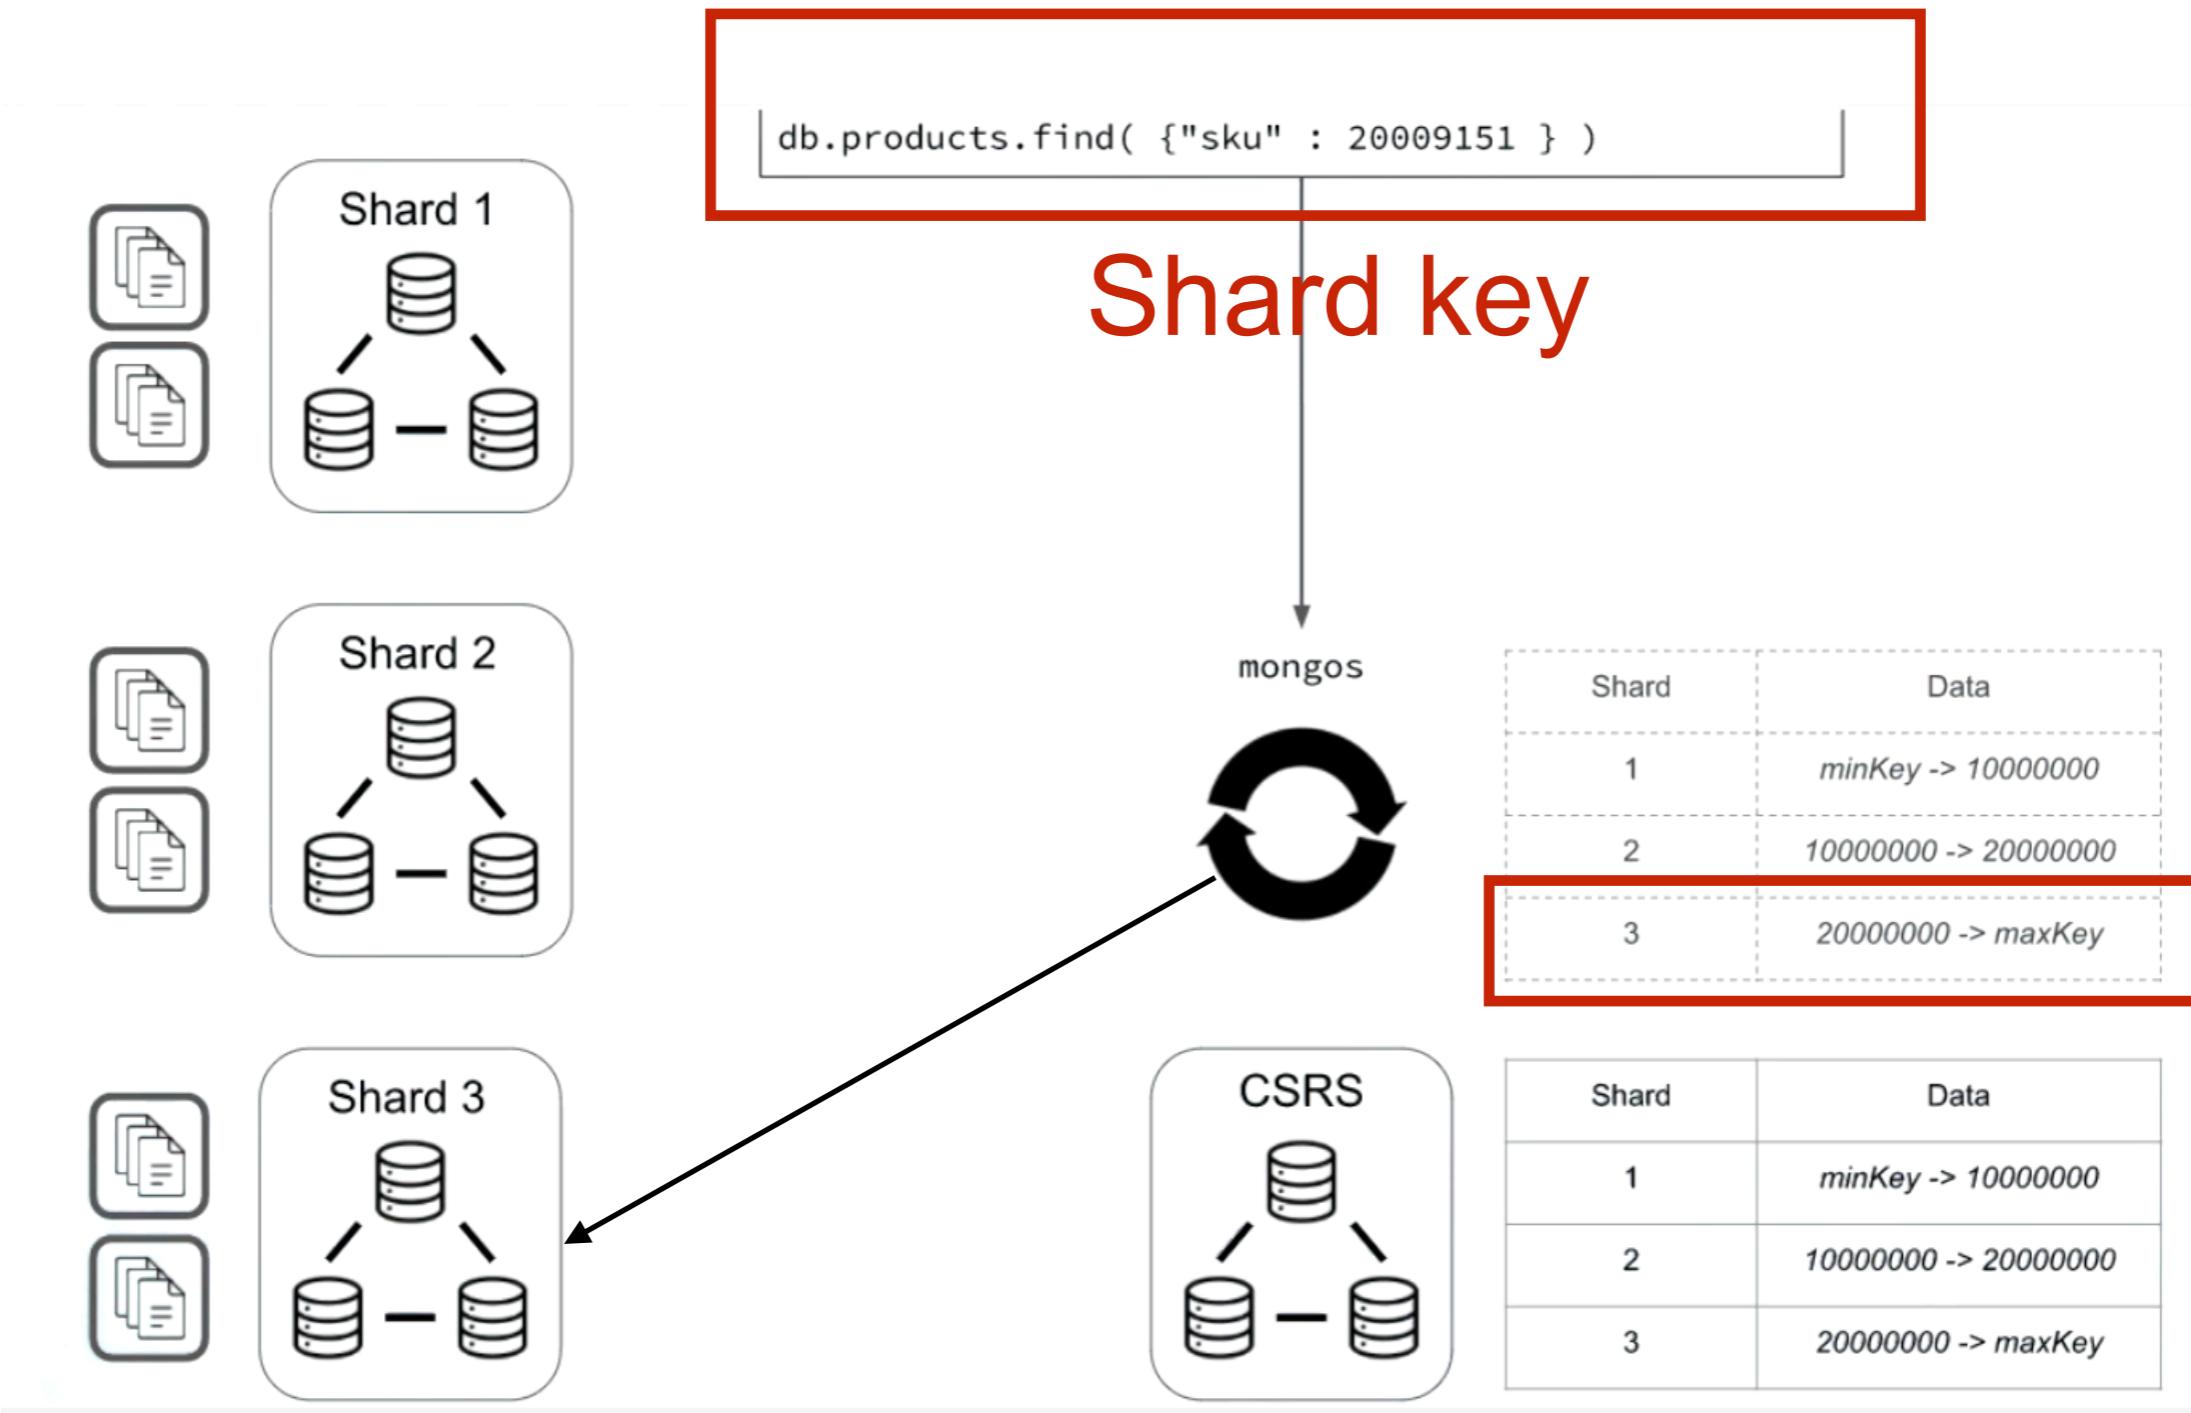
The mongos performs the skip against the merged set of results

Targeted vs Scatter gather query

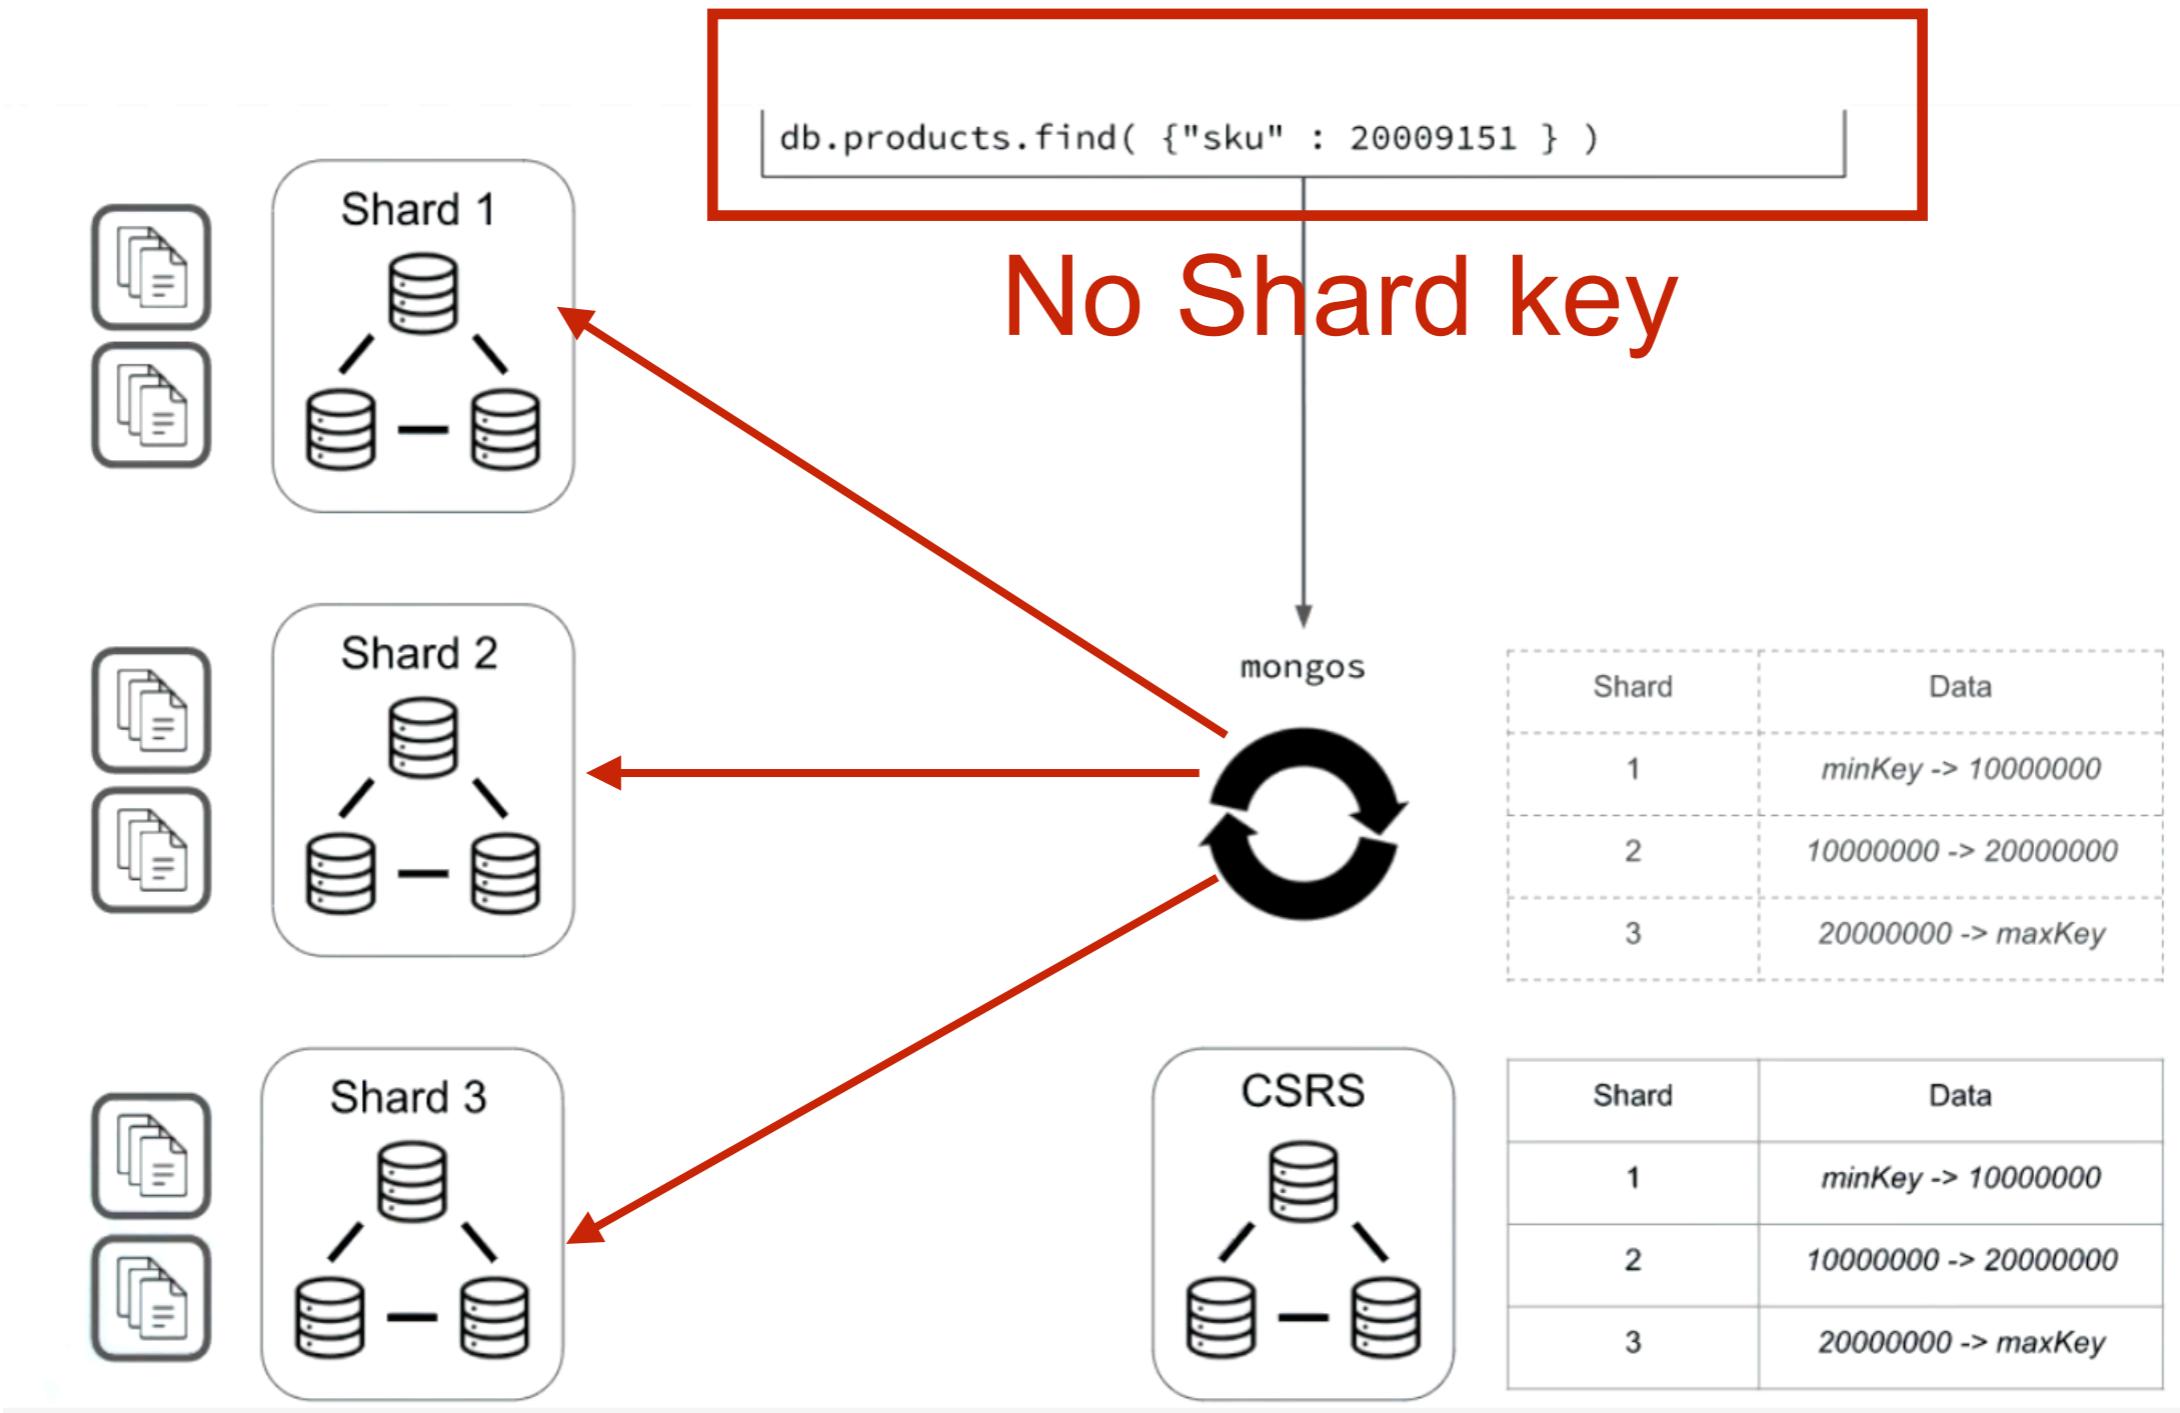
Targeted queries requires the shard key

Without shard key, the mongos must perform a
scatter-gather query

Targeted query



Scatter gather query



Workshop

Security

Security in MongoDB

User authentication

Role-Based Access Control (RBAC)

Internal communication authentication

Authentication vs Authorization



Working with MongoDB

Create user document

```
db.createUser(  
  {  
    user: "accountAnalytics",  
    pwd: "LionelMessi1987",  
    roles: [  
      { role: "read", db: "accounts" },  
      { role: "readWrite", db: "analytics" }  
    ]  
  }  
)
```

Working with MongoDB

Create roles of user

```
db.createUser(  
  {  
    user: "accountAnalytics",  
    pwd: "LionelMessi1987",  
    roles: [  
      { role: "read", db: "accounts" },  
      { role: "readWrite", db: "analytics" }  
    ]  
  }  
)
```

Quality Assurance feedback

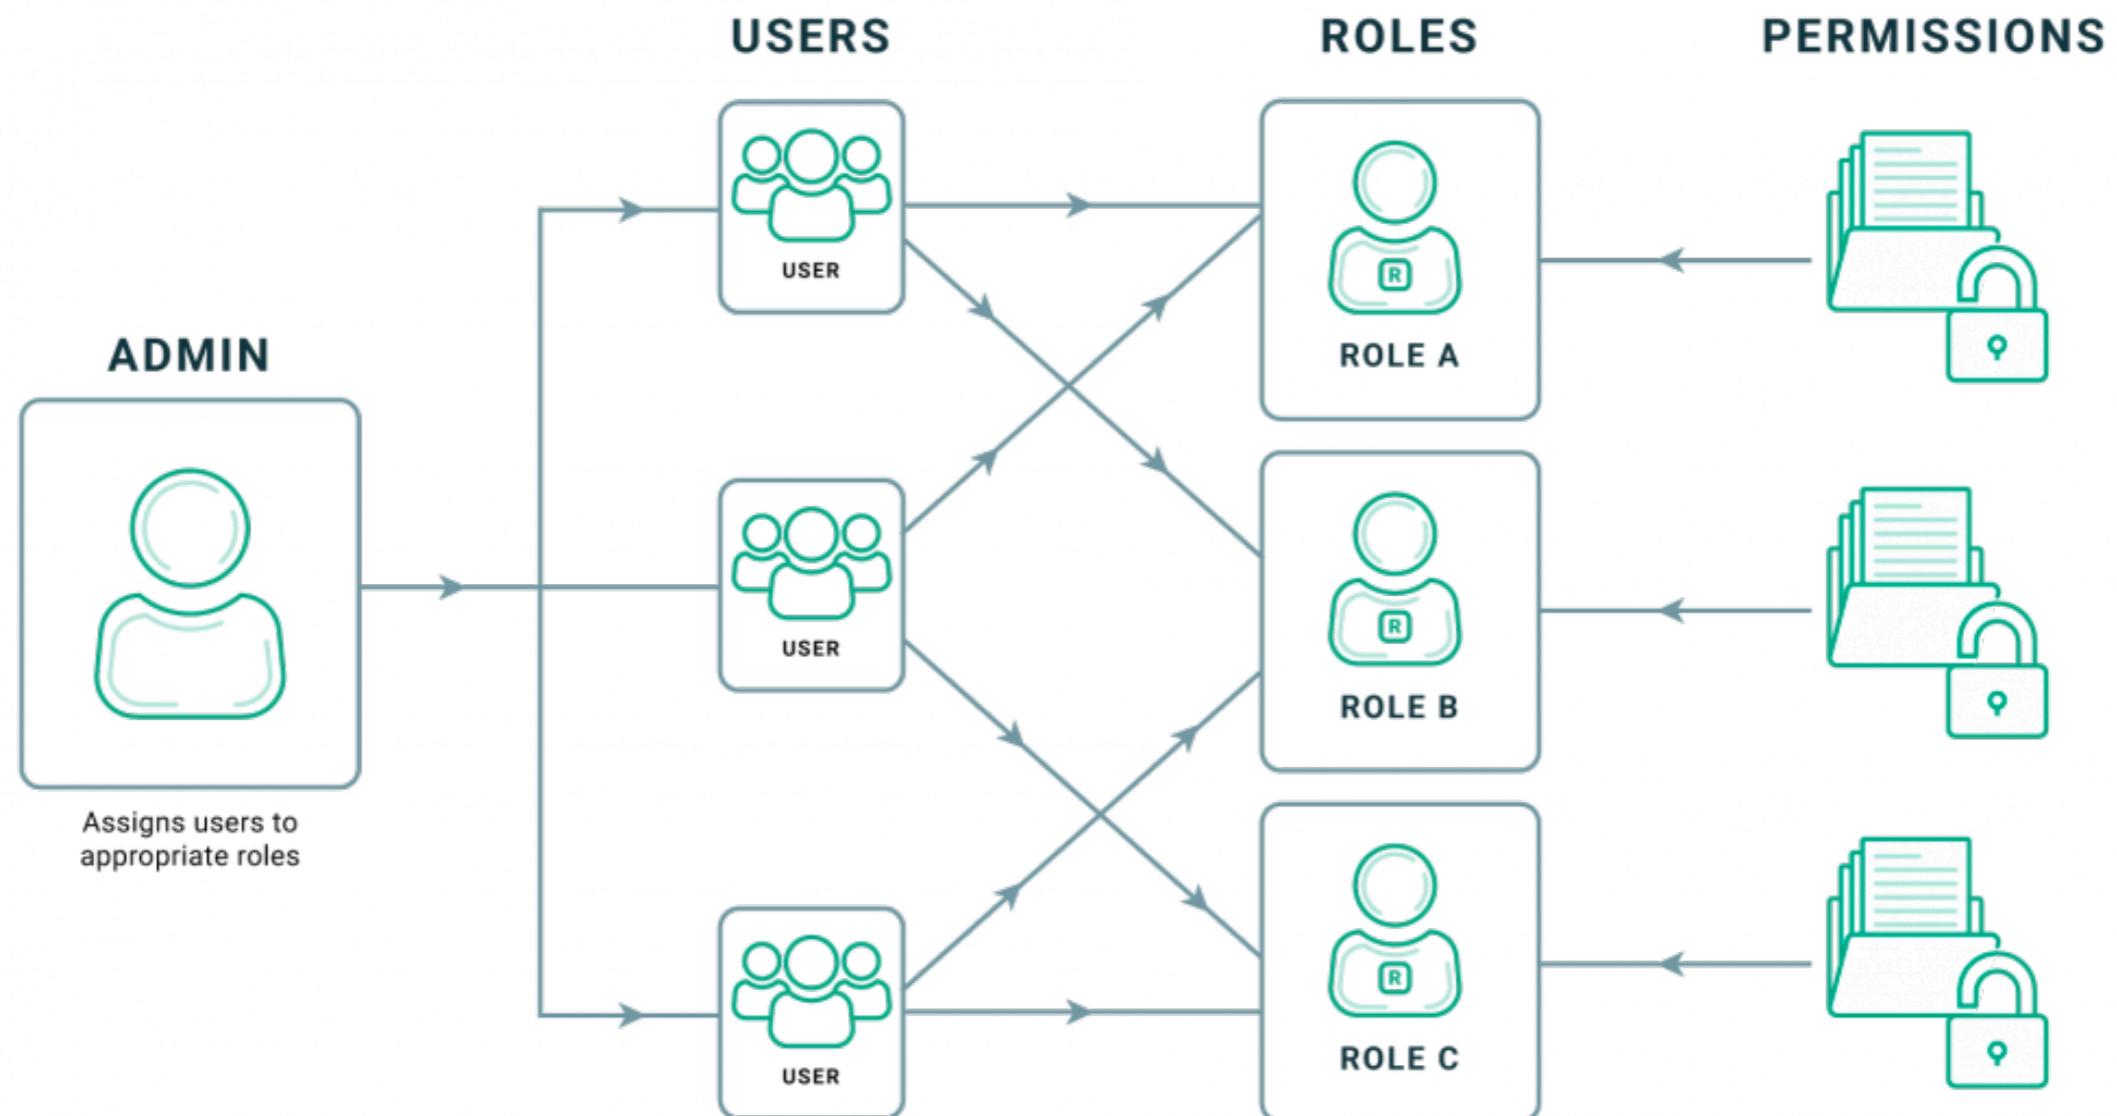
Database is secure ?

Database allows users to perform commands
without authenticating first !!

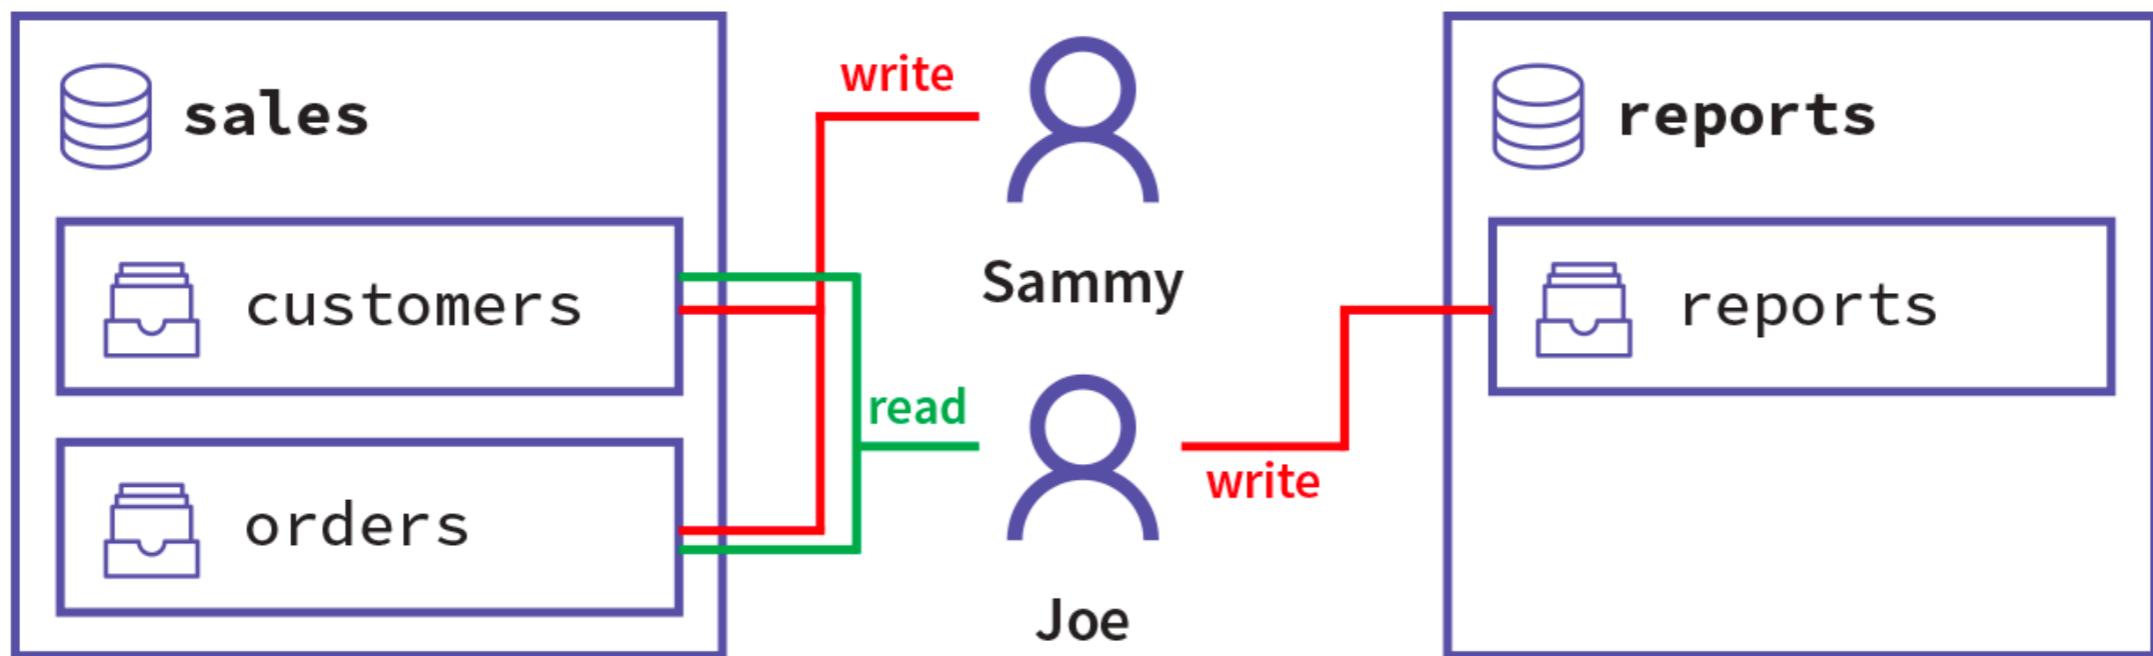


Workshop

Role-Based Access Control



Role-Based Access Control



Role-Based Access Control

Build-in roles

User-defined roles

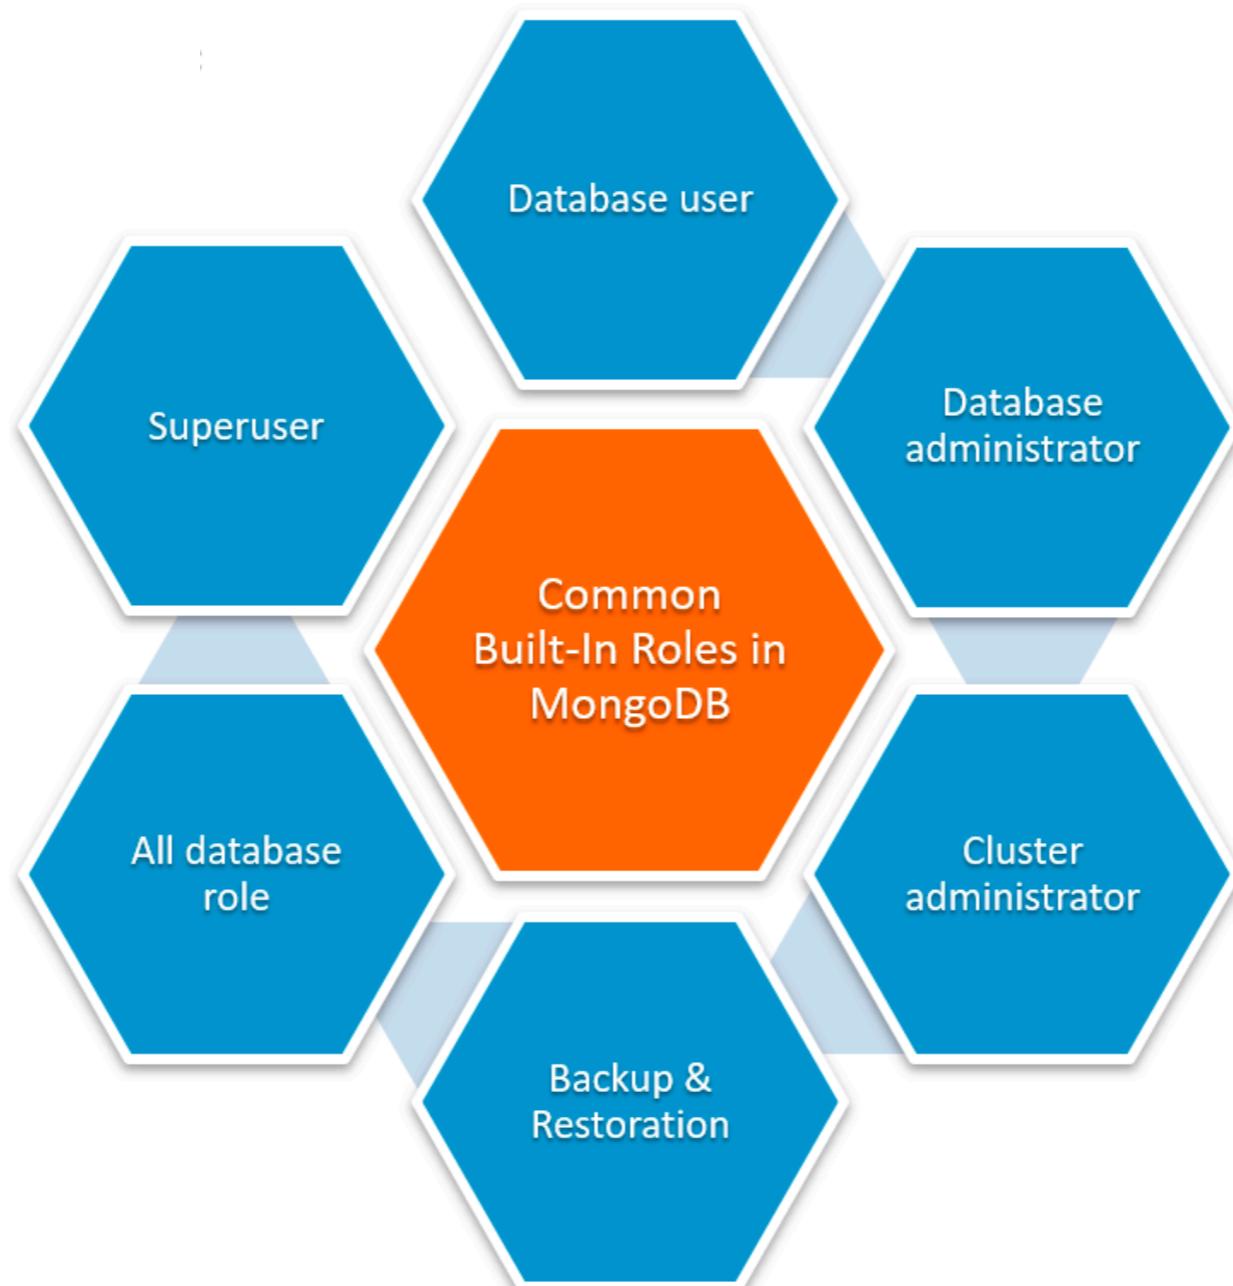
<https://docs.mongodb.com/manual/core/authorization/>

Build-in roles

MongoDB provides build-in roles to offer a set of privileges that are commonly needed in a database system

<https://docs.mongodb.com/manual/reference/built-in-roles/>

Build-in roles



Build-in roles

Group name	Roles
Database user	read readWrite
Database administration	dbAdmin userAdmin dbOwner
Cluster admin	clusterManager, clusterMonitor hostManager, clusterAdmin
Backup and restore	backup restore
All databases	readAnyDatabase, readWriteAnyDatabase userAdminAnyDatabase, dbAdminAnyDatabase
Superuser	dbOwner, userAdmin userAdminAnyDatabase

Workshop

User-defined roles

If build-in roles do not provide all the expected privileges, database administrators can define custom roles using the `createRole()`

<https://docs.mongodb.com/manual/core/security-user-defined-roles/>

Why create user-defined roles ?

Fine-grained control to user administrators

Build-in roles grant too many privileges

User admin can adhere
to the **Principle of Least Privilege (PoLP)**

Principle of Least Privilege



Principle of Least Privilege

Users should have the least privilege required
for their intended purpose

https://en.wikipedia.org/wiki/Principle_of_least_privilege

User-defined roles

Method name	Description
db.createRole()	Create a role and its privileges
db.updateRole()	Update the user-defined role
db.dropRole()	Delete a user-defined role
db.grantPrivilegesToRole()	Assigns new privileges to a role
db.revokePrivilegesFromRole()	Removes privileges from a roles

Build-in role : userAdmin

List of privilege actions

changeCustomData

grantRole

changePassword

revokeRole

createRole

setAuthenticationRestriction

createUser

viewRole

dropRole

viewUser

dropUser

Build-in role : userAdmin

Need only 3 privilege actions for this role

changeCustomData

grantRole

changePassword

revokeRole

createRole

setAuthenticationRestriction

createUser

viewRole

dropRole

viewUser

dropUser

Create new role

For all databases and collections

```
db.createRole(  
  {  
    role: "grantRevokeRolesAnyDatabase",  
    privileges: [  
      {  
        resource: { db: "", collection: "" },  
        actions: [ "grantRole", "revokeRole", "viewRole" ]  
      }  
    ],  
    roles: []  
  }  
)
```

Create new role

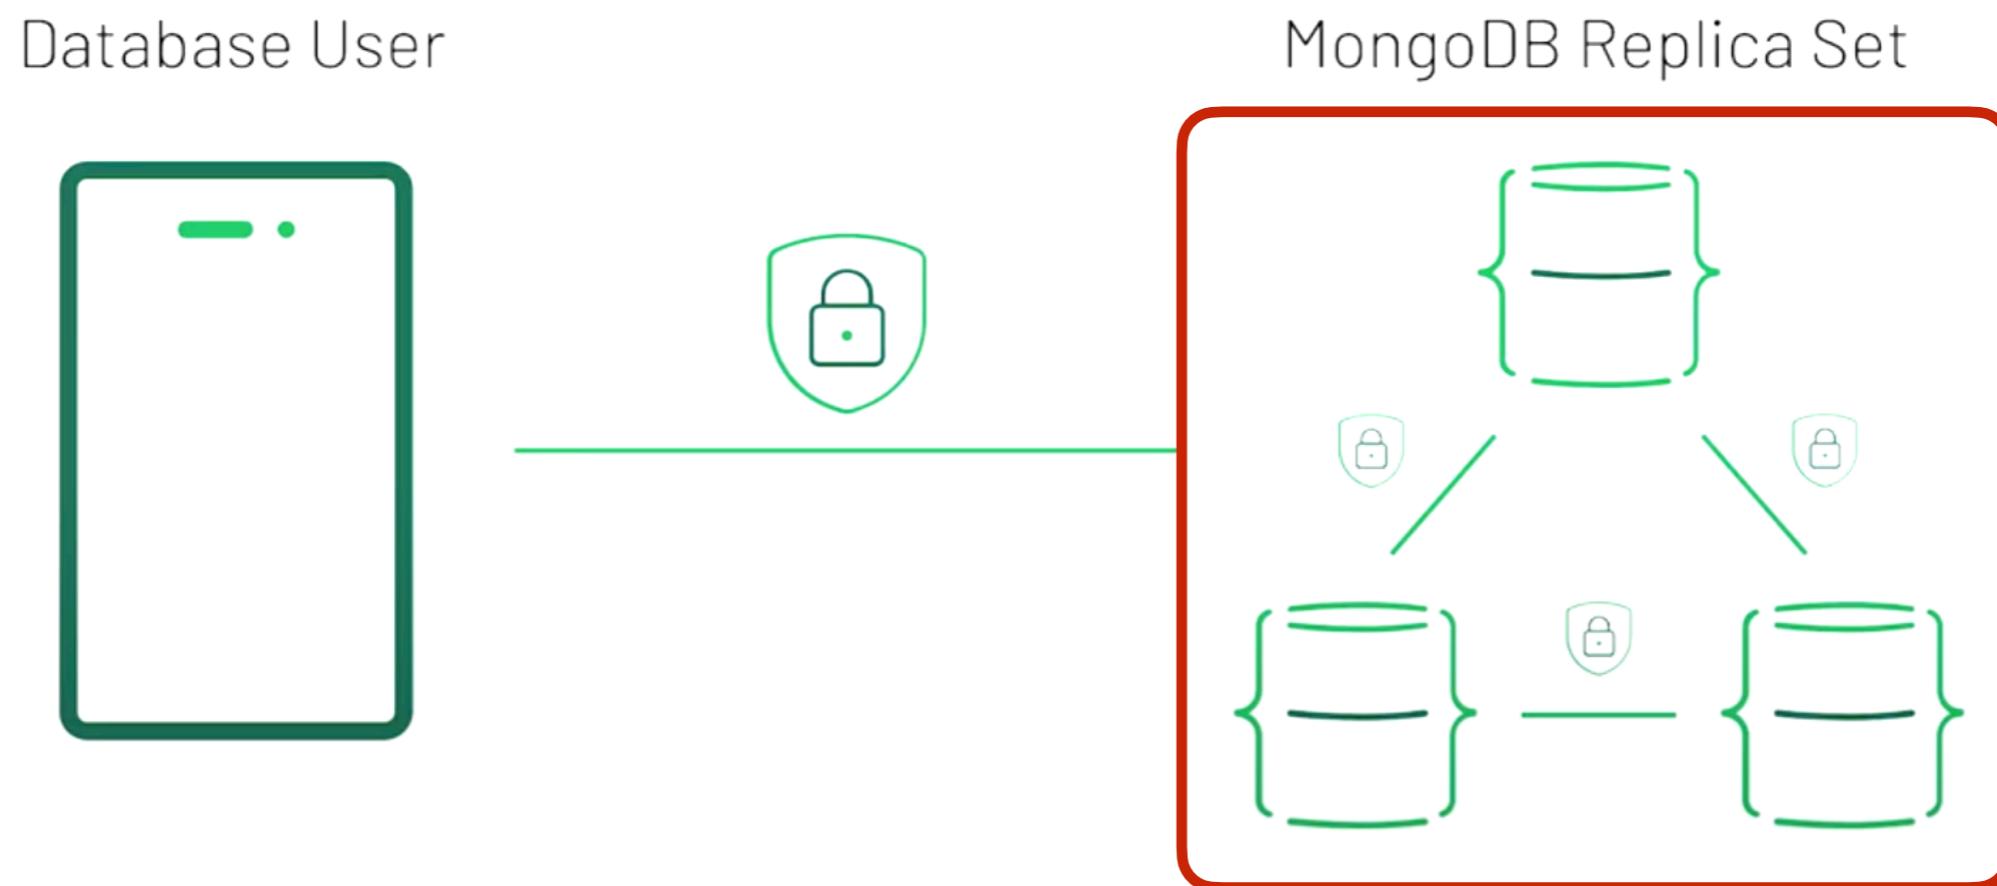
Add privilege actions to this role

```
db.createRole(  
  {  
    role: "grantRevokeRolesAnyDatabase",  
    privileges: [  
      {  
        resource: { db: "", collection: "" },  
        actions: [ "grantRole", "revokeRole", "viewRole" ]  
      }  
    ],  
    roles: []  
  }  
)
```

Workshop

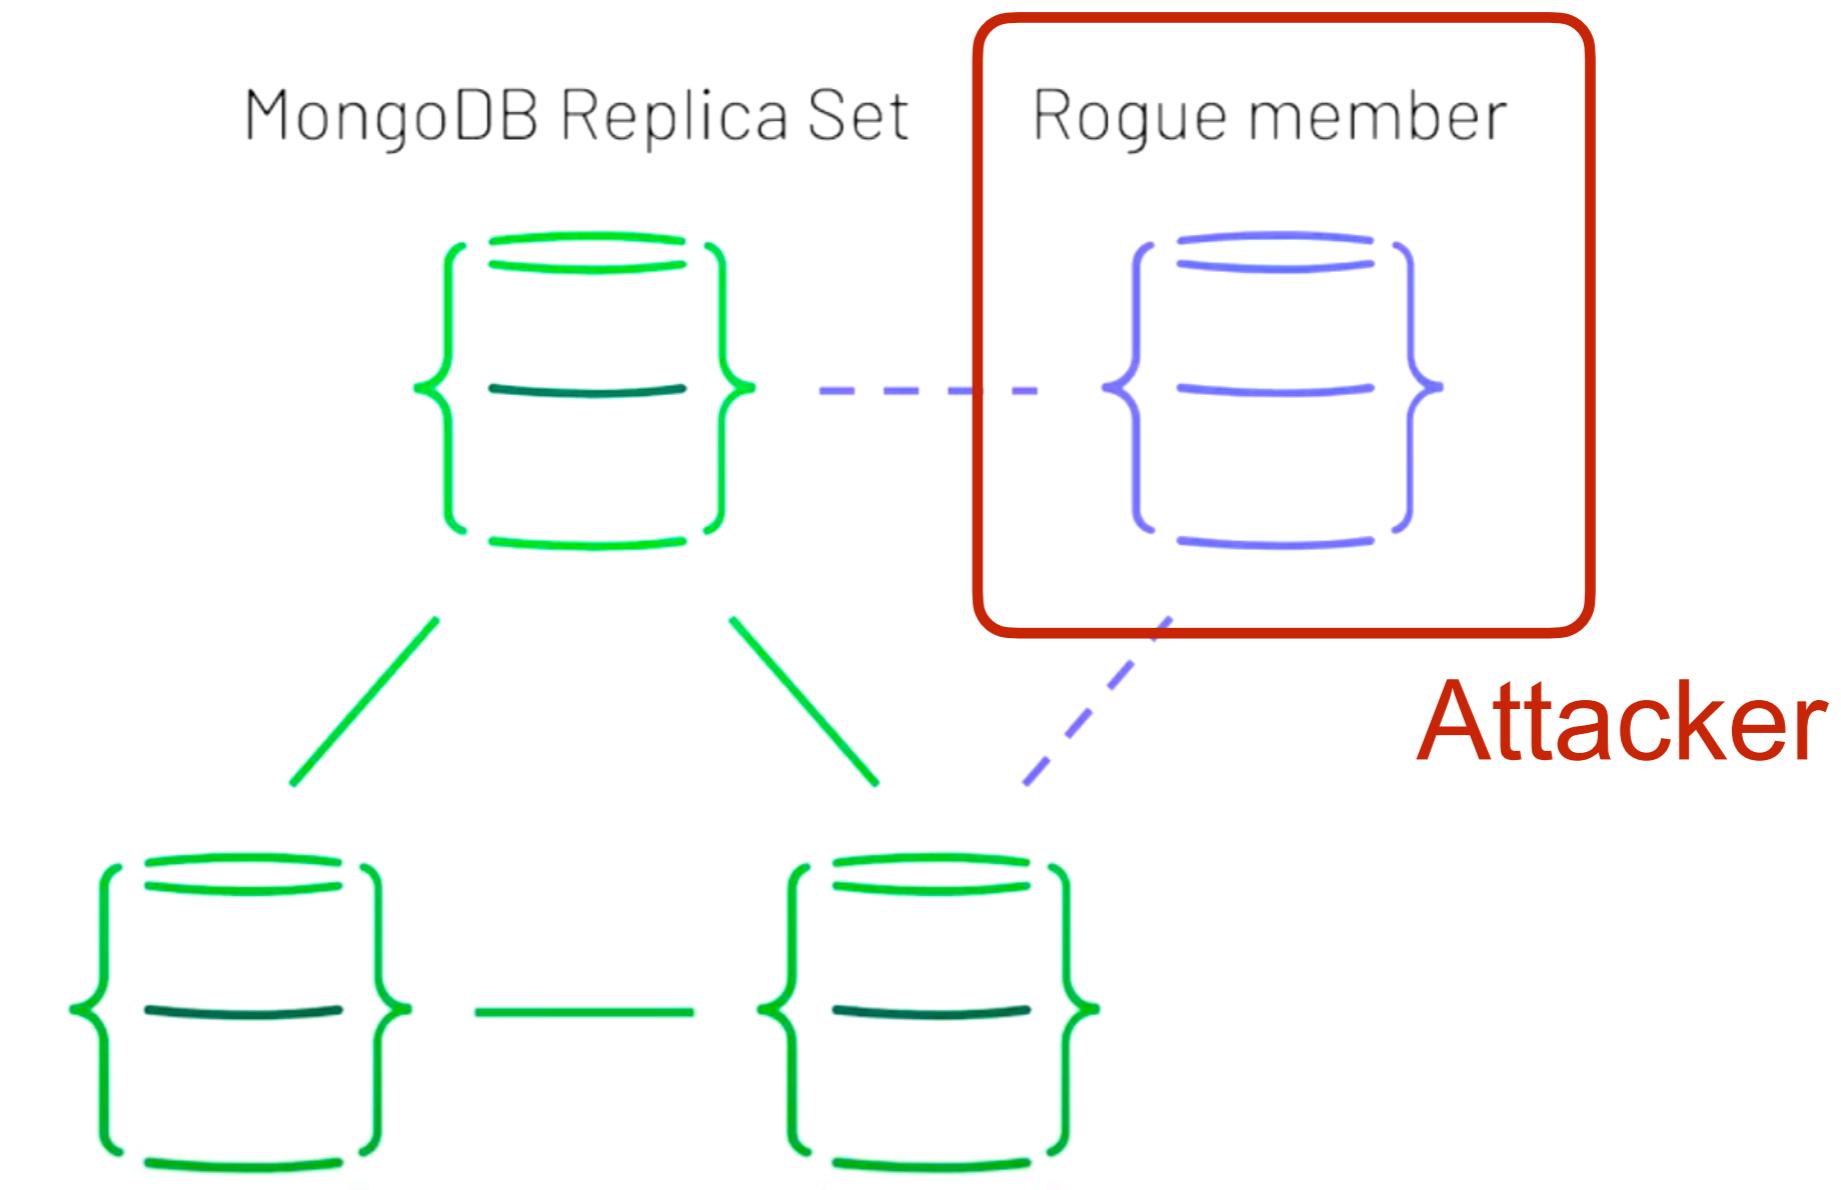
Internal Authentication ?

To verify the identity of one MongoDB instance to another



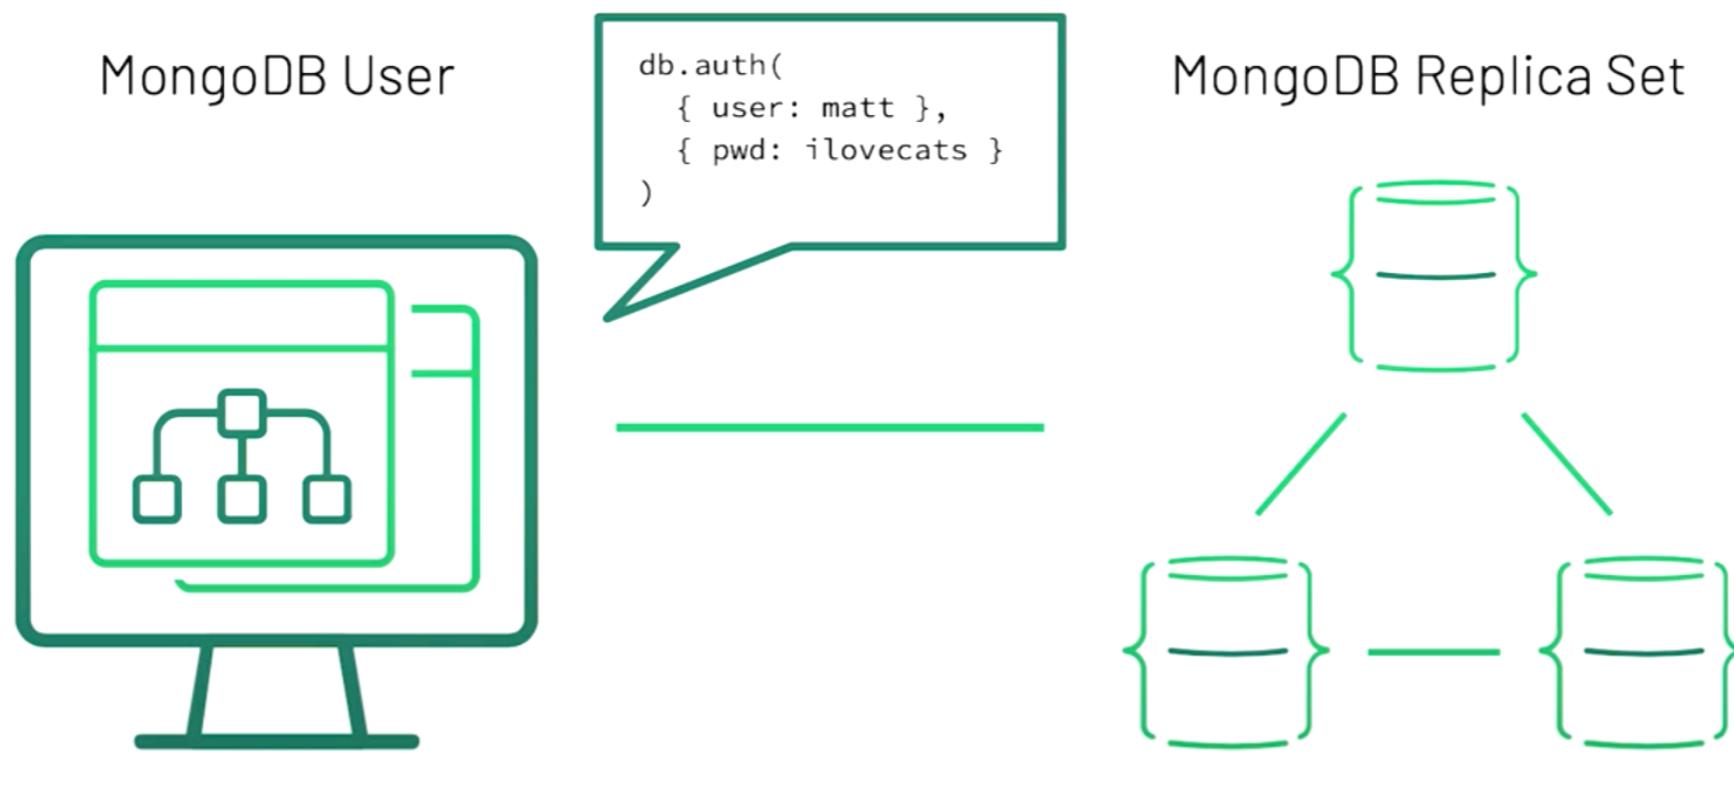
<https://docs.mongodb.com/manual/core/security-internal-authentication/>

Replica sets without internal communication



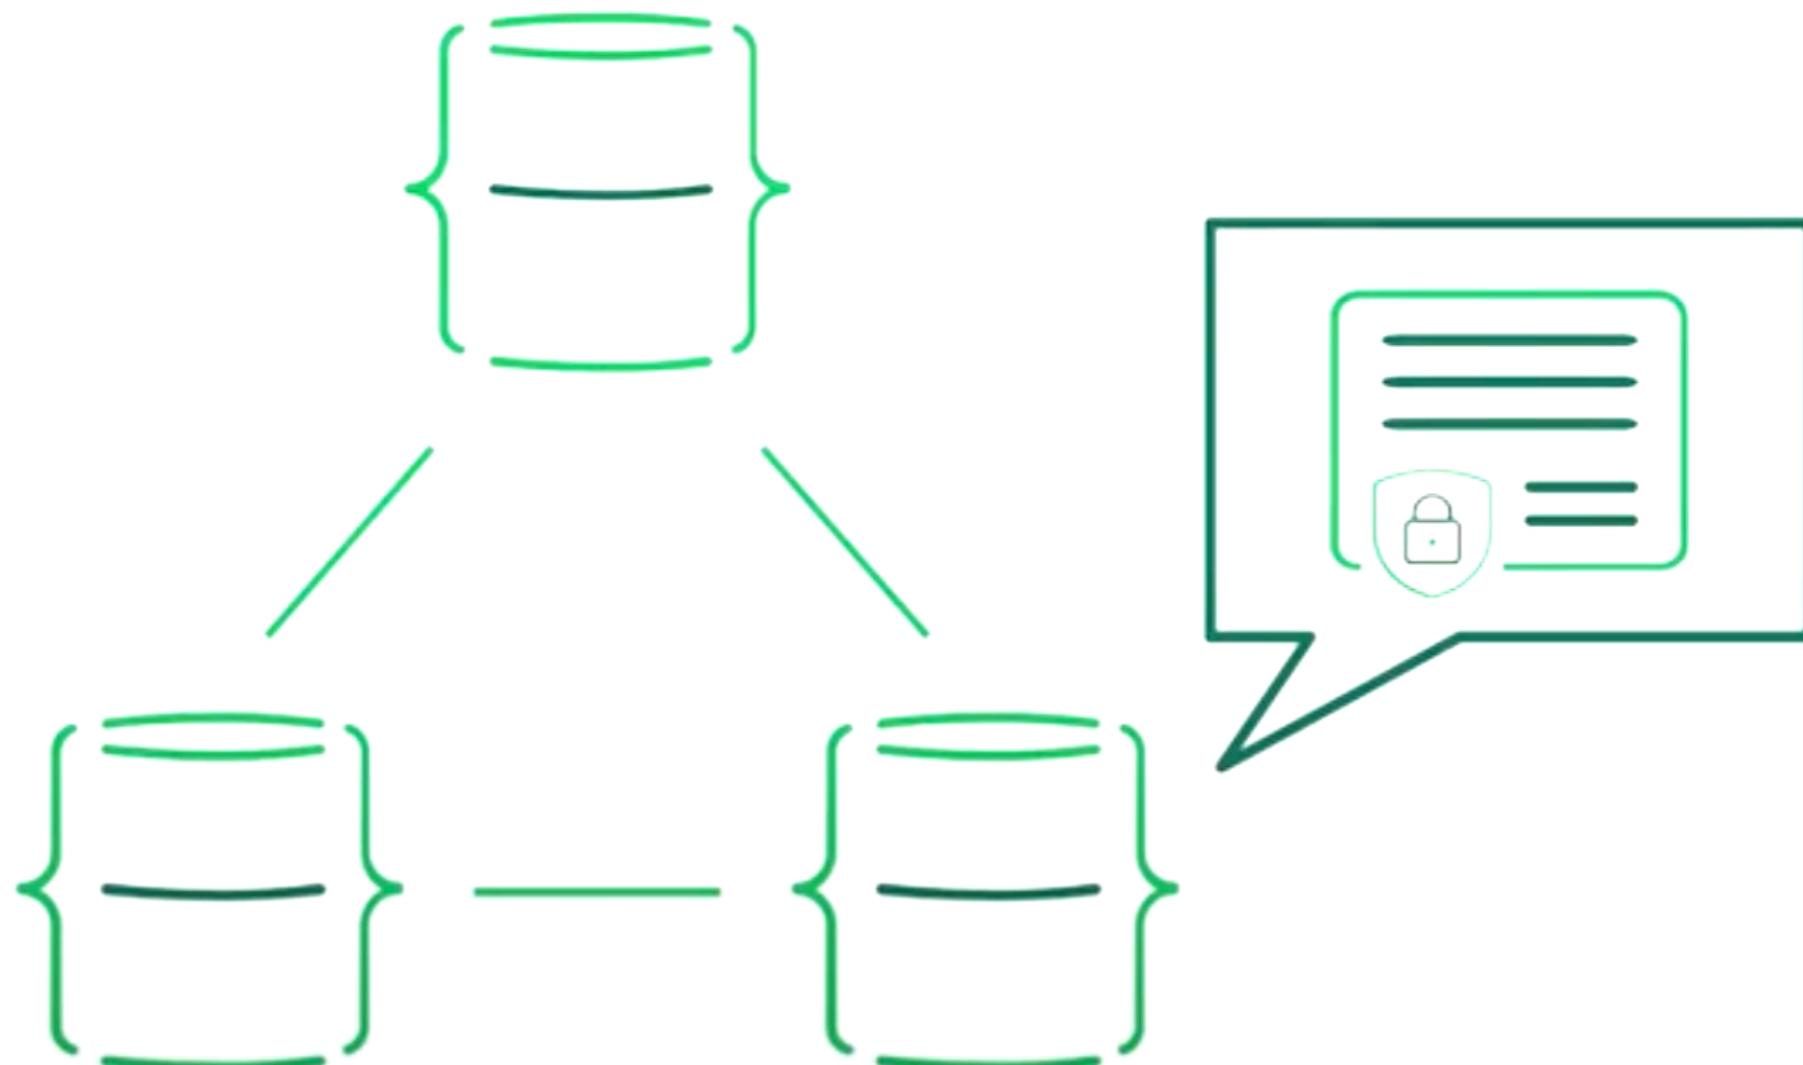
SCRAM

Use keyfile
Default authentication for MongoDB



<https://docs.mongodb.com/manual/core/security-scram/>

x.509 certificates



<https://docs.mongodb.com/manual/core/security-x.509/>

x.509 certificates

Verify members of replica set are authenticating to each other

Encrypt sent over the network

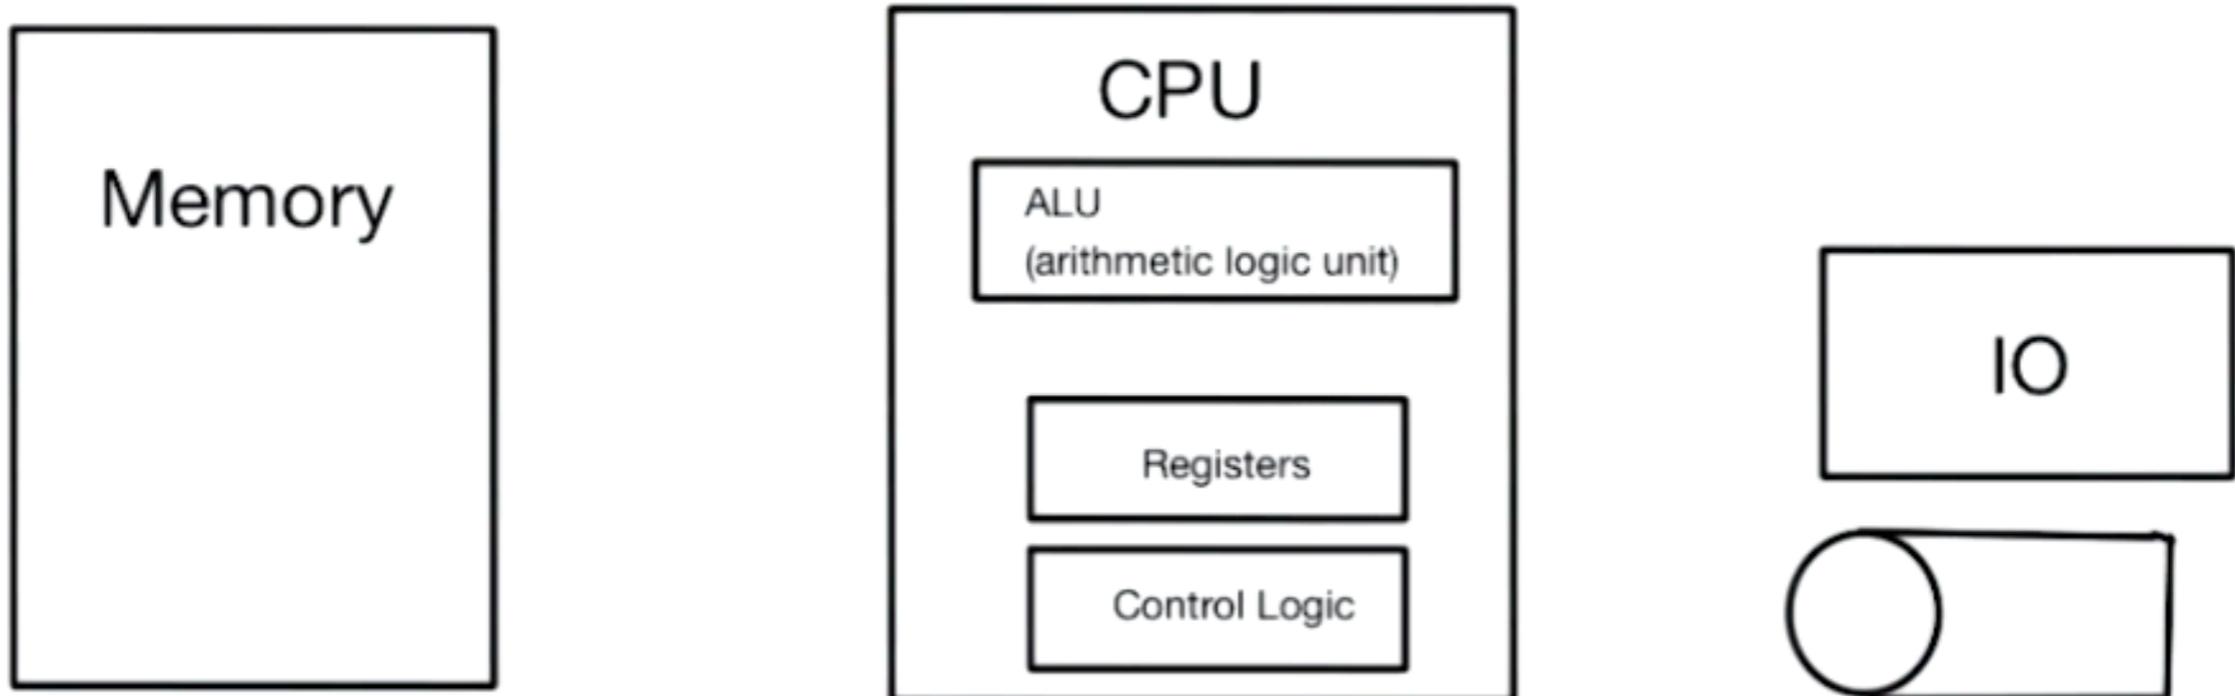
Workshop

Performance

Performance

Hardware
Configuration
Indexes
CRUD optimization
Performance on cluster

Hardware



Memory with MongoDB

Aggregation
Index traversing
Write operations
Query engine
Connections

CPU with MongoDB

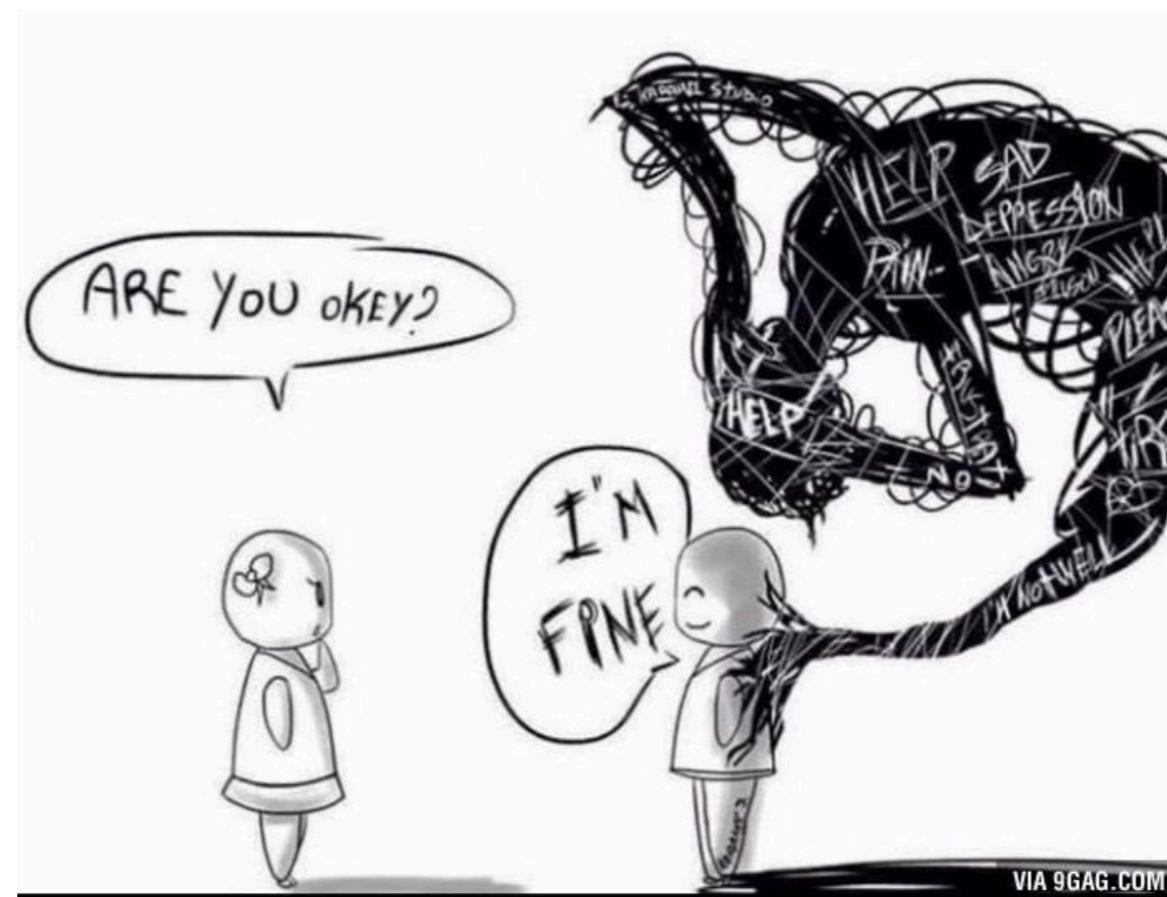
Storage engine
Concurrency model

Indexed

<https://docs.mongodb.com/manual/indexes/>

Indexes

Make query more efficient
Most impactful ways to improve query performance



Index types

Single
Compound
Multikey
Geospatial
Text
Hashed

Better indexes ?

Build the right indexes for your criteria

Learn how MongoDB picks a index

Improve query performance with **explain**

Learn about other **index types**

Avoid sorting !!

<https://docs.mongodb.com/manual/tutorial/sort-results-with-indexes>

Working with Distributed system

Consider latency

Data is spread across different nodes

Read and Write operations

Increase write performance

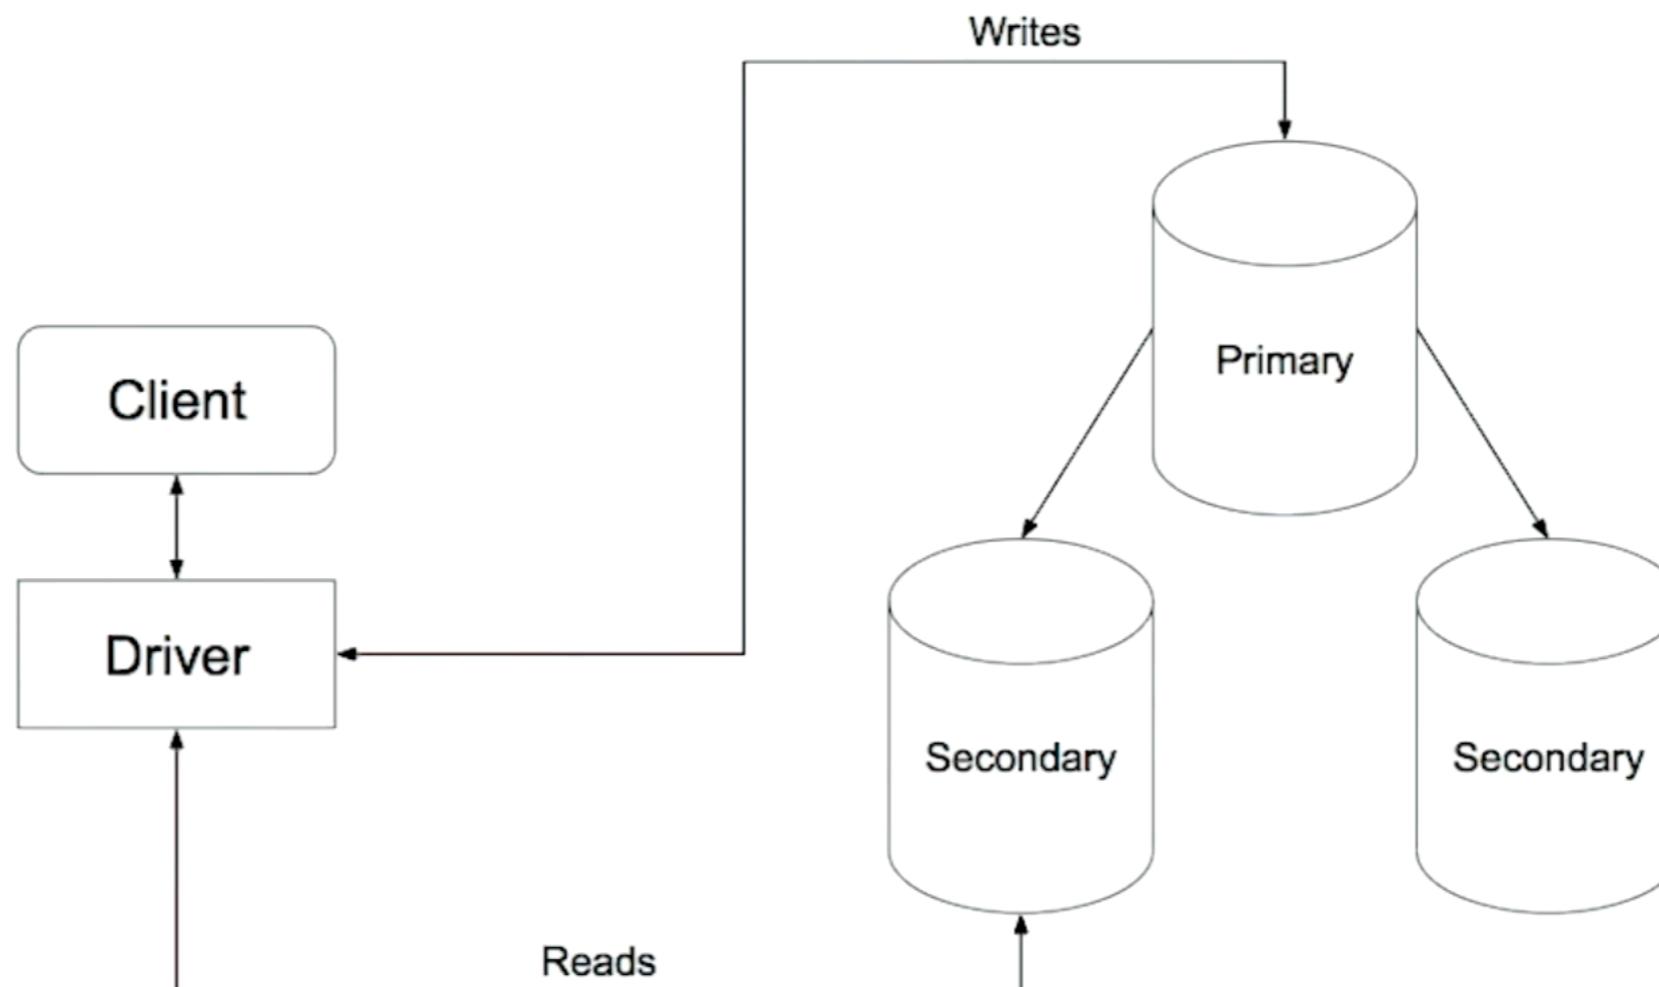
Vertical vs Horizontal scale

Shard key rules

Bulk write

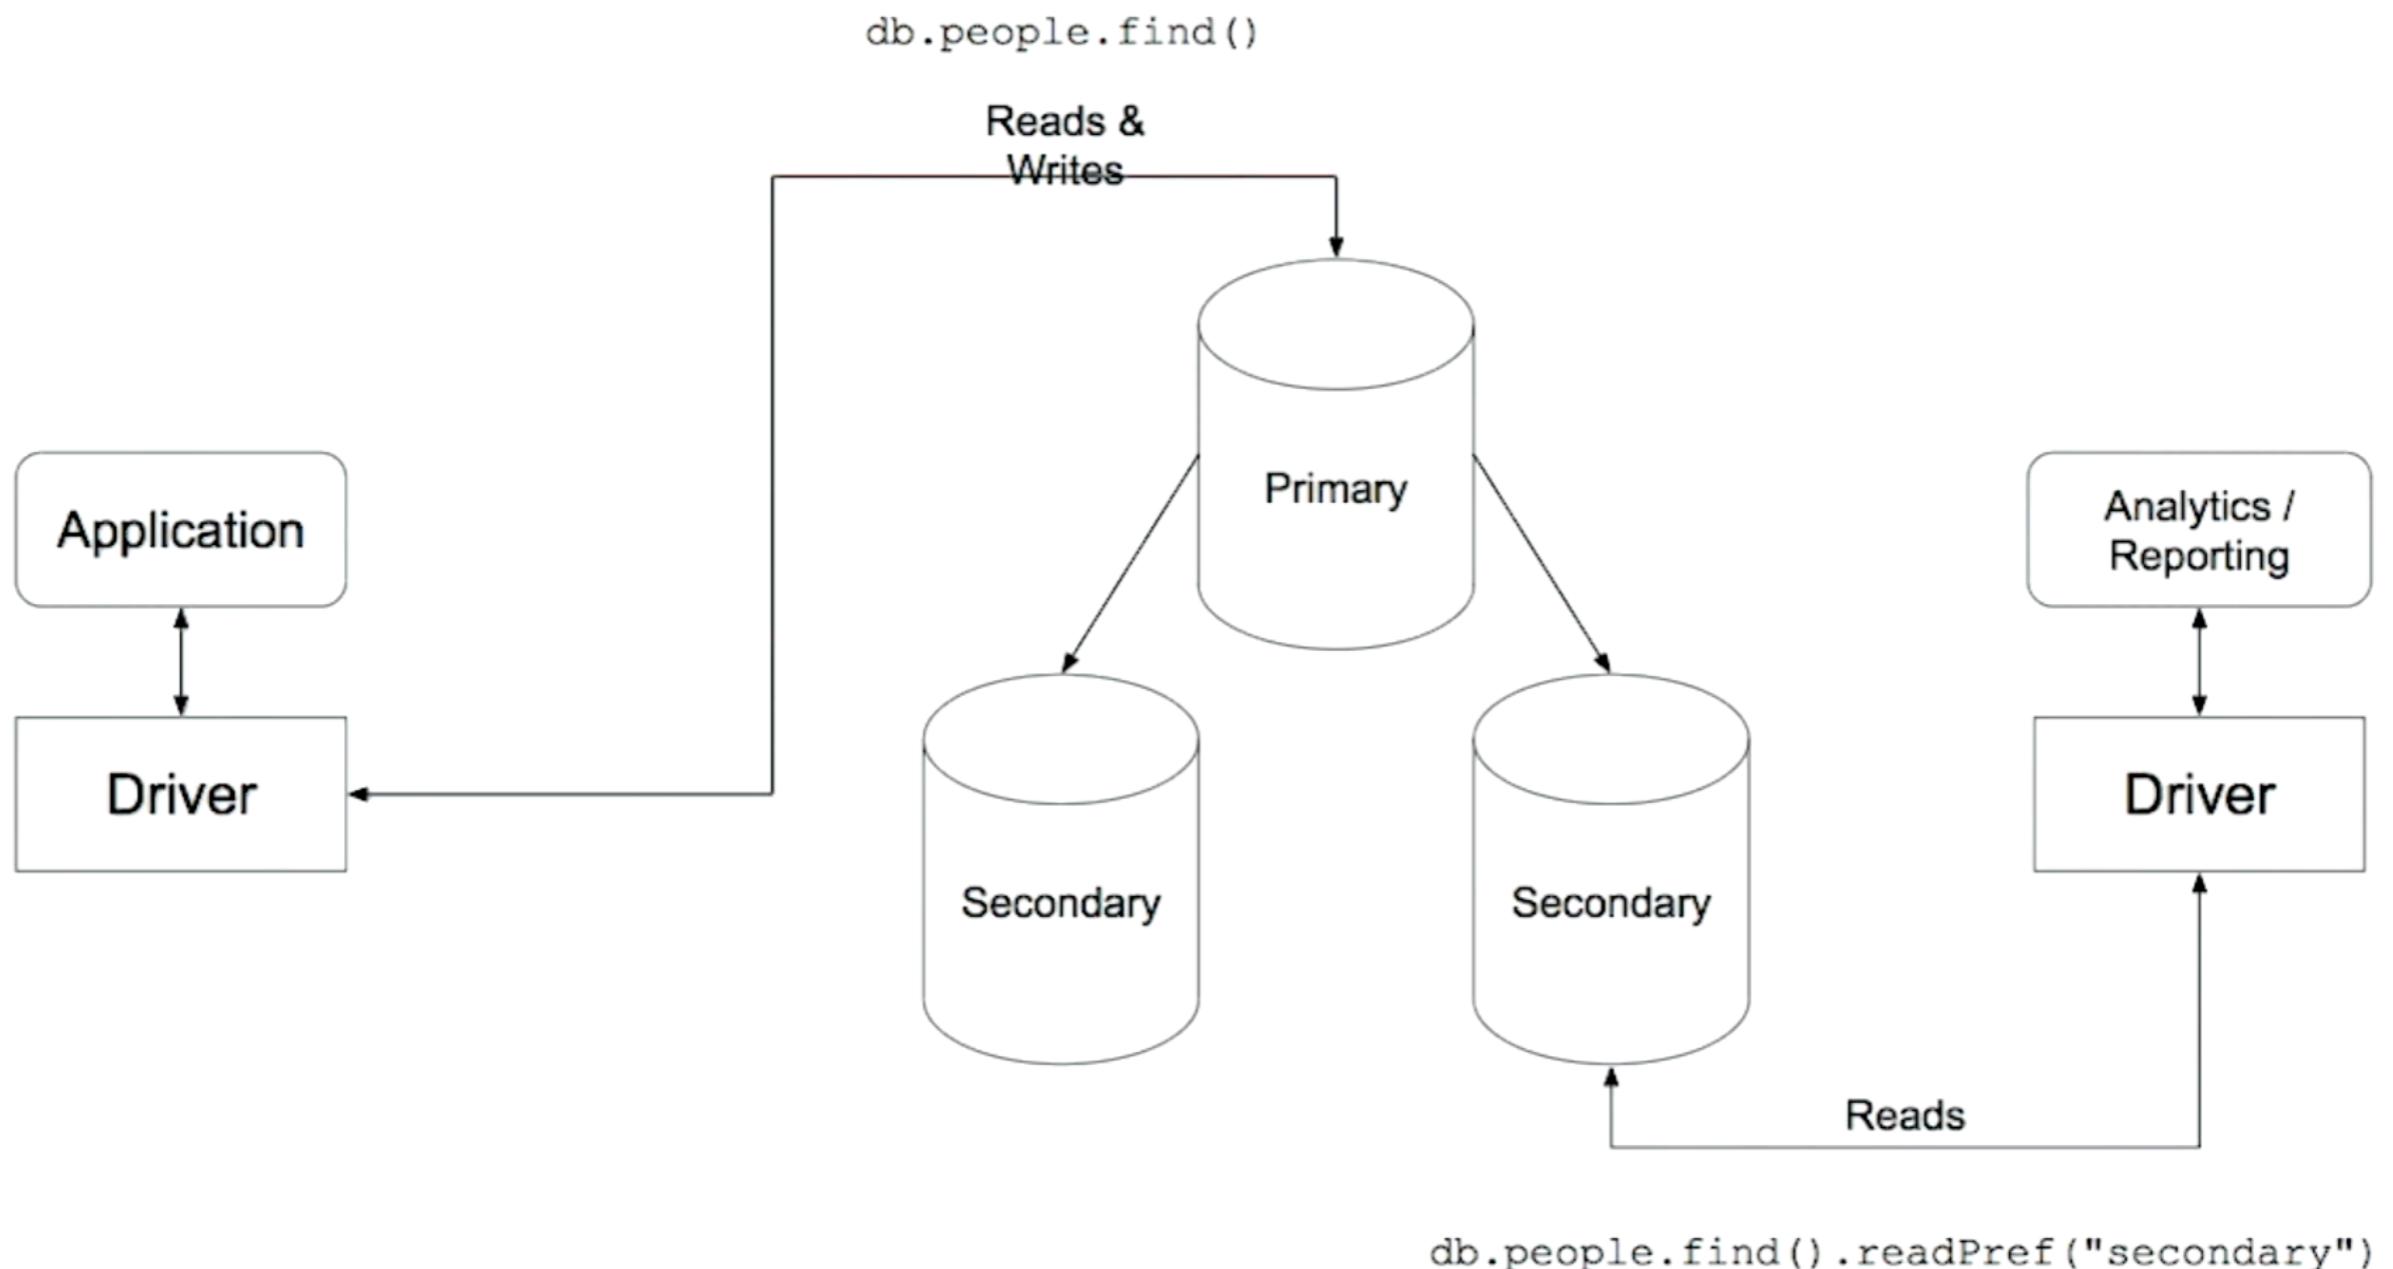
Read from secondary

Eventually consistency

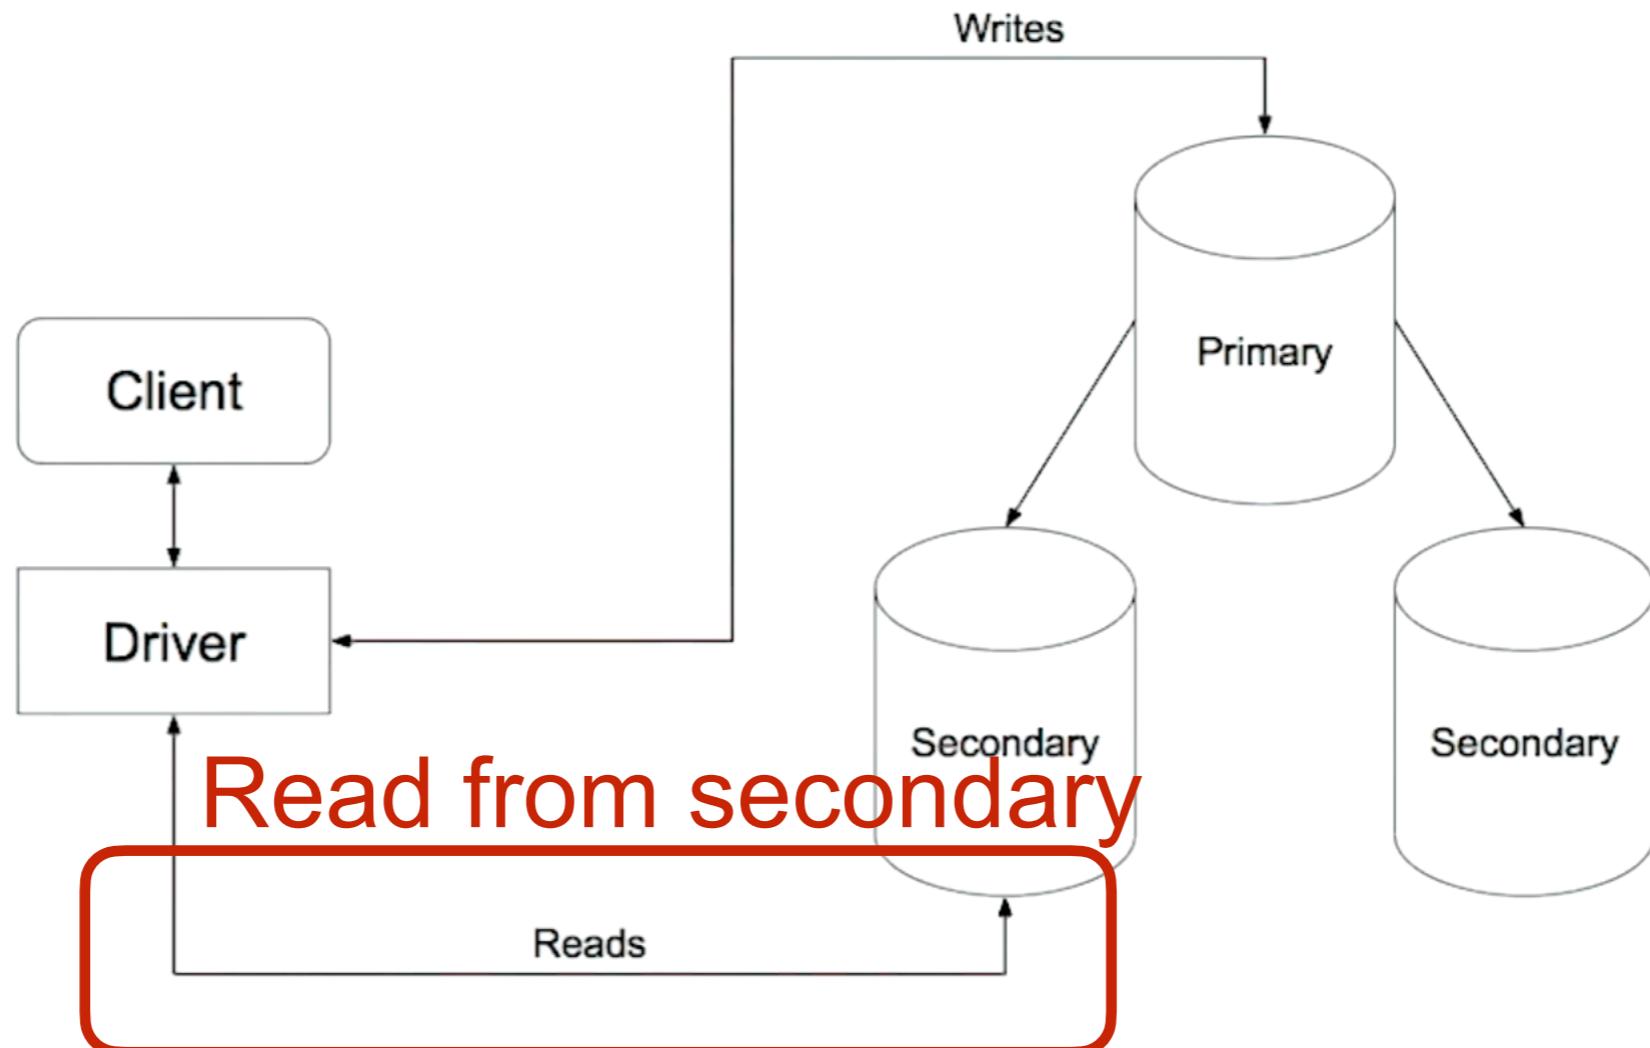


```
db.people.find().readPref("secondary")
```

Use case



Be careful !!



```
db.people.find().readPref("secondary")
```

Workshop

Q/A